

Client-Side Web Development

Joel Ross and Mike Freeman

October 28, 2017

Contents

1	Getting Setup	7
1.1	Web Browser	7
1.2	Code Editors	8
1.3	Bash (Command Line)	9
1.4	Git and GitHub	10
1.5	Node and npm	11
1.5.1	package.json	12
2	Client-Side Development	15
2.1	Client-Side File Types	16
2.2	HTTP Requests and Servers	17
3	HTML Fundamentals	21
3.1	HTML Elements	21
3.2	Nesting Elements	25
3.3	Web Page Structure	28
	The <head> Section	29
3.4	Web Page Template	30
4	CSS Fundamentals	33
4.1	Why Two Different Languages?	33
4.2	CSS Rules	34
4.3	The Cascade	39
5	Standards and Accessibility	43
5.1	Web Standards	43
5.2	Why Accessibility	45
5.3	Supporting Accessibility	47
6	More CSS	55
6.1	Compound Selectors	55
6.2	Property Values	60
7	CSS Layouts	67

7.1	Block vs. Inline	67
7.2	The Box Model	68
7.3	Changing the Flow	72
7.4	Flexbox	73
8	Responsive CSS	79
8.1	Mobile-First Design	79
8.2	Media Queries	82
9	CSS Frameworks	89
9.1	Using a Framework	91
9.2	Bootstrap	94
10	JavaScript Fundamentals	103
10.1	Programming with JavaScript	103
10.2	Running JavaScript	105
10.3	Writing Scripts	107
10.4	Variables	109
10.5	Control Structures	118
10.6	Functions	119
11	Functional Programming in JS	123
11.1	Functions ARE Variables	123
11.2	Object Functions	125
11.3	Callback Functions	126
11.4	Functional Looping	130
11.5	Pure Functions	136
12	Document Object Model (DOM)	139
12.1	The DOM API	139
12.2	DOM Manipulation	142
12.3	Listening for Events	148
13	JavaScript Libraries	153
13.1	Including a Library	154
13.2	Example: jQuery	156
14	AJAX Requests	163
14.1	AJAX	163
14.2	Fetching Data	166
14.3	Asynchronous Programming	167
15	ES6 and Beyond	173

About this Book

This book covers the the skills and techniques necessary for creating sophisticated and accessible interactive web applications. It focuses on the client-side languages, tools, and libraries that professionals use to build the web sites you use every day. It assumes a basic background in computer programming (e.g., one course in Java, and some concepts from the technical foundations of informatics). These materials were developed for the **INFO 343: Client-Side Web Development** course taught at the University of Washington Information School; however they have been structured to be an online resource for anyone who wishes to learn modern web programming techniques.

Some content has been adapted from tutorials by David Stearns.

This book is currently in **alpha** status. Visit us on [GitHub](#) to contribute improvements.



This book is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

Chapter 1

Getting Setup

This course will cover a wide variety of tools and techniques used in modern web development, including different software programs that are used to write, manage, and execute the code for your web application. This chapter explains how to install and use some of the software you will need to utilize.

Note that iSchool lab machines should have all appropriate software already installed and ready to use.

1.1 Web Browser

The first thing you'll need is a web browser for viewing the web pages you make! We recommend you install and utilize **Chrome**, which comes with an effective set of built-in developer tools that will be especially useful in this class.

- You can access the Chrome Developer tools by selecting **View > Developer > Developer Tools** from Chrome's main menu (Cmd+Option+I on a Mac, Ctrl+Shift+I on Windows). You will almost always want to have these tools open when doing web development, particularly when including interactivity via JavaScript.

Other modern browsers such as **Firefox** or **Microsoft Edge** will also function in this class and include their own versions of the required development tools. Note that different browsers may and will render code in different ways which will be discussed extensively throughout the course.

1.2 Code Editors

In order to write web code, you need somewhere to write it. There are a variety of code editors and IDEs (Integrated Development Environments) that are specialized for web development, providing syntax highlighting, code completion, and other useful functionality. There are lots of different code editors out there, all of which have slightly different appearances and features. You only need to download and use one of the following programs (we recommend **Visual Studio Code** as a default), but feel free to try out different ones to find something you like (and then evangelize about it to your friends!)

Visual Studio Code

Visual Studio Code (or VS Code; not to be confused with Visual Studio) is a free, open-source editor developed by Microsoft—yes, really. It focuses on web programming and JavaScript, though also supports many other languages and provides a number of community-built extensions for adding even more features. Although fairly new, it is updated regularly and has become my main editors for programming. VS Code is actually a stand-alone web application, so it's written in the same HTML, CSS, and JavaScript you'll learn in this course!

To install VS Code, follow the above link and Click the “Download” button to download the installer (e.g. `.exe`) file, then double-click on that to install the application.

Once you've installed VS Code, the trick to using it effectively is to get comfortable with the Command Palette. If you hit `Cmd+Shift+P`, VS Code will open a small window where you can search for whatever you want the editor to do. For example, if you type in `markdown` you can get list of commands related to Markdown files (including the ability to open up a preview). The `Format Code` option is particularly useful.

For more information about using VS Code, see the documentation, which includes videos if you find them useful. The documentation for programming in HTML, CSS, and especially JavaScript also contain lots of tips and tricks.

Atom

Atom is a text editor built by the folks at GitHub and has been gaining in popularity. It is very similar to VS Code in terms of features, but has a somewhat different interface and community. It has a similar *command-palette* to VS Code, and is arguably even nicer about editing Markdown specifically. The document you are reading was authored in Atom.

Brackets

Brackets is a coding editor authored by Adobe specialized for client-side web developers. It has some intriguing features that are not yet in Visual Studio Code, as well as possibly the nicest interface of this list.

Sublime Text

Sublime Text is a very popular text editor with excellent defaults and a variety of available extensions (though you'll need to manage and install extensions to achieve the functionality offered by other editors out of the box). While the software can be used for free, every 20 or so saves it will prompt you to purchase the full version. This is my application of choice for when I just write to write a plain text file.

1.3 Bash (Command Line)

Many of the software tools used in professional web development are used on the **command-line**: a text-based interface for controlling your computer. While the command-line is harder to learn and figure out, it is particularly effective for doing web development. Command-line automation is powerful and efficient enough to handle the dozens of repeated tasks across hundreds of different source files (split across multiple computers) commonly found in web programming. You will need to be comfortable using the command-line in order to utilize the software for this course.

While there are multiple different **command shells** (command line interfaces), this course is based on the Bash shell, which provides a particular common set of commands common to Mac and Linux machines.

On a Mac you'll want to use the built-in app called **Terminal**. You can open it by searching via Spotlight (hit Cmd (⌘) and Spacebar together, type in "terminal", then select the app to open it), or by finding it in the Applications/Utilities folder.

On Windows, we recommend using **Git Bash**, which you should install along with `git` (see below). Open this program to open the command-shell.

- Note that Windows does come with its own command-prompt, called the *DOS Prompt*, but it has a different set of commands and features. *Powershell* is a more powerful version of the DOS prompt if you really want to get into the Windows Management Framework. But Bash is more common in open-source programming like we'll be doing, and so we will be focusing on that set of commands.

- Alternatively, the 64-bit Windows 10 Anniversary update (August 2016) *does* include a beta version of an integrated Bash shell. You can access this by enabling the subsystem for Linux and then running `bash` in the command prompt. This is currently (May 2017) “beta” technology, but will suffice for our purposes if you can get it running.

This course expects you to already be familiar with basic command-line usage. For review, see The Command Line in the *INFO 201* course reader.

1.4 Git and GitHub

Professional web development involves many different people working on many different files. **git** is a collaborative version control system that provides a set of commands that allow you to manage changes to written code, particularly when collaborating with other programmers.

You will need to download and install the software. If you are on a Mac, **git** should already be installed. If you are using a Windows machine, then installing **git** will also install Git Bash, a command shell (described above).

Note that **git** is a command-line application: you can test that it is installed by running the command:

```
git --version
```

While **git** is the software used to manage versions of code, **GitHub** is a website that is used to store copies of computer code that are being managed with **git** (think “Imgur for code”).

In order to use GitHub, you’ll need to create a free GitHub account, if you don’t already have one. You should register a username that is identifiable as you (e.g., based on your name or your UW NetID). This will make it easier for others to determine out who contributed what code, rather than needing to figure out who ‘LeetDesigner2099’ is. This can be the start of a professional account you may use for the rest of your career!

- Note that you can have **git** save your GitHub password on your local machine so you don’t have to type it repeatedly. See Authenticating with GitHub from Git.

This course expects you to already be familiar with utilizing Git and GitHub. For review, see Git and GitHub, Git Branches, and Git Collaboration in the *INFO 201* course reader. Note that students in the INFO 343 course will be using GitHub and Pull Requests to turn in programming assignments.

1.5 Node and npm

Node.js (commonly just “Node”) is a command-line runtime environment for the JavaScript programming language—that is, a program that is used to *interpret* and *execute* programming instructions written in JavaScript. Although client-side development usually involves running JavaScript in the browser (see Chapter: JavaScript), Node provides a platform for installing and running a wide variety of “helper” programs that are frequently used in web development.

To install Node, visit the download page and select the installer for your operating system (you probably want the `.msi` for Windows and the `.pkg` for Mac). For this course you will want to install the **latest version** of Node (8.6+), so you should update it if you haven’t in a while. Node is a command-line application, so you can test that it is installed and available to your command shell (e.g., Terminal or Git Bash) with:

```
node --version
```

Installing Node also installs an additional command-line program called **npm**. **npm** is a **package manager**, or a program used to “manage” other programs—think of it as a command-line “app store” for developer tools and libraries. **npm** is the most common way of installing and running a large number of tools used in professional web development. At the time of writing, the **npm** “registry” lists around 500,000 different packages.

If you have an older version of Node and **npm** (before Node 8.6 or **npm** 5.1), you will need to update them. You can re-install the latest version of Node, or just upgrade **npm** using the command **npm install npm@latest -g**. See below for an explanation of this command.

Managing packages with npm

You can use the **npm** program to download and install command-line programs by name:

```
npm install -g PACKAGE-NAME
```

For example, you can install the *live-server* utility (a simple program that runs a local web server and will automatically “refresh” the browser when your code changes) using

```
npm install -g live-server
```

- Once the program is done installing, you can run it from the command-line by using the command `live-server`. This program will serve all of the content from the current directory. See Chapter 2 for details.

Importantly, note the included **-g** option. This tells **npm** that the package

should be installed **globally**, making it available across the entire computer, rather than just from a particular folder. Because you want to be able to use a command-line utility like *live-server* from any folder (e.g., for any project), command-line utilities are always installed globally with the `-g` option.

It is also possible to omit that option and install a package *locally*. For example:

```
npm install lodash
```

Will download the `lodash` code library (a set of useful JavaScript functions). This package will be placed into a new folder *in the current project directory* called `node_modules/`, and can be imported and used in the current directory's code. (It's called a local install because the package is only available to the "local" project). You will of course need to install local packages once per project.

Because node packages can be very large, and projects can have lots of them, you want to be sure to **not** commit the `node_modules/` folder to version control. Make sure that the folder is listed in your `.gitignore` file!

1.5.1 package.json

As projects become large, it is common for them to build up many *dependencies*: packages that must be installed in order for the program to work. In other words, there needs to be a certain set of packages in the project's `node_modules/` folder. *npm* is able to keep track of these dependencies by recording them in a specialized file called `package.json` that can be placed inside the project directory. A `package.json` file is a text file containing a JSON list of information about your project. For example:

```
{
  "name": "example",
  "version": "1.0.0",
  "private": true,
  "description": "A project with an example package.json",
  "main": "index.js",
  "scripts": {
    "test": "jest"
  },
  "author": "Joel Ross",
  "license": "ISC",
  "dependencies": {
    "lodash": "^4.17.4",
    "moment": "^2.18.1"
  },
  "devDependencies": {
    "html-validator": "^2.2.2"
  }
}
```

```
}  
}
```

(You can create one of these files by using the command `npm init` in the current project directory, and then following the instructions to fill in the fields).

Notice that there are two packages listed under "dependencies": `lodash` and `moment` (the `^4.17.4` indicates which version of `lodash`). You can use `npm` to automatically install all of packages listed under "dependencies" (as well as "devDependencies") using the command:

```
npm install
```

Thus using `npm install` without any arguments means “install all of the requirements that have been listed for this project”. This is a good first step *any time* you download a project or checkout a repository from GitHub.

When installing specific packages, you can have `npm` add them to the dependencies list by using the `--save` option:

```
npm install --save lodash
```

will install `lodash` locally, and list it in the `package.json` file as a dependency.

Similarly, the `--save-dev` option will instead save the package in the "devDependencies" list, which are dependencies needed only for development (writing the program's code) and not for execution (running the program).

If you want to use a *globally* installed package in your local project (e.g., have it be a dependency but not have to download and install it again), you can use `npm link` to “include” the global package locally. For example, the below will allow you to use a globally installed version of `lodash` so you don't have to download a copy for each project:

```
npm link lodash
```

Be sure to `link` any global packages before you run `npm install` so you don't any packages from `package.json` that you already have!

You can uninstall packages using `npm uninstall`, or can remove packages from the dependencies lists simply by editing the `package.json` file (e.g., with VS Code).

To sum up, you will use three commands with `npm` to install packages:

1. `npm install -g PACKAGE-NAME` to *globally* install command-line programs.
2. `npm install` to *locally* install all of the dependencies for a project you check out.
3. `npm install --save PACKAGE-NAME` to *locally* install a new code package and record it in the `package.json` file.

While *npm* is the most popular package manager (and the one utilized in this course), there are others as well. For example, **Yarn** is a package manager created by Facebook that is compatible with *npm* and is quickly growing in popularity.

Resources

Links to the recommended software are collected here for easy access:

- Chrome
- git (and Git Bash)
 - GitHub (sign up)
 - optional: Bash on Windows
- Visual Studio Code
- Node.js (and npm)
 - npm documentation

Chapter 2

Client-Side Development

Web development is the process of implementing (programming) web sites and applications that users can access over the internet. However, the internet is a network involving *many* different computers all communicating with one another. These computers can be divided into two different groups: **servers** store (“host”) content and provide (“serve”) it to other computers, while **clients** request that content and then present it to the human users.

Consider the process of viewing a basic web page, such as the Wikipedia entry on Informatics. In order to visit this page, the user types the web address (<https://en.wikipedia.org/wiki/Informatics>) into the URL bar, or clicks on a link to go to the page. In either case, user’s computer is the **client**, and their browser takes that address or link and uses it to create an **HTTP Request**—a *request* for data sent following the *HyperText Transfer Protocol*. This request is like a letter asking for information, and is sent to a different computer: the **web server** that contains that information.




Figure 2.1: A diagram of client/server communication.

The web server will receive this request, and based on its content (e.g., where it was sent to) will decide what information to send as a **response** to the client. In general, this response will be made up of lots of different files: the text content of the web page, styling information (font, color) for how it should look, instructions for respond to user interaction (button clicks), images or other assets to show, and so forth.

The client's web browser will then take all of these different files in the response and use them to *render* the web page for the user to see: it will determine what text to show, what font and color to make that text, where to put the images, and is ready to do something else when the user clicks on one of those images. Indeed, a web browser is just a computer program that is able to send HTTP requests on behalf of the user, and then render the resulting response.

Given this interaction, **client-side web development** involves implementing programs (writing code) that is interpreted by the *browser*, and so is executed by the client. It is authoring the code that is sent in the server's response. This code specifies how websites should appear and how the user should interact with them. On the other hand, **server-side web development** involves implementing programs that the *server* uses to determine which client-side code is delivered. As an example, a server-side program contains the logic to determine which cat picture should be sent along with the request, while a client-side program contains the logic about where and how that picture should appear on the page.

This course focuses on *client-side web development*, or developing programs that are executed by the browser (generally as a response to a web server request). While we will cover how client-side programs can interact with a server, many of the concepts discussed here can also be run inside a browser without relying on an external server (called “running locally”).

2.1 Client-Side File Types

It is the web browser's job to interpret and render the source code files sent by a server as part of an HTTP response. As a client-side web programmer, your task is to write this source code for the browser to interpret. There are multiple different types of source code files, including:

- **.html** files containing code written in HTML (HyperText Markup Language). This code will specify the textual and *semantic* content of the web page. See the chapter HTML Fundamentals for details on HTML.
- **.css** files containing code written in CSS (Cascading Style Sheets). This code is used to specify styling and *visual appearance* properties (e.g., color and font) for the HTML content. See the chapter CSS Fundamentals for details on CSS.

- **.js** files containing code written in JavaScript. This code is used to specify *interactive behaviors* that the website will perform—for example, what should change when the user clicks a button. Note that JavaScript code are “programs” that sent over by the web server as part of the response, but are *executed* on the client’s computer. See the chapter JavaScript Fundamentals for details on JavaScript.

HTTP responses may also include additional **asset** files, such as images (**.png**, **.jpg**, **.gif**, etc), fonts, video or music files, etc.

2.2 HTTP Requests and Servers

Modern web browsers are able to *render* (interpret and display) all of these types of files, combining them together into the modern, interactive web pages you use every day. In fact, you can open up almost any file inside a web browser, such as by right-clicking on the file and selecting “Open With”, or dragging the file into the browser program. HTML files act as the basis for web pages, so you can open a **.html** file inside your web browser by double-clicking on it (the same way you would open a **.docx** file in MS Word):




Figure 2.2: An very simple HTML file. See Chapter 3 for source code.

Consider the URL bar in the above browser. The URL (Uniform Resource Locator) is actually a specialized version of a **URI (Uniform Resource Identifier)**. URIs act a lot like the *address* on a postal letter sent within a large organization such as a university: you indicate the business address as well as the department

and the person, and will get a different response (and different data) from Alice in Accounting than from Sally in Sales.

- Note that the URI is the **identifier** (think: variable name) for the resource, while the **resource** is the actual *data* value (the file) that you want to access.

Like postal letter addresses, URIs have a very specific format used to direct the request to the right resource.



Figure 2.3: The format (schema) of a URI.

The parts of this URI format include:

- **scheme** (also **protocol**): the “language” that the computer will use to send the request for the resource (file).

In the example browser window above, the protocol is **file**, meaning that the computer is accessing the resource from the file system. When sending requests to web servers, you would use **https** (secure HTTP). *Don’t use insecure http!*

Web page *hyperlinks* often include URIs with the **mailto** protocol (for email links) or the **tel** for phone numbers.

- **domain**: the address of the web server to request information from. You can think of this as the recipient of the request letter.

In the browser window example, there is no domain because the **file** protocol doesn’t require it.

- **port** (*optional*): used to determine where to connect to the web server. By default, web requests use port 80, but some web servers accept connections on other ports—e.g., 8080, 8000 and 3000 are all common on development servers, described below.
- **path**: which resource on that web server you wish to access. For the **file** protocol, this is the *absolute path* to the file. But even when using **https**, for many web servers, this will be the *relative path* to the file, starting from the “root” folder of that server (which may not be the computer’s root folder). For example, if a server used `/Users/joelross/` as its root, then the *path* to the above HTML file would be `Desktop/index.html` (e.g., `https://domain/Desktop/index.html`).

Important! If you don't specify a path, most web servers will serve the file named `index.html` from that server's root folder (i.e., the path “defaults” to `index.html`). As such, this is the traditional name for the HTML file containing a websites home page.

As in any program, you should always use **relative** paths in web programming, and these paths are frequently (but not always!) relative to the web server's *root folder*.

- **query** (*optional*): extra **parameters** (arguments) included in the request about what resource to access. The leading `?` is part of the query.
- **fragment** (*optional*): indicates which part (“fragment”) of the resource to access. This is used for example to let the user “jump” to the middle of a web page. The leading `#` is part of the query.

Development Servers

As noted above, it is possible to request a `.html` file (open a web page) using the `file` protocol by simply opening that file directly in the browser. This works fine for testing most client-side programs. However, there are a few client-side interactions that for security reasons only work if a web page is requested from a web server (e.g., via the `http` protocol).

For this reason, it is recommended that you develop client-side web applications using a **local development web server**. This is a web server that you run from your own computer—your machine acts as a web server, and you use the browser to have your computer send a request *to itself* for the webpage. Think about mailing yourself a letter. Development web servers can help get around cross-origin request restrictions, as well as offer additional benefits to speed development—such as *automatically reloading the web browser when the source code changes*.

There are many different ways to run a simple development server from the command-line (such as using the Python `http.server` module). These servers, when started, will “serve” files using the current directory as the “root” folder. So again, if you start a server from `/Users/joelross`, you will be able to access the `Desktop/index.html` file at `http://127.0.0.1:port/Desktop/index.html` (which port will depend on which development server you use).

- The address `127.0.0.1` is the IP address for `localhost` which is the domain of your local machine (the “local host”). Most development servers, when started, will tell you the URL for the server's root directory.
- Most commonly, you will want to start the web server from the root directory of your *project*, so that the relative path `index.html` finds the file you expect.

- You can usually stop a command-line development server with the universal `ctrl-c` cancel command. Otherwise, you'll want to leave the server running in a background terminal as long as you are working on your project.

If you use the recommended **live-server** utility, it will open a web browser to the root folder and *automatically reload the page* whenever you **save** changes to a file in that folder. This will make your life much, much better.

Chapter 3

HTML Fundamentals

A webpage on the internet is simply a set of files that the browser *renders* (shows) in a particular way, allowing you to interact with it. The most basic way to control how a browser displays content (e.g., words, images, etc) is by *encoding* that content in HTML.

HTML (**H**yper**T**ext **M**arkup **L**anguage) is a language that is used to give meaning to otherwise plain text, which the browser can then use to determine how to display that text. HTML is not a programming language but rather a *markup language*: it adds additional details to information (like notes in the margin of a book), but doesn't contain any logic. HTML is a “hypertext” markup language because it was originally intended to mark up a document with hyperlinks, or links to other documents. In modern usage, HTML describes the **semantic meaning** of content: it marks what content is the a *heading*, what content is a *paragraph*, what content is a *definition*, what content is an *image*, what content is a *hyperlink*, and so forth.

- HTML serves a similar function to Markdown, but is much more expressive and powerful.

This chapter provides an overview and explanation of HTML's syntax (how to use it to annotate content). HTML's syntax is very simple, and generally only takes someone a few days to learn—though using it effectively can require more practice.

3.1 HTML Elements

HTML content is normally written in `.html` files. By using the `.html` extension, your editor, computer, and browser should automatically know that this file will contain content marked up in HTML.

As mentioned in Chapter 2, most web servers will by default serve a file named **index.html**, and so that filename is traditionally used for a website’s home page.

As with all programming languages, **.html** files are really just plain text files with a special extension, so can be created in any text editor. However, using a coding editor such as VS Code provides additional helpful features that can speed up your development process.

HTML files contain the **content** of your web page: the text that you want to show on the page. This content is then annotated (marked up) by surrounding it with **tags**:



Figure 3.1: Basic syntax for an HTML element.

The **opening/start tag** comes before the content and tell the computer “I’m about to give you content with some meaning”, while the **closing/end tag** comes after the content to tell the computer “I’m done giving content with that meaning.” For example, the `<h1>` tag represents a top-level heading (equivalent to one `#` in Markdown), and so the open tag says “here’s the start of the heading” and the closing tag says “that’s the end of the heading”.

Tags are written with a less-than symbol `<`, then the name of the tag (often a single letter), then a greater-than symbol `>`. An *end tag* is written just like a *start tag*, but includes a forward slash `/` immediately after the less-than symbol—this indicates that the tag is closing the annotation.

- HTML tag names are not case sensitive, but you should always write them in all lowercase.
- Line breaks and white space around tags (including indentation) is ignored. Tags may thus be written on their own line, or *inline* with the content. These two uses of the `<p>` tag (which marks a *paragraph* of content) are equivalent:

```
<p>  
  The itsy bitsy spider went up the water spout.  
</p>
```

```
<p>The itsy bitsy spider went up the water spout.</p>
```

Taken together, the tags and the content they *contain* are called an **HTML Element**. A website is made of a bunch of these elements.

Some Example Tags

The HTML standard defines lots of different elements, each of which marks a different meaning for the content. Common elements include:

- `<h1>`: a 1st-level heading
- `<h2>`: a 2nd-level heading (and so on, down to `<h6>`)
- `<p>`: a paragraph of text
- `<a>`: an “anchor”, or a hyperlink
- ``: an image
- `<button>`: a button
- ``: emphasized content. Note that this doesn’t mean *italic* (which is not semantic), but *emphasized* (which is semantic). The same as `_text_` in Markdown.
- ``: important, strongly stated content. The same as `**text**` in Markdown
- ``: an unordered list (and `` is an ordered list)
- ``: a list item (an item in a list)
- `<table>`: a data table
- `<form>`: a form for the user to fill out
- `<svg>`: a Scalable Vector Graphic (a “coded” image)
- `<circle>`: a circle (in an `<svg>` element)
- `<div>`: a division (section) of content. Also acts as an empty *block* element (followed by a line break)
- ``: a span (section) of content. Also acts as an empty *inline* element (not followed by a line break)

Comments

As with every programming language, HTML includes a way to add comments to your code. It does this by using a tag with special syntax:

```
<!-- this is a comment -->  
<p>this is is not a comment</p>
```

Because that syntax is somewhat awkward to type, most source-code editors will let you comment-out the currently highlighted text by pressing `cmd + /` (or `ctrl + /` on Windows). If you’re using a code editor, try placing your cursor on a line and using that keyboard command to comment and un-comment the line.

Comments can appear anywhere in the file. Just as in other languages, they are ignored by any program reading the file (with a few interesting exceptions), but they do remain in the page and are visible when you view the page source.

Attributes

The start tag of an element may also contain one or more **attributes**. These are similar to attributes in object-oriented programming: they specify *properties*, options, or otherwise add additional meaning to an element. Like named parameters in R or HTTP query parameters, attributes are written in the format `attributeName=value` (no spaces are allowed around the `=`); values of attributes are almost always strings, and so are written in quotes. Multiple attributes are separated by spaces:

```
<tag attributeA="value" attributeB="value">
  content
</tag>
```

For example, a hyperlink anchor (`<a>`) uses a `href` (“**h**ypertext **r**eference”) attribute to specify where the content should link to:

```
<a href="https://ischool.uw.edu">iSchool homepage</a>
```

- In a hyperlink, the *content* of the tag is the displayed text, and the *attribute* specifies the link’s URL. Contrast this to the same link in Markdown:

```
[iSchool homepage](https://ischool.uw.edu)
```

Similarly, an image (``) uses the `src` (**s**ource) attribute to specify what picture it is showing. The `alt` attribute contains alternate text to use if the browser can’t show images—such as with screen readers (for the visible impaired) and search engine indexers.

```

```

- Note that because an `` has no textual content, it is an *empty element* (see below).

There are also a number of global attributes that can be used on any element. For example:

- Every HTML element can include an **id** attribute, which is used to give them a unique identifier so that we can refer to them later (e.g., from CSS or JavaScript). `id` attributes are named like variable names, and must be **unique** on the page.

```
<h1 id="title">My Web Page</h1>
```

The `id` attribute is most commonly used to create “bookmark hyperlinks”, which are hyperlinks to a particular location on a page (i.e., that

cause the page to scroll down). You do this by including the `id` as the **fragment** of the URI to link to (e.g., after the `#` in the URI).

```
<a href="index.html#nav">Link to element on `index.html` with `id="nav"` </a>  
<a href="#footnote">Link to element on current page with `id="footnote"` </a>
```

Note that the title attribute does NOT contain the `#` symbol, but the URI to link to does.

- The `lang` attribute is used to indicate the language in which the element's content is written. Programs reading this file might use that to properly index the content, correctly pronounce it via a screen reader, or even translate it into another language:

```
<p lang="sp">No me gusta</p>
```

Specify the `lang` attribute for the `<html>` element (see below) to define the default language of the page; that way you don't need to mark the language of every element. **Always include this attribute.**

```
<html lang="en">
```

Empty Elements

A few HTML elements don't require a closing tag because they *can't* contain any content. These tags are often used for inserting media into a web page, such as with the `` tag. With an `` tag, you can specify the path to the image file in the `src` attribute, but the image element itself can't contain additional text or other content. Since it can't contain any content, we leave off the end tag entirely:

```

```

Older versions of HTML (and current related languages like XML) required you to include forward slash `/` just before the greater-than symbol. This “end” slash indicated that the element was complete and expected no further content:

```

```

This is no longer required in HTML5, so feel free to omit that forward slash (though some purists, or those working with XML, will still include it).

3.2 Nesting Elements

Web pages are made up of multiple (hundreds! thousands!) of HTML elements. Moreover, HTML elements can be **nested**: that is, the content of an HTML element can contain *other* HTML tags (and thus other HTML elements):



Figure 3.2: An example of element nesting: the `` element is nested in the `<h1>` element’s content.

The semantic meaning indicated by an element applies to *all* its content: thus all the text in the above example is a top-level heading, and the content “(with emphasis)” is emphasized in addition.

Because elements can contain elements which can *themselves* contain elements, an HTML document ends up being structured as a “**tree**” of elements:

In an HTML document, the “root” element of the tree is always an `<html>` element. Inside this we put a `<body>` element to contain the document’s “body” (that is, the shown content):

```
<html lang="en">
  <body>
    <h1>Hello world!</h1>
    <p>This is <em>conteeeeeent</em>!</p>
  </body>
</html>
```

This model of HTML as a tree of “nodes”—along with an API (programming interface) for manipulating them— is known as the **Document Object Model (DOM)**. See Chapter 12 for details.

Caution! HTML elements have to be “closed” correctly, or the semantic meaning may be incorrect! If you forget to close the `<h1>` tag, then *all* of the following content will be considered part of the heading! Remember to close your inner tags *before* you close the outer ones. Validating your HTML can help with this.

Block vs. Inline Elements

All HTML elements fall into one of two categories:

- **Block elements** form a visible “block” on a page—in particular, they will be on a new line from the previous content, and any content after it will also be on a new line. These tend to be structural elements for a page: headings (`<h1>`), paragraphs (`<p>`), lists (``), etc.

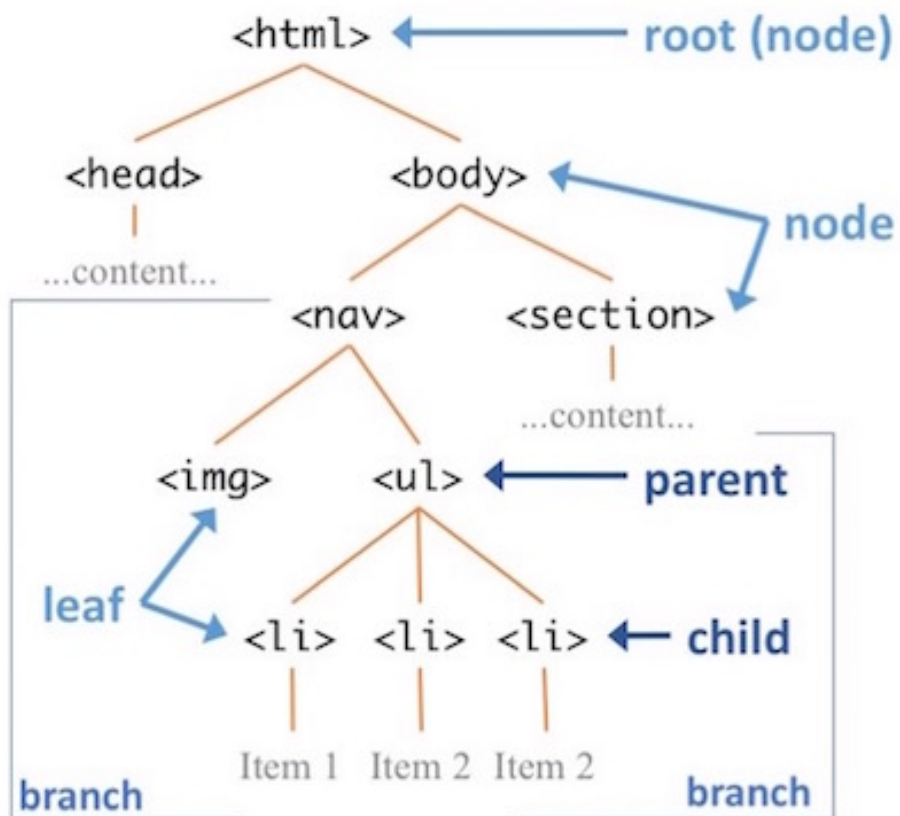


Figure 3.3: An example DOM tree (a tree of HTML elements).

```
<div>Block element</div>
<div>Block element</div>
```



Figure 3.4: Two block elements rendered on a page.

- **Inline elements** are contained “in the line” of content. These will *not* have a line break after them. Inline elements are used to modify the content rather than set it apart, such as giving it emphasis () or declaring that it to be a hyperlink (<a>).

```
<span>Inline element</span>
<span>Other inline element</span>
```




Figure 3.5: Two inline elements rendered on a page.

Inline elements go inside of block elements, and it’s common to put block elements inside of the other block elements (e.g., an inside of a , or a <p> inside of a <div>). However, it is invalid to nest a block element inside of an inline element—the content won’t make sense, and probably won’t look right.

Some elements have further restrictions on nesting. For example, a (unordered list) is *only* allowed to contain elements—anything else is invalid markup.

3.3 Web Page Structure

Now that you understand how to specify HTML elements, you can begin making real web pages! However, there are a few more tags you need to know and include for a valid, modern web page.

Doctype Declaration

All HTML files start with a document type declaration, commonly referred to as the “Doctype.” This tells the rendering program (e.g., the browser) what format and syntax your document is using. Since you’re writing pages with HTML 5, you can declare it as follows:

```
<!DOCTYPE html>
<html lang="en">
...
</html>
```

`<!DOCTYPE>` isn't technically an HTML tag (it's actually XML). While modern browsers will perform a “best guess” as to the Doctype, it is best practice to specify it explicitly. Always include the DOCTYPE at the start of your HTML files!

The `<head>` Section

In addition to the `<body>` element that defines the displayed content, you should also include a `<head>` element that acts as the document “header” (the `<head>` is nested inside the `<html>` at the same level as the `<body>`). The content of the `<head>` element is *not* shown on the web page—instead it provides extra (meta) information *about* the document being rendered.

There are a couple of common elements you should include in the `<head>`:

- A `<title>`, which specifies the “title” of the webpage:

```
<title>My Page Title</title>
```

Browsers will show the page title in the tab at the top of the browser window, and use that as the default bookmark name if you bookmark the page. But the title is *also* used by search indexers and screen readers for the blind, since it often provides a strong signal about what the page's subject. Thus your title should be informative and reflective of the content.

- A `<meta>` tag that specifies the character encoding of the page:

```
<meta charset="UTF-8">
```

The `<meta>` tag itself represents “metadata” (information about the page's data), and uses an attribute and value to specify that information. The most important `<meta>` tag is for the character set, which tells the browser how to convert binary bits from the server into letters. Nearly all editors these days will save files in the UTF-8 character set, which supports the mixing of different scripts (Latin, Cyrillic, Chinese, Arabic, etc) in the same file.

- You can also use the `<meta>` tag to include more information about the author, description, and keywords for your page:

```
<meta name="author" content="your name">
<meta name="description" content="description of your page">
<meta name="keywords" content="list,of,keywords,separate,by,commas">
```

Note that the `name` attribute is used to specify the “variable name” for that piece of metadata, while the `content` attribute is used to specify the “value” of that metadata. `<meta>` elements are *empty elements* and have no content of their own.

Again, these are not visible in the browser window (because they are in the `<head>`!), but will be used by search engines to index your page.

- *At the very least, always include author information for the pages you create!*
- We will discuss additional elements for the `<head>` section throughout the text, such as using `<link>` to include CSS and using `<script>` to include JavaScript.

3.4 Web Page Template

Putting this all together produces the following “template” for making a web page:

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="utf-8">
  <meta name="author" content="your name">
  <meta name="description" content="description of your page">
  <title>My Page Title</title>
</head>
<body>
  ...
  Content goes here!
  ...
</body>
</html>
```

You can use this to start off every web page you ever create from now on!

Resources

- [Getting Starting with HTML](#)
- [General HTML 5 Reference](#)
- [Alphabetical HTML Tag Reference](#)
- [Dive into HTML5 - Free book on HTML](#)
- [W3C HTML Validation Service](#)

Also remember you can view the HTML page source of *any* webpage you visit. Use that to explore how others have developed pages and to learn new tricks and techniques!

Chapter 4

CSS Fundamentals

CSS (Cascading **Style Sheets**) is a declarative language used to alter the appearance or *styling* of a web page. CSS is used to define a set of formatting **rules**, which the browser applies when it renders your page. Thus CSS can tell the browser to use a particular *font* for the page text, a certain *color* for the first paragraph in an article, or a picture for the page's *background*.

Files of CSS rules (called **stylesheets**) thus act kind of like Styles or Themes in PowerPoint, but are way more powerful. You can control nearly every aspect of an element's appearance, including its overall placement on the page.

- To give you some idea of just how much control you have, check out the examples in the CSS Zen Garden. Every one of those examples uses the exact same HTML page, but they all look completely different because each one uses a different CSS stylesheet.

This chapter will explain how to include CSS in your web page and the overall syntax for declaring basic CSS rules. Additional details and options can be found in the following chapter

4.1 Why Two Different Languages?

If you are new to web programming, you might be wondering why there are two different languages: HTML for your page content; and CSS for formatting rules. Why not just include the formatting right in with the content?

There is an old, tried-and-true principle in programming referred to as “**separation of concerns**”. Well-designed software keeps separate things separate, so that it's easy to change one without breaking the other. And one of the most common forms of separation is to keep the **data** (content) in a program separate from the **presentation** (appearance) of that data.

By separating content (the HTML) from its appearance (the CSS), you get a number of benefits:

- The same content can easily be presented in different ways (like in the CSS Zen Garden). In web development, you could allow the user to choose different “themes” for a site, or you could change the formatting for different audiences (e.g., larger text for vision-impaired users, more compact text for mobile users, or different styles for cultures with different aesthetic sensibilities).
- You can have several HTML pages to all share the same CSS stylesheet, allowing you to change the look of an entire web site by only editing one file. This is an application of the Don’t Repeat Yourself (DRY) principle.
- You can also dynamically adjust the look of your page by applying new style rules to elements in response to user interaction (clicking, hovering, scrolling, etc.)
- Users who don’t care about the visual appearance (e.g., blind users with screen readers, automated web indexers) can more quickly and effectively engage with the content without needing to determine what information is “content” and what is just “aesthetics”.

Good programming style in web development thus keeps the **semantics** (HTML) separate from the **appearance** (CSS). Your HTML should simply describe the meaning of the content, not what it looks like!

For example, while browsers might normally show `` text as italic, we can use CSS to instead make emphasized text underlined, highlighted, larger, flashing, or with some other appearance. The `` says nothing about the visual appearance, just that the text is emphatic, and it’s up to the styling to determine how that emphasis should be conveyed *visually*.

4.2 CSS Rules

While it’s possible to write CSS rules directly into HTML, the best practice is to create a separate CSS **stylesheet** file and connect that to your HTML content. These files are named with the `.css` extension, and are typically put in a `css/` folder in web page’s project directory, as with the following folder structure:

```
my-project/  
|-- css/  
    |-- style.css  
|-- index.html
```

(`style.css`, `main.css`, and `index.css` are all common names for the “main” stylesheet).

You connect the stylesheet to your HTML by adding a **<link>** element to your page's **<head>** element:

```
<head>
  <!--... other elements here...-->

  <link rel="stylesheet" href="css/style.css">
</head>
```

The **<link>** element represents a connection to another resource. The tag includes an attribute indicating the **relation** between the resources (e.g., that the linked file is a stylesheet). The **href** attribute should be a *relative path* from the **.html** file to the **.css** resource. Note also that a **<link>** is an empty element so has no closing tag.

- It is also possible to include CSS code directly in your HTML by embedded it in a **<style>** element in the **<head>**, but this is considered bad practice (keep concerns separated!) and should only be used for quick tests.

Overall Syntax

A CSS stylesheet lists **rules** for formatting particular elements in an HTML page. The basic syntax looks like:

```
/* this is pseudocode for a CSS rule */
selector {
  property: value;
  property: value;
}

/* this would be another, second rule */
selector {
  property: value;
}
```

A CSS **rule** starts with a **selector**, which specifies which elements the rule applies to. The selector is followed by a pair of braces **{ }**, inside of which is a set of formatting **properties**. Properties are made up of the property *name* (e.g., **color**), followed by a colon (**:**), followed by a *value* to be assigned to that property (e.g., **purple**). Each name-value pair must end with a semi-colon (**;**).

- If you forget the semi-colon, the browser will likely ignore the property and any subsequent properties—and it does so silently without showing an error in the developer tools!

Like most programming languages, CSS ignores new lines and whitespace. However, most developers will use the styling shown above, with the brace on the same line as the selector and indented properties.

As a concrete example, the below rule applies to any `h1` elements, and makes them appear in the ‘Helvetica’ font as white text on a dark gray background:

```
h1 {  
  font-family: 'Helvetica';  
  color: white;  
  background-color: #333; /*dark gray*/  
}
```

Note that CSS **comments** are written using the same block-comment syntax used in Java (`/* a comment */`), but *cannot* be written using inline-comment syntax (`//a comment`).

When you modify a CSS file, you will need to *reload the page in your browser* to see the changed appearance. If you are using a program such as `live-server`, this reloading should happen automatically!

CSS Properties

There are many, many different CSS formatting properties you can use to style HTML elements. All properties are specified using the `name:value` syntax described above—the key is to determine the name of the property that produces the appearance you want, and then provide a valid value for that property.

Pro-tip: modern editors such as VS Code will provide auto-complete suggestions for valid property names and values. Look carefully at those options to discover more!

Below is a short list of common styling properties you may change with CSS; more complex properties and their usage is described in the following chapter.

- `font-family`: the “font” of the text (e.g., `'Comic Sans'`). Font names containing white space *must* be put in quotes (single or double), and I tend to quote any specific font name as well.

Note that the value for the `font-family` property can also be a *comma-separated list* of fonts, with the browser picking the first item that is available on that computer:

```
/* pick Helvetic Nue if exists, else Helvetica, else Arial, else the default  
   sans-serif font */  
font-family: 'Helvetica Nue', 'Helvetica', 'Arial', sans-serif;
```

- `font-size`: the size of the text (e.g., `12px` to be 12 pixels tall). The value must include units (so `12px`, not `12`). See the next chapter for details on Units & Sizes.
- `font-weight`: boldness (e.g., `bold`, or a numerical value such as `700`).

- **color**: text color (e.g., either a named color like `red` or a hex value like `#4b2e83`. See the next chapter for details on colors. The `background-color` property specifies the background color for the element.
- **border**: a border for the element (see “Box Model” in the next chapter). Note that this is a short-hand property which actually sets multiple related properties at once. The value is thus an *ordered* list of values separated by **spaces**:

```
/* border-width should be 3px, border-style should be dashed, and border-color  
   should be red */  
border: 3px dashed red;
```

Read the documentation for an individual property to determine what options are available!

Note that not all properties or values be effectively or correctly supported by all browsers. Be sure and check the browser compatibility listings!

CSS Selectors

Selectors are used to “select” which HTML elements the css rule should apply to. As with properties, there are many different kinds of selectors (and see the following chapter), but there are three that are most common:

Element Selector

The most basic selector, the **element selector** selects elements by their element (tag) name. For example, the below rule will apply the all `<p>` elements, regardless of where they appear on the page:

```
p {  
  color: purple;  
}
```

You can also use this to apply formatting rules to the entire page by selecting the `<body>` element. Note that for clarity/speed purposes, we generally do *not* apply formatting to the `<html>` element.

```
body {  
  background-color: black;  
  color: white;  
}
```

Class Selector

Sometimes you want a rule to apply to only *some* elements of a particular type. You will most often do this by using a **class selector**. This rule will select elements with a `class` attribute that contains the specified name. For example, if you had HTML:

```
<!-- HTML -->
<p class="highlighted">This text is highlighted!</p>
<p>This text is not highlighted</p>
```

You could color just the correct paragraph by using the class selector:

```
/* CSS */
.highlighted {
    background-color: yellow;
}
```

Class selectors are written with a single dot (`.`) preceding the *name of the class* (not the name of the tag!) The `.` is only used in the CSS rule, not in the HTML `class` attribute.

Class selectors also let us apply a single, consistent styling to multiple different types of elements:

```
<!-- HTML -->
<h1 class="alert-flashing">I am a flashing alert!</h1>
<p class="alert-flashing">So am I!</p>
```

CSS class names should start with a letter, and can contain hyphens, underscores, and numbers. Words are usually written in lowercase and separated by hyphens rather than camelCased.

Note that HTML elements can contain **multiple classes**; each class name is separate by a **space**:

```
<p class="alert flashing">I have TWO classes: "alert" and "flashing"</p>
<p class="alert-flashing">I have ONE class: "alert-flashing"</p>
```

The class selector will select any element that *contains* that class in its list. So the first paragraph in the above example would be selected by either `.alert` **OR** `.flashing`.

You should always strive to give CSS classes **semantic names** that describe the purpose of element, rather than just what it looks like. `highlighted` is a better class name than just `yellow`, because it tells you what you're styling (and will remain sensible even if you change the styling later). Overall, seek to make your class names *informative*, so that your code is easy to understand and modify later.

There are also more formal methodologies for naming classes that you may wish to utilize, the most popular of which is BEM (Block, Element, Modifier).

Class selectors are often commonly used with `<div>` (block) and `` (inline) elements. These HTML elements have *no* semantic meaning on their own, but can be given appearance meaning through their `class` attribute. This allows them to “group” content together for styling:

```
<div class="cow">
  <p>Moo moo moo.</p>
  <p>Moooooooooooooooooooooo.</p>
</div>

<div class="sheep">
  <p>Baa baa <span class="dark">black</span> sheep, have you any wool?</p>
</div>
```

Id Selector

It is also possible to select HTML elements by their `id` attribute by using an **id selector**. Every HTML element can have an `id` attribute, but unlike the `class` attribute the value of the `id` must be unique within the page. That is, no two elements can have the same value for their `id` attributes.

Id selectors start with a `#` sign, followed by the value of the `id`:

```
<div id="sidebar">
  This div contains the sidebar for the page
</div>

/* the one element with id="sidebar" */
#sidebar {
  background-color: lightgray;
}
```

The `id` attribute is more specific (it’s always just one element!) but less flexible than the `class` attribute, and makes it harder to “reuse” your styling across multiple elements or multiple pages. Thus you should *almost always use a class selector instead*, unless you are referring to a single, specific element.

4.3 The Cascade

CSS is called **Cascading** Style Sheets because multiple rules can apply to the same element (in a “cascade” of style!)

CSS rules are *additive*—if multiple rules apply to the same element, the browser will combine all of the style properties when rendering the content:

```
/* CSS */
p { /* applies to all paragraphs */
  font-family: 'Helvetica'
}

.alert { /* applies to all elements with class="alert" */
  font-size: larger;
}

.success { /* applies to all elements with class="success" */
  color: #28a745; /* a pleasant green */
}

<!-- HTML -->
<p class="alert success">
  This paragraph will be in Helvetica font, a larger font-size, and green color,
  because all 3 of the above rules apply to it.
</p>
```

CSS styling apply to *all* of the content in an element. And since that content can contain other elements that may have their own style rules, rules may also in effect be *inherited*:

```
<div class="content"> <!-- has own styling -->
  <div class="sub-sec"> <!-- has own styling + .content styling -->
    <ol class="demo-list"> <!-- own styling (ol AND .demo-list rules) + .sub-sec + .
      <!-- li styling + .demo-list + .sub-sec + .content -->
        <li>Item 1</li>
        <li>Item 2</li>
        <li>Item 3</li>
      </ol>
    </div>
  </div>
```

We call these inherited properties, because the child elements inherit the setting from their ancestor elements. This is a powerful mechanism that allows you to specify properties only once for a given branch of the DOM element tree. In general, try to set these properties on the highest-level element you can, and let the child elements inherit the setting from their ancestor.

Rule Specificity

Important! Rules are applied in the order they are defined in the CSS file. If you link multiple CSS files to the same page, the files are processed in order they are linked in the HTML. The browser selects elements that match the rule and applies the rule's property. If a later rule selects the same element and applies a different value to that property, the previous value is *overridden*. So in general, all things being equal, **the last rule on the page wins**.

```
/* two rules, both alike in specificity */
p { color: red; }
p { color: blue; }
```

```
<p>This text will be blue, because that rule comes last!</p>
```

However, there are some exceptions when CSS treats rules as *not* equal and favors earlier rules over later ones. This is called Selector Specificity. In general, more specific selectors (`#id`) take precedence over less specific ones (`.class`, which is more specific than `tag`). If you notice that one of your style rules is not being applied, despite your syntax being correct, check your browser's developer tools to see if your rule is being overridden by a more specific rule in an earlier stylesheet. Then adjust your selector so that it has the same or greater specificity:

```
/* css */
.alert { color: red; }
div { color: blue; }
```

```
<!-- html -->
<div class="alert">This text will be red, even though the `div` selector is last,
because the `.alert` selector has higher specificity so is not overridden.</div>
```

Precedence rules are **not** a reason to prefer `#id` selectors over `.class` selectors! Instead, you can utilize the more complex selectors described in Chapter 6 to be able to create reusable rules and avoid duplicating property declarations.

Resources

- Getting started with CSS (MDN)
- CSS Tutorial (w3schools)
- CSS Reference (MDN) a complete alphabetical reference for all CSS concepts.
- CSS Selectors Reference a handy table of CSS selectors.
- CSS Properties Reference a table of CSS properties, organized by category.
- CSS-Tricks a blog about tips for using CSS in all kinds of ways. Contains many different useful guides and explanations.

- W3C CSS Validation Service

Chapter 5

Standards and Accessibility

This chapter details considerations and techniques to ensure that your web pages and applications can be used by *everyone*, regardless of their choice of web browser or their own physical abilities.

5.1 Web Standards

As discussed in Chapter 2, the HTML and CSS files you author are delivered to clients upon request. The code within these files is *interpreted* by the **web browser** in order to create the visual presentation that the user can see and interact with.

A web browser is thus any piece of software capable of rendering these `.html` and `.css` files (and sending HTTP requests to fetch them in the first place). And there are many different web browsers in the world:

- Note that while there may be some clear “winners” in terms of browser popularity, do not dismiss less popular browsers. For example, if 0.34% of users use Internet Explorer 8 (July 2017), that’s roughly 1,300,000 people worldwide.

These web browsers are all created by different developers, working for different (often rival!) organizations. How is it that they are all able to read and interpret the same code, and produce the same rendered output? *Standards*.

Web Standards are agreed-upon specifications for how web page source code should be rendered by the browser. Web standards detail both the language syntax (e.g., how to write HTML tags) and the language semantics (e.g., which HTML tags to use), so that it can be understood by any browser that follows (agrees to) that standard. And since as a developer you want your pages to




Figure 5.1: *Desktop* browser market share July 2017, from netmarketshare.com. See also caniuse.com’s usage table for information on mobile and older browsers.

render the same across all browsers, web standards give the requirements for how you should write your code so that your pages render correctly.

Modern web standards are created and maintained by a huge group of stakeholders known as the World Wide Web Consortium (W3C), which includes major browser developers such as Google and Mozilla. However, this group has no enforcement powers: and so browsers often deviate from the published standards! A browser may ignore a standard to “help out” developers (e.g., making a “best guess” to render malformed HTML content), or to introduce new features (e.g., a new CSS property that produces some special effect). Some browsers are better at conforming to the accepted standards than others. Internet Explorer—IE 6 in particular—is notorious for not meeting standards and requiring extra effort from developers to make pages work. This is part of what IE has such a bad reputation and gets so much scorn from developers (though Microsoft Edge, the browser in Windows 10, is awesome about being standards compliant).

Getting so many people to agree on a standard of communication takes time; thus web standards change relatively slowly: the HTML 4 standard was adopted in 1997, but the **HTML 5** standard (that this course teaches) wasn’t finalized until 2014. The **CSS 3** standard is broken up into more than 50 different *modules* that are developed and introduced independently, and so is continuously being adopted piece-wise.

When introducing new or experimental CSS properties, browsers historically utilized vendor prefixes in naming the properties. As standards are being decided upon, each browser may treat that property in a slightly different way, thus forcing developers who want to use the “latest and greatest” to provide a different definition or value for each vendor. Prefixes are a naming convention that will cause the property to only be understood and applied by a particular browser; e.g., `-webkit-hyphens` would use the Webkit version of `hyphens` property (for controlling word breaks), while `-ms-hyphens` would use the IE version. This

practice is currently discouraged and being phased out, though prefixes may be required when supporting older browsers. Tools such as Autoprefixer can help automatically manage prefixes.

In general, as long as you program webpages that conform to web standards, they will work on all “*modern browsers*”—though there may still be a few cross-browser compatibility concerns to notice. You should thus test and **validate** your code against the standards. Luckily, there the W3C provides online tools that can help validate code:

- W3C HTML Validation Service
- W3C CSS Validation Service
- W3C Developer Tools for a complete list of validators

To use these services, simply enter your web pages URL (or copy and paste the contents of your `.html` or `.css` files), and then run the validation. You will definitely want to fix any *errors* you get. *Warnings* should be considered; however, it is possible to get false positives. Be sure and read the warning carefully and consider whether or not it is actually a “bug” in your code!

Supporting *really* old browsers that are not standards compliant is left as an exercise to the reader.

Develop for other people’s browsers, not your own! Just because it “works” for you, doesn’t mean it works for anyone else. Test your code against standards and automated systems; don’t just look at the rendered result in a single browser!

5.2 Why Accessibility

Consider the following hypothetical webpage user:

Tracy is a 19-year-old college student and was born blind. In high school she did well as she could rely on audio tapes and books and the support of tutors, so she never bothered to really learn Braille. She is interested in English literature and is very fond of short stories; her dream is to become an audiobook author. Tracy uses the Internet to share her writing and to connect with other writers through social networks. She owns a laptops and uses a screen reader called JAWS: a computer program which reads her screen out loud to her in an artificial voice. (Adapted from here)

One of the most commonly overlooked limitations on a website’s usability is whether or not it can be used by people with some form of disability. There are many different forms of disabilities or impairments that may affect whether or not a person can access a web page, including:

- *Vision Impairments*: About 2% of the population is blind, so use alternate mediums for reading web pages. Farsightedness and other vision problems are also very common (particularly among older adults), requiring larger and clearer text. Additionally, about 4.5% the population is color-blind.
- *Motor Impairments*: Arthritis occurs in about 1% of the population, and can restrict people's ease at using a mouse, keyboard or touch screen. Other impairments such as tremors, motor neuron conditions, and paralysis may further impact people's access.
- *Cognitive Impairments*: Autism, dyslexia, and language barriers may cause people to be excluded from utilizing your website.

If you fail to make your website accessible, you are locking out 2% or more of users, reducing the availability and use of your website. Indeed, companies are finally seeing this population as an important but excluded market; for example, Facebook has an entire team devoted to accessibility and supporting users with disabilities. “*Accessibility Engineers*” have good job prospects.

Supporting users with disabilities is not just the morally correct thing to do, it's also *the law*. US Courts have ruled that public websites are subject to Title III of the Americans with Disabilities Act (ADA), meaning that is a possible and not uncommon occurrence for large organizations to be sued for discrimination if their websites are not accessible. So far, “accessibility” has legally meant complying with the W3C's Web Content Accessibility Guidelines (WCAG) (see below), a standard that is not overly arduous to follow if you consider accessibility from the get-go.

Finally, designing your website (or any system) with accessibility in mind will not just make it more usable for those with disabilities—it will make it more usable for *everyone*. This is the principle design **Universal Design** (a.k.a. *universal usability*): designing for *accessibility*—being usable by all people no matter their ability (physical or otherwise)—benefits not just those with some form of limitation or disability, but **everyone**. The classic real-world example of universal design are curb cuts: the “slopes” built into curbs to accommodate those in wheelchairs. However, these cuts end up making life better for *everyone*: people with rollerbags, strollers, temporary injuries, small children learning to ride a bicycle, etc.

Universal design applies to websites as well:

- If you support people who can't see, then you may also support people who can't see *right now* (e.g., because of a bad glare on their screen).
- If you support people with motor impairments, then you may also support people trying to use your website without a mouse (e.g., from a laptop while on a bumpy bus).
- If you support people with cognitive impairments, then you may also support people who are temporarily impaired (e.g., inebriated or lacking sleep).

If you make sure that your web page is well-structured and navigable by those

with screen readers, it will ensure that it is navigable by *other* machines, such as search engine indexers. Or for unusual or future browsers (a virtual reality browser perhaps).

Thus supporting accessibility in client-side web development is important both for helping a population that is often overlooked (a form of social justice), as well as for supporting new technologies and systems. This fact is increasingly being acknowledged by companies as key to usability, and thus it is important that you apply it to your own design work.

In addition to individual capabilities, people are also reliant on a large amount of existing *infrastructure* to ensure that they have an internet connection and their requests can reach your web server. The lack of such access is often tied to economic or social inequalities, forming what is called the *digital divide*. Considering the availability of network access and other infrastructural needs is vitally important when developing information technologies, but is difficult to manage through client-side development (though see Responsive Design for some things you can do).

5.3 Supporting Accessibility

In this course, we will primarily aim to support accessibility for users with visual impairments such as those using screen readers. A **screen reader** is a piece of software that is able to synthesize and “read” content on a computer’s screen out loud through the speakers, so that users are able to navigate and control the computer without needing to see the screen. Screen readers are often combined with *keyboard controls*, so that users use just the keyboard to control the computer and not the mouse (almost like a command-line interface!).

There are a number of different screen reader software packages available:

- Macs have had VoiceOver built into the operating system since 2005, though it has been refined with each new OS version.
- On Windows, the most popular screen readers are JAWS and NDVA. Windows also has a built-in screen reader called Microsoft Narrator, which with Windows 10 is beginning to reach parity with the 3rd-party offerings.

You should try out this software! Follow the above links to learn how to turn on the screen reader for your computer, and then try using it to browse the internet *without looking at the screen*. This will give you a feel for what it is like using a computer while blind.

Screen readers are just software that interpret the HTML of a website in order to allow the user to hear and navigate the content—they are basically *non-visual web browsers*. As such, supporting screen readers just means implementing your web site so it works with this browser. In particular, this means making

sure your page conforms to the Web Accessibility Content Guidelines (WCAG). This is a list of principles and techniques that you should utilize when authoring HTML document in order to make sure that they are accessible. The guidelines are driven by 4 main principles:

1. **Perceivable:** Information and user interface components must be presentable to users in ways they can perceive.
2. **Operable:** User interface components and navigation must be operable.
3. **Understandable:** Information and the operation of user interface must be understandable.
4. **Robust:** Content must be robust enough that it can be interpreted reliably by a wide variety of user agents, including assistive technologies.

More concretely, accessible web pages are those that can be **navigated** by screen readers (so people can easily get to the information they care about), and have content can be **perceived** by screen readers (so is not just presented visually).

Don't worry; there are simple, specific implementation steps you can to follow these principles, the most common of which are described below.

Semantic HTML

A major step in making sure your web pages are accessible is to make sure that your use of HTML elements is *semantic*. HTML elements should be used to describe the meaning or form of their content, *not* to give that content a visual appearance!

For example, an `<h1>` element is used to indicate a top-level heading, such as the title of the page. But by default, browsers will give `<h1>` elements a different visual appearance than unmarked content (usually by making it larger and bold). It is possible to achieve a similar visual effect just using CSS:

```
<!-- html -->
<h1>This is a real heading!</h1>
<div class="fake-header">
    This just LOOKS like a heading.
</div>
```

```
/* css */
.fake-header {
    font-size: 2em;
    font-weight: bold;
}
```

Your HTML should **always** be semantic! A screen reader will (mostly) ignore the CSS, so if something is supposed to be a top-level heading, you need to tag it as such for the screen reader to understand that. Similarly, you should **only** use elements when they are semantically appropriate—don't use an `<h3>`

This is a real heading!

This just LOOKS like a heading.

Figure 5.2: A real heading and a fake heading

element just because you want text to look bigger if it isn't actually a third-level heading (use CSS to do the styling instead).

HTML is just for semantics. CSS is for appearance.

This is the same reason that you should never use the `<i>` tag to make elements *italic*. The `<i>` tag has no valuable semantic meaning (it just means “italic”, or “different”); instead you want to use a more meaningful semantic tag. Usually this means ``, since italic text is usually being *emphasized*.

In fact, HTML 5 includes a large number of semantic elements that can be used to indicate “formatted” text that has a particular meaning, including ones such as: `<abbr>`, `<address>`, `<cite>`, ``, `<dfn>`, `<mark>`, and `<time>`. Be sure and check the documentation for each of these elements to determine when they are supposed to be used (as well as what additional attributes they may support).

Never use the `<table>` tag to structure documents and position content. The `<table>` tag should **only** be used for content that is semantically a table (e.g., a data table). If you want to lay out content in a grid, use a CSS system such as Flexbox (see Chapter 7).

ARIA

The WCAG is developed by the Web Accessibility Initiative (WAI), who *also* have authored an additional **web standard** for supporting accessibility. This standard is called the Accessible Rich Internet Applications Suite, or **ARIA**. ARIA specifies an *extension* to HTML, defining additional HTML *attributes* that can be included in elements in order to provide additional support for screen readers. For example, ARIA provides additional support to help screen readers navigate through a page (for people who can't just visually scroll), as well as providing a mechanism by which “rich” interactive web apps (such as those you will create with JavaScript) can communicate their behavior to screen readers.

In short, where the HTML standard falls short in supporting accessibility, ARIA steps in.

For example, the ARIA specification defines a **role** attribute that allows non-semantic elements (such as `<div>` elements) to specify their semantic purpose.

For example:

```
<div class="btn" role="button">Submit</div>
```

would define a unsemantic `div` element that is *styled* (via a developer-defined CSS class) to look like a button. But in order to tell the screen reader that this element can be clicked like a button, it is given the `role` of `"button"`. There are a large number of different roles (following a potentially complex classification scheme), and include roles for interactive widgets such as a `popup dialog`, a `progressbar`, or a `menu`. These are important when making complex interactions, though for most simpler pages you can stick with the standard semantic elements (e.g., a `<button>` instead of a custom `<div>`).

The other important type of ARIA `role` are landmark roles. These are roles that are used to indicate specific locations in the document, to allow the screen reader to easily “jump” to different sections of the content (thus making the page *navigable*). Landmark roles include:

- **banner**: the “banner” at the top of the page—that is, content related to the website rather than to the specific page itself. Alternatively, use the `<header>` element described below (in which case you don’t need to include the `role="banner"` attribute, as it is implied).
- **navigation**: a section of links for moving around the page, such as a navigation bar. This allows screen readers to quickly jump to the “quick links” and move around the page as quickly as possible. Note though that listening to all of the links read can be time consuming, so allowing the user to jump to the rest of the content is also important. Alternatively, use the `<nav>` element described below.
- **main**: the start of the “main” content (usually after the banner). This is useful for allowing the screen reader to jump to the meat of the page, past the banner and navigation sections. Alternatively, use the `<main>` element described below.
- **contentinfo**: a section that contains information about the webpage, such as author contact info and copyright information. This is usually found at the bottom of the page. Alternatively, use the `<footer>` element described below.

Considering your content in terms of these roles is a good design trick for figuring out how to organize your page.

Page Structure (Navigable)

Because screen readers cannot take in a web page’s content “at a glance”, accessible pages need to be explicitly structured so that their content can be easily navigated—e.g., so that the user can quickly “scroll” to a particular blog post.

The most important way to provide this structure is by the considered use of **heading** elements (`<h1>`, `<h2>`, etc). Screen readers automatically generate a “table of contents” based on these headings, allowing users to easily move through large amounts of content. In order to be sure that the headings are useful, they need to be **meaningful** (actually marking section headings) and **hierarchical** (they don’t skip levels: every `<h3>` has an `<h2>` above it). The former is just good HTML usage that you will be doing anyway; the later may take some consideration.

The `<h#>` heading elements are part of the original HTML specification, and so will be supported by *all* screen reader systems. However, HTML 5 introduced additional elements that can be used to help organize web page content in order to make its structure more explicit. These are often referred to as semantic elements or “sectioning elements”. These are all **block-level** elements that produce no visual effects on their own, but provide semantic structuring to web content. You can think of them as specialized `<div>` elements.

One metaphor is that nested HTML elements into a `<div>` is like putting that content into a envelope to “group” it together (e.g., for styling purposes). In that case, a semantic element such as `<header>` is just an envelope with a unique color that makes it easier to find in the filing cabinet.

- `<header>` represents the “header” or introduction part of the page, such as the title or banner image. It may also include common page elements such as navigation or search bars. This corresponds to the `role="banner"` landmark role; you do not need to include that role if you use this element.

Note that a `<header>` is different from a heading (`<h1>`) which is different from the `<head>`! An *heading* is element (e.g., `<h1>`) that includes a title or subtitle. The *header* is a “grouping” element that can contain multiple elements, and usually has banner/logo information. The `<head>` is an element that is NOT part of the `<body>` (so is not shown in the web page), and contains *metadata* about that page.

- `<nav>` contains navigation links, usually for navigating around the site (think like a navigation bar). Not all links need to be in a `<nav>`; this is for “sections” of the webpage that are purely navigational. This element corresponds to the `role="navigation"` landmark role.
- `<main>` represents the “main” content of the document. This usually comes *after* the `<header>` (but not inside—in fact, `<main>` *cannot* be a descendant of `<header>`). Note that a web page can only have a single `<main>` element. This element corresponds to the `role="main"` landmark role.
- `<footer>` represents the “footer” of the page, usually containing information about the page. This element corresponds to the `role="contentinfo"` landmark role.
- `<section>` represents a standalone section of content (e.g., that might

have a subheading such as `<h3>`). A `<section>` can also contain its own `<header>`, `<footer>`, and `<nav>` elements relevant to that section.

Similarly, an `<article>` element also represents standalone content, but content that might be published independently (such as a news article or a blog post). Note that a `<section>` may group together multiple `<article>` elements (such as a blog roll), and an `<article>` might potentially contain more than one `<section>`. Think about how a newspaper is structured (with a “Sports Section” that contains articles, which may themselves have different sections).

Overall, utilizing these semantic sectioning elements will help organize your content for screen readers (so visually impaired users can easily navigate the page), as well as making your content more clearly structured.

Visual Information (Perceivable)

Webpages often contain a significant amount of *primarily visual* information. In addition to obvious media such as images or video, interactive elements may be labeled visually (e.g., a search button that shows a magnifying glass). In order to make web pages accessible, screen readers need a way to **perceive** and correctly interpret such content.

The most common form of visual information are images (created with the `` element). In order to make images accessible, you should always include an `alt` attribute that gives **alternate** text for when the images cannot be displayed (e.g., on screen readers, but also if the image fails to load):

```

```

This will be read by screen readers as “a cute baby, image”. Note that the “alt-text” should not include introductory text such as “*a picture of*”, as screen readers will already report that something is an image!

Every image should include an `alt` attribute!

For more complex images (such as charts or infographics), you can instead provide a *link* to a longer description by using the `longdesc` attribute. This attribute takes a value that is a URI (relative or absolute; what you would put in the `href` attribute of a hyperlink) referring to where the description can be found. Screen readers will prompt the user with the option to then navigate to this long description.

```
 <!-- link to other page with text description
```

Of course, including descriptive text is good for *any* user, not just the visually impaired! If you wish to add a caption to an image, you can do so accessibly by

placing the image inside a `<figure>` element, and then using a `<figcaption>` element to semantically designate the caption.

```
<figure>
  
  <figcaption>
    A caption for the above figure. It provides the same information,
    but in a text format.
  </figcaption>
</figure>
```

(The `<figure>` element is another good example of how multiple HTML elements may be nested and combined to produce complex but accessible page structure!)

On the other hand, some images (or other elements) are purely decorative: company logos, icons that accompany text descriptions on buttons, etc. You can cause a screen reader to “skip” these elements by using an ARIA attribute `aria-hidden`. The element will still appear on the page, it just will be ignored by screen readers.

```
<!-- a search button with an icon -->
<!-- the icon will not be read, but the button text will be -->
<button>Search</button>
```

It is also possible to use ARIA attribute to provide an equivalent to an `alt` tag for elements other than ``, such as a `<div>` that may be styled in a purely visual way (perhaps with a background image). The **aria-label** acts like the `alt` attribute for any element, specifying what text should be read *in place of the normal content*.

```
<div class="green-rect" aria-label="a giant green rectangle"></div>
```

For longer descriptions, use the **aria-describedby** attribute to include a reference to the `id` of a different element *on the same page* that contains the textual description. `aria-describedby` takes as a value a fragment reference (similar to a bookmark hyperlink) to the element containing the description.

```
<div class="green-rect" aria-describedby="#rectDetail"></div>
<p id="rectDetail">The above rectangle is giant and green.</p>
```

As with bookmark hyperlinks, notice that the `aria-describedby` value starts with a `#`, but the target element’s `id` does not.

Additionally, note that while `aria-label` attributes *replaces* its elements content, the `aria-describedby` attribute will be read *after* the element’s content (see also here).

ARIA provides a number of other attributes that can be used to control screen readers, but these are the most common and useful.

Overall, the way to make purely visual information accessible is to *not have purely visual information*. Always include textual descriptions and captions for images, prefer text labels for buttons, etc.

In conclusion, creating accessible content basically means writing proper and semantic HTML (what you would do doing anyway), with the small extra step of making sure that visual content is also labeled and perceivable. By meeting these standards, you will ensure that your website is usable by everyone.

Resources

- Caniuse.com details about browser support for emerging web features
- W3C Validation Tools
- Website Accessibility and the Law summative blog post, May 2017
- **Teach Access Tutorial** for creating accessible web pages. Provides lots of examples and details. Be sure and check out the code checklist and the design checklist.
- UW Accessibility Checklist a thorough and detailed list of considerations for developing accessible systems put together by UW. See also their complete list of tools and resources.
- ARIA Resources (MDN)
- Using ARIA (w3c) official guidance for using the ARIA specification
- WAVE Accessibility Evaluation Service

Chapter 6

More CSS

Chapter 4 explained the basic syntax and usage of CSS, enough to let you create and style your own web pages. This chapter provides more details about additional selectors and properties to use when defining CSS rules; the following chapter discusses particular properties that can be used to further style the layout of your page's content.

6.1 Compound Selectors

As described in the previous chapter, the core selectors used in CSS are the **element selector**, **class selector**, and **id selector**. However, CSS does offer ways to combine these selectors in order to specify rules only for particular elements or groups of elements.

Group Selector

The **group selector** allows you to have a single rule apply to elements matched by lots of different selectors. To do this, separate each selector with a comma (,); the properties defined in the rule will then apply to any element that is matched by *any* of the selectors. For example, if you want to have a single style for all headings, you might use:

```
/* applies to h1, h2, AND h3 tags */
h1, h2, h3 {
    font-family: Helvetica;
    color: #4b2e83; /* UW purple */
}
```

The comma-separated selectors can be **any** kind of selector, including `.class` or `#id` selectors (or any of the compound selectors described below):

```
/* can also include class or id selectors */
/* this rule applies to h2 elements, "menu" classed elements, and the
   #sidebar element */
h2, .menu, #sidebar {
    background-color: gray;
}
```

Note that since later rules override earlier ones, you can use a group select to apply a property to multiple different elements, but then include additional rules to add variations. For example, you can have one rule that applies “general” styling to a large class of elements, with further rules then customizing particular elements.

```
/* all headings are Helvetica, bold, and purple */
h1, h2, h3 {
    font-family: Helvetica;
    font-weight: bold;
    color: #4b2e83; /* UW purple */
}

/* h2 elements are not bolded, but italic */
h2 {
    font-weight: normal; /* not bold, overrides previous rule */
    font-style: italic;
}
```

Combined Selectors

It is also possible to combine element, class, and id selectors together to be more specific about where a rule applies. You do this by simply putting the class or id selector *immediately after* the previous selector, without a comma or space or anything between them:

```
/* Selects only p elements that have class="alert"
   Other p elements and "alert" classed elements not affected */
p.alert {
    color: red;
}

/* Selects only h1 elements that have id="title" */
/* Note that this is redundant, since only one element can have the id! */
h1#title {
    color: purple;
}
```



```

/* Selects elements that have class "alert" AND class "success" */
.alert.success {
  color: green;
  font-size: larger;
}

/* And can combine with group selector */
/* applies to <p class="highlighted"> and <li class="selected"> */
p.highlighted, li.marked {
  background-color: yellow;
}

```

This specificity can allow you to reuse class names (e.g., for shared semantics and readability purposes) but have them work differently for different elements. So a “highlighted” paragraph `p.highlighted` might look different than a “highlighted” heading `h1.highlighted`.

Note that putting a space between the selectors parts instead specifies a **descendant selector**, which has a totally separate meaning. Every character matters!

Descendant Selector

So far, all selectors mentioned will apply to matching elements regardless of where they are in the HTML element tree. But sometimes you want to be more specific and style only a set of elements that exist within a particular parent or ancestor element, and not all the other matching elements elsewhere in the page. You can do this form of targeted selecting using a **descendant selector**. This is written by putting a blank space () between selectors. Elements are only selected if they have *parents that match the selectors that precede them*:

```

<header>
  <h1>Welcome to the page</h1>
  <p>I am a special paragraph</p>
</header>
<section>
  <p>some other paragraph</p>
</section>

```

```

/*
  Selects p elements that exist within header elements
  Other p elements will not be affected
*/
header p {
  /* ... */
}

```

```
}
```

You can have as many “levels” of a descendant selector as you want, and each level can be made up of any kind of selector. However, it is best to not have more than 2 or 3 levels. If you need to be more specific than that, then perhaps defining a new `.class` is in order.

```
/* selects elements with class="logo"
   contained within <p> elements
   contained within <header> elements */
header p .logo {
    /* ... */
}
```

Note that descendant selectors will select matching descendant elements *anywhere* lower in the tree branch, not just direct children, so the `.logo` elements here could be nested several layers below the `<p>` element (perhaps inside a ``). This is usually a good idea because you may introduce new nesting layers to your page as you go along, and don’t want to modify the CSS. But if you really want to select only *direct* children, you can use a variant known as a **child selector**, which uses a `>` symbol to indicate direct descendants only:

```
<body>
  <p>Body content</p>
  <section>
    <p>Section content</p>
  </section>
</body>

/* Selects page content (immediately within body),
   not section content (immediately within section) */
body > p {
    color: blue;
}
```

Pseudo-classes

The last kind of selector you will commonly use in web development is the application of what are called **pseudo-classes**. These select elements based on what **state** the element is in: for example, whether a link has been visited, or whether the mouse is hovering over some content. You can almost think of these as pre-defined classes built into the browser, that are added and removed as you interact with the page.

Pseudo-classes are written by placing a colon (`:`) and the name of the pseudo-class immediately after a basic selector like an element selector. You’ll see this most commonly with styling hyperlinks:

```
/* style for unvisited links */
a:link { /*...*/ }

/* style for visited links */
a:visited { /*...*/ }

/* style for links the user is hovering over with the mouse */
a:hover { /*...*/ }

/* style for links that have keyboard focus */
a:focus { /*...*/ }

/* style for links as they are being 'activated' (clicked) */
a:active { /*...*/ }
```

Remember to always set both `hover` and `focus`, to support accessibility for people who cannot use a mouse. Additionally, `a:hover` *must* come after `a:link` and `a:visited`, and `a:active` must come after `a:hover` for these states to work correctly.

Note that there are many additional pseudo-classes, including ones that consider specific element attributes (e.g., if a checkbox is `:checked`) or where an element is located within its parent (e.g., if it is the `:first` or `:last-child`, which can be useful for styling lists).

Attribute Selectors

Finally, it is also possible to select element that have a particular attribute by using an **attribute selector**. Attribute selectors are written by placing brackets `[]` after a basic selector; inside the brackets you list the attribute and value you want to select for using `attribute=value` syntax:

```
/* select all p elements whose "lang=sp" */
p[lang="sp"] {
    color: red;
    background-color: orange;
}
```

It is also possible to select attributes that only “partially” match a particular value; see the documentation for details.

Note that it is most common to use this selector when styling form inputs; for example, to make checked boxes appear different than unchecked boxes:

```
/* select <input type="checked"> that have the "checked" state */
input[type=checkbox]:checked {
```

```
color: green;
}
```

6.2 Property Values

This section of the guide provides further details about the possible *values* that may be assigned to properties in CSS rules. These specifics are often relevant for multiple different properties.

Units & Sizes

Many CSS properties affect the **size** of things on the screen, whether this is the height of the text (the **font-size**) or the width of a box (the **width**; see the next chapter). In CSS, you can use a variety of different **units** to specify sizes.

CSS uses the following **absolute units**, which are the same no matter where they are used on the page (though they are dependent on the OS and display).

Unit	Meaning
px	pixels ($\frac{1}{96}$ of an inch, even on high-dpi “retina” displays)
in	inches (OS and display dependent, but maps to physical pixels in some way)
cm, mm	centimeters or millimeters, respectively
pt	points (defined as $\frac{1}{72}$ of an inch)

Although technically based on in as a standard, it is considered best practice to always use px for values with absolute units.

CSS also uses the following **relative units**, which will produce sizes based on (relative to) the size of other elements:

Unit	Meaning
em	Relative to the current element’s font-size. Although originally a typographic measurement, this unit will not change based on font-family .
%	Relative to the parent element’s font-size <i>or</i> dimension. For font-size, use em instead (e.g., 1.5em is 150% the parent font-size).
rem	Relative to the root element’s font-size (i.e., the font-size of the root html or body element). This will often be more consistent than em.

Unit	Meaning
vw, vh	Relative to the viewport (e.g., the browser window). Represents 1% of the viewport width and height, respectively. This unit is not supported by older browsers.

Note that most browsers have a default font size of **16px**, so **1em** and **1rem** will both be initially equivalent to **16px**.

In general, you should specify font sizes using *relative units* (e.g., **em**)—this will support accessibility, as vision-impaired users will be able to increase the default font-size of the browser and all your text will adjust appropriately. Absolute units are best for things that do not scale across devices (e.g., image sizes, or the maximum width of content). However, using relative sizes will allow those components to scale with the rest of the page.

- Font-sizes should always be relative; layout dimensions may be absolute (but relative units are best).

Colors

Colors of CSS properties can be specified in a few different ways.

You can use one of a list of 140 predefined color names:

```
p {  
  color: mediumpurple;  
}
```

While this does not offer a lot of flexibility, they can act as useful placeholders and starting points for design. The list of CSS color names also has a fascinating history.

Alternatively, you can specify a color as a “red-green-blue” (RGB) value. This is a way of representing *additive color*, or the color that results when the specified amount of red, green, and blue light are aimed at a white background. RGB values are the most common way of specifying color in computer programs.

```
p {  
  color: rgb(147, 112, 219); /* medium purple */  
}
```

This value option is actually a *function* that takes a couple of parameters representing the amount of red, green, and blue respectively. Each parameter ranges from 0 (none of that color) to 255 (that color at full). Thus **rgb(255,0,0)** is pure bright red, **rgb(255,0,255)** is full red and blue but no green (creating magenta), **rgb(0,0,0)** is black and **rgb(255,255,255)** is white.

Note that if you want to make the color somewhat transparent, you can also specify an alpha value using the `rgba()` function. This function takes a 4th parameter, which is a decimal value from 0 (fully transparent) to 1.0 (fully opaque):

```
p {  
  background-color: rgba(0,0,0,0.5); /* semi-transparent black */  
}
```

CSS also supports `hsl()` and `hsla()` functions for specifying color in terms of a hue, saturation, lightness color model.

Finally, and most commonly, you can specify the RGB value as a hexadecimal (base-16) number.

```
p {  
  color: #9370db; /* medium purple */  
}
```

In this format, the color starts with a #, the first two characters represent the red (ranging from 00 to FF, which is hex for 255), the second two characters represent the green, and the third two the blue:



Figure 6.1: How to reading a hex value, from Smashing Magazine.

This is a more compact and efficient way to describe the RGB of a color, and is how most digital artists convey color. See this article for more details about encoding colors.

Fonts and Icons

As mentioned in the previous chapter, you specify the typographical “font” for text using the **font-family** property. This property takes a *comma-separated list* of fonts to use; the browser will render the text using the first font in the list that is available (you can think of the rest as “back-ups”).

```
p {
  font-family: 'Helvetica Nue', 'Helvetica', 'Arial', sans-serif;
}
```

The last font in the list should always be a generic family name. These are a list of “categories” that the browser can draw upon even if the computer doesn’t have any common fonts available. In practice, the most common generic families used are `serif` (fonts with serifs, e.g., “Times”), `sans-serif` (fonts *without* serifs, e.g., “Arial”), and `monospace` (fonts with equal width characters, e.g., “Courier”).




Figure 6.2: Serif vs. Sans-Serif fonts. From 99designs.ca.

It is also possible to include specific fonts in your web page, which will be delivered to the browser by the web server in case the client doesn’t have the previously available. You do this manually by using the `@font-face` rule, and specifying the url for the font file (usually in `.woff2` format).

However, it is usually easier to instead include a stylesheet that *already has this rule in place*. For example, the Google Fonts collection provides more than 800 different freely available fonts that you can include directly in your web page:

```
<head>
  <!-- ... -->

  <!-- load stylesheet with font first so it is available -->
  <link href="https://fonts.googleapis.com/css?family=Encode+Sans" rel="stylesheet">

  <!-- load stylesheet next -->
  <link href="css/style.css" rel="stylesheet">
</head>
```

```
body {
    font-family: 'Encode Sans', sans-serif; /* can now use Encode Sans */
}
```

Notice that the `<link>` reference can be to an external file on a different domain! This is common practice when utilizing fonts and CSS frameworks.

Important Note that when using Google Fonts, you’ll need to specify if you also want variations such as **bold** or *italic* typesets. For example, the Encode Sans font is available in font weights (what you would set with `font-weight`) from 100 to 900, but you need to specify these in the linked resource:

```
<!-- includes normal (400) and bold (700) weights -->
<link href="https://fonts.googleapis.com/css?family=Encode+Sans:400,700" rel="stylesheet">
```

If you don’t include the different set of glyphs for the bolded font, then setting the text in that font to bold won’t have any effect (because the browser doesn’t now how to show text in “Bold Encode Sans”)!

You can also use external font styles to easily add **icons** (symbols) to your web page by using an *dingbat (icon) font*. These are similar in concept to the infamous Wingdings font: they are fonts where instead of the letter “A” looking like two bent lines with a bar between, it instead looks like a heart or a face or some other symbol. By including such a font, you get access to a large number of symbols that can easily be styled (read: colored and sized) to suit your needs.

Emoji are defined using a different set of Unicode values, and are browser and operating-system instead of being available through a font.

The most popular icon set is called Font Awesome. You can include the Font Awesome font by linking to it as any other stylesheet:

```
<link href="https://maxcdn.bootstrapcdn.com/font-awesome/4.7.0/css/font-awesome.min.css">
```

- This will load the stylesheet from a Content Delivery Network (CDN), which is a web service intended to quickly serve files commonly used by multiple websites. In particular, browsers will *cache* (save) copies of these files locally, so that when you visit a different website (or the same website for a second time), the file will already be downloaded. We will load most of the CSS and JavaScript frameworks used in this class from a CDN.

The `.min` before the `.css` extension in the filename is a convention to indicate that the file is *minimized*: all extraneous spaces, comments, etc. have been removed to make the file as small and quick to download as possible.

Font Awesome icons can then be included in your HTML by including an element with an appropriate CSS class (that sets the style of that content to be the correct font with the correct “character” content). For example, you can include a Universal Access icon with


```
<i class="fa fa-universal-access" aria-hidden="true"></i>
```

- Because the icon has no semantic meaning beyond an appearance, it is often included with a semantically-deficient (empty) `<i>` element. Remember to include an `aria-hidden` attribute so that a screen-reader won't try to read the strange letter!

Other icon-based fonts include Glyphicons and Octicons.

Backgrounds and Images

You have previously seen how to use the `background-color` property to color the background of a particular element. However, CSS supports a much wider list of background-related properties. For example, the `background-image` property will allow you to set an image as the background of an element:

```
header {
  background-image: url('../img/page-banner.png');
}
```

This property takes as a value a `url()` data type, which is written like a function whose parameter is a string with the URI of the resource to load. These URIs can be absolute (e.g., `http://...`), or relative **to the location of the stylesheet** (not to the web page!—you may need to use `..` if your `.css` file is inside a `css/` directory).

There are additional properties used to customize background images, including where it should be positioned in the element (e.g., centered), how large the image should be, whether it should repeat, whether it should scroll with the page, etc.

```
header {
  background-image: url('../img/page-banner.png');
  background-position: center top; /* align to center top */
  background-size: cover; /* stretch so element is filled; preserves ratio (img may be cropped) */
  background-repeat: no-repeat; /* don't repeat */
  background-attachment: fixed; /* stay still when window scrolls */
  background-color: beige; /* can still have this for anything the image doesn't cover
                             (or for transparent images) */
}
```

This is a lot of properties! To understand all the options for their values, read through the documentation and examples (also here and here).

To try and make things easier, CSS also includes a **shorthand property** called just **background**. Shorthand properties allow you to specify multiple properties at once, in order to keep your code more compact (if somewhat less readable).

Shorthand properties values are written as a *space-separated list of values*; for example, the above is equivalent to:

```
header {
  background: url('../img/page-banner.png') top center / cover no-repeat fixed be
}
```

- The `background-position` and `background-size` are separated by a `/` since they both can have more than one value.
- You can include some or all of the available background properties in the shorthand. Unlike most shorthand properties, the `background` properties can go in any order (though the above is recommended).

Note that a shorthand property is interpreted as writing out all of the rules it replaces; so will *replace* any previous properties within the same rule:

```
body {
  background-color: gold;
  background: mediumpurple; /* later property override previous ones */
                          /* page will have a purple background */
}
```

Additionally, all of the `background` properties support multiple backgrounds by using a *comma-separated* list of values. This can enable you to easily overlay partially-transparent images on top of each other, similar to using layers in Photoshop.

There are many, many other properties you can use as well to style your page. Be sure and look through all the documentation and examples, and explore the source code of existing pages to see how they achieved particular effects!

Resources

- CSS Diner a fun game for practicing with different CSS selectors
- CSS Units and Values (MDN)
- The Code Side of Color
- CSS Backgrounds (MDN)

Chapter 7

CSS Layouts

The previous chapters have discussed how to use CSS to specify the appearance of individual html elements (e.g., text size, color, backgrounds, etc). This chapter details how to use CSS to declare where HTML elements should *appear* on a web page!

7.1 Block vs. Inline

Without any CSS, html elements follow a default **flow** on the page based on the order they appear in the HTML. Layout is based on whether the element is a *block element* or an *inline element*.

As mentioned in Chapter 3, **inline** elements (e.g., ``, `<a>`, ``) are put next to each other on a single line (left to right, unless you specify a right-to-left language). **Block** elements (`<p>`, ``, `<div>`) are placed on subsequent “lines”, from top to bottom.

```
<div>Block element</div>
<div>Block element</div>
```



Figure 7.1: Example of block elements, placed on top of each other.

```
<span>Inline element</span>
<span>Other inline element</span>
```

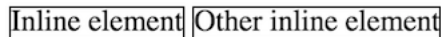


Figure 7.2: Example of inline elements, next to each other (even if the code is on separate lines).

However, you can force an element to be either **block** or **inline** by declaring the **display** CSS property:

```
.inlined {
  display: inline;
}


```

All HTML element (*including text!*) include an imaginary **box** around their content. Elements are laid out with their boxes next to each other (side by side for **inline**, stacked on top for **block**). These boxes are normally just large enough to contain the content inside the element, but you can use CSS to alter the size of and spacing between these boxes in order to influence the layout.

First off, you can set the **width** and **height** of elements explicitly, though be careful when you do this. If your **width** and **height** are too small, the element's content will be clipped by default (a behavior controlled by the **overflow** property). It's generally best to set only the width **or** the height, but not both. You can also specify a **min-width** or **min-height** to ensure that the width or height is at least a particular size. Conversely, you can use **max-width** and **max-height** to constrain the size of the element.

In order to adjust the spacing between boxes, you can manipulate one of 3 properties:




Figure 7.4: A diagram of the box model properties

Padding

The **padding** is the space between the content and the border (e.g., the edge of the box).

It is possible to specify the padding of each side of the box individually, or a uniform padding for the entire element:




Figure 7.5: An element's padding.

```
/* specify each side individually */
div {
  padding-top: 1em;
  padding-bottom: 1em;
  padding-left: 2em;
  padding-right: 0; /* no units needed on 0 */
}

/* specify one value for all sides at once */
div {
  padding: 1.5em;
}

/* specify one value for top/bottom (first)
   and one for left/right (second) */
div {
  padding: 1em 2em;
}
```

Border

The **border** (edge of the box) can be made visible and styled in terms of its width, color, and “style”, listed in that order:

```
.boxed {
  border: 2px dashed black; /* border on all sides */
}

.underlined {
  border-bottom: 1px solid red; /* border one side */
}
```

```
.something { /* control border properties separately */
  border-top-width: 4px;
  border-top-color: blue;
  border-top-style: dotted;
  border-radius: 4px; /* rounded corners! */
}
```

Margin

Finally, the **margin** specifies the space *between* this box and other nearby boxes. **margin** is declared in an equivalent manner to **padding**.




Figure 7.6: An element's margins.

Note that browsers typically collapse (overlap) the margins of adjacent elements. For example, if you have two paragraphs on top of one another, and you set **margin-bottom** on the first and **margin-top** on the second, most browsers will overlap those margins and just use the larger of the two values to determine the spacing.

Box-Sizing

An element's **padding**, **border**, and **margin** can be used to put space between element content on the page. However, when you assign an explicit **width** or **height** to an element, the dimension you specify **does not include** the padding or border when calculating the size of the element on the page! That is, if you have an element with the properties

```
.my-box {
  width: 100px;
  padding: 10px; /* includes both left and right */
}
```

Then the element will take up 120px on the screen: the width plus the left and right padding.

However, when specifying more complex or responsive layouts, it's often useful to have `width` represent the entire width of the box, and not need to account for the border and padding separately in calculations. You can do this using the `box-sizing` property—a value of `border-box` will indicate that specified *size* of the box (e.g., the `width`) should include the size of the padding and border when determining the content area.

It's common to want to apply this property to **all** of the elements on the page, which you can do with the `*` selector (like a wildcard from the command-line!):

```
* {  
    box-sizing: border-box; /* all elements include border and padding in size */  
}
```

This is a common enough change that you may wish to include it in *all* of your `.css` files!

7.3 Changing the Flow

Specifying the `display` style and `box` properties will adjust the layout of HTML elements, but they are still following the browser's default *flow*. The layout rules will still apply, and elements are influenced by the amount of content and the size of the browser (e.g., for when inline elements “wrap”).

However, it is possible to position elements outside of the normal flow by specifying the `position` property.

For example, giving an element a `position:fixed` property will specify a “fixed” position of the element *relative to the browser window*. It will no move no matter where it appears in the HTML or where the browser scrolls. See this [Code Pen](#) for an example.

- In order to specify the location for a **fixed** element, use the `top`, `left`, `bottom`, and/or `right` properties to specify distance from the appropriate edge of the browser window:

```
/* make the <nav> element fixed at the top of the browser window */  
nav {  
    position: fixed;  
    top: 0; /* 0px from the top */  
    left: 0; /* 0px from the left */  
    width: 100%; /* same as parent, useful for spanning the page */  
}
```


You can also specify an element's `position` to be **relative**, meaning *relative to its normal position*. Note that this leaves the element within normal flow (e.g., for how its padding affects other elements around it). See this Code Pen to explore this option.

Finally, you can specify an element's `position` to be **absolute**, meaning *relative to it (positioned) parent element*. If the parent has not been explicitly positioned (has a declared `position` property), the element is positioned relative to the root element. Try it out in this Code Pen. This mostly remove an element from normal flow, though its parent may still be part of the flow and thus may influence the absolutely positioned element's location.

Floating

You can also remove an element from its normal position in the *flow* by making it **float**. This is commonly done with pictures, but can be done with any element (such as `<div>`). A floated element is shoved over to one side of the screen, with the rest of the content wrapping around it:

```
.floating-image {  
  float: right;  
  margin: 1em; /* for spacing */  
}
```

Content will continue to sit along side a floated element until it “clears” it (gets past it to the next line). You can also force content to “clear” a float by using the `clear` property. An element with this property *cannot* have other elements floating to the indicated side:

```
.clear-float {  
  clear: both; /* do not allow floating elements on either side */  
}
```

The `float` property is good for when you simply want some content to sit off to the side. But you should **not** try to use this property for more complex layouts (e.g., multi-column structures); there are better solutions for that.

7.4 Flexbox

The `position` and `float` properties allow you to have individual elements break out of the normal page flow. While it is possible to combine these to produce complex effects such as **multi-column layouts**, this approach is fraught with peril and bugs due to browser inconsistencies. In response, CSS3 has introduced new standards specifically designed for non-linear layouts called **Flexbox**. The Flexbox layout allows you to efficiently lay out elements inside a container (e.g.,

columns inside a page) so that the space is *flexibly* distributed. This provides additional advantages such as ensuring that columns have matching heights.

Flexbox is a new standard that is now supported by most modern browsers; it has a buggy implementation in Microsoft IE, but is supported in the standards-compliant Edge. For older browsers, you can instead rely on a grid system from one of the popular CSS Frameworks such as Bootstrap.

Despite its capabilities, Flexbox still is designed primarily for one-directional flows (e.g., having one row of columns). To handle true grid-like layouts, browsers are adopting *another* emerging standard called **Grid**. The Grid framework shares some conceptual similarities to Flexbox (configuring child elements inside of a parent container), but utilizes a different set of properties. Learning one should make it easy to pick up the other. Note that the grid framework is less well supported than even Flexbox (it is not supported by IE, Edge, or common older Android devices), so should be used with caution.

To use a Flexbox layout, you need to style *two* different classes of elements: a **container** (or **parent**) element that acts as a holder for the **item** (or **child**) elements—the child elements are *nested* inside of the parent:

```
<div class="flex-container"> <!-- Parent -->
  <div class="flex-item">Child 1</div>
  <div class="flex-item">Child 2</div>
  <div class="flex-item">Child 3</div>
</div>
```




Figure 7.7: An example of a simple Flexbox layout.

Note that the “outer” parent element has one class (e.g., `flex-container`, but you can call it whatever you want), and the “inner” child elements share another (e.g., `flex-item`). Creating an effective Flexbox layout requires you to specify properties for *both* of these classes. You most often use `<div>` elements for the parent and child elements (and the child of course can have further content, including more divs, nested within it).

In order to use a Flexbox layout, give the *parent* element the **`display:flex`** property. This will cause the contents of that parent element to be laid out according to a “flex flow”:

```
.flex-container { /* my flexbox container class */
  display: flex;
}
```

A flex flow will lay out items *horizontally* (even if they are block elements!), though you can adjust this by specifying the *parent's* `flex-direction` property.

By default, a flex container will fill its parent element (the whole page if the container is in the `<body>`), and the child items will be sized based on their content like normal. While it is possible to then use dimensional properties such as `width` and `height` to size the children within the horizontal layout, Flexbox provides further options that make it more *flexible*.

Any *immediate child* of the flexbox container can use additional properties to define how that particular item should be layed out within the container. There are three main properties used by flex **items**:

```
* { box-sizing: border-box; } /* recommended for item sizing */

.flex-item {
  flex-grow: 1; /* get 1 share of extra space */
  flex-shrink: 0; /* do not shrink if items overflow container */
  flex-basis: 33%; /* take up 33% of parent initially */
}
```

- **flex-grow** specifies what “share” or ratio of any extra space in the container the item should take up. That is, if the container is `500px` wide, but the items only takes up `400px` of space, this property determines how much of the remaining `100px` is given to the item.

The value is a unitless number (e.g., 1 or 2, defaulting to 0), and the amount of remaining space is divided up *proportionally* among the items with a `flex-grow`. So an item with `flex-grow:2` will get twice as much of the remaining space as an item with `flex-grow:1`. If there are 4 items and `100px` of space remaining, giving each item `flex-grow:1` will cause each item to get `25px` ($100/4$) of the extra space. If one of the items has `flex-grow:2`, then it will get `40px` ($\frac{2}{1+1+1+2} = \frac{2}{5} = 40\%$) of the extra space, while the other three will only get `20px`.

In practice, you can give each item a property `flex-grow:1` to have them take up an equal amount of space in the container.

- **flex-shrink** works similar to `flex-grow`, but in reverse. It takes as a value a number (default to 1), which determine what “share” or ratio it should shrink by in order to accommodate any overflow space. If the specified dimensions of the items exceeds the dimensions of the container (e.g., 4 `100px` items in a `300px` container would have `100px` of “overflow”), the `flex-shrink` factor indicates how much size needs to be “taken off” the item. A higher number indicates a greater amount of shrinkage.

In practice, you will often leave this property at default (by not specifying it), *except* when you want to make sure that an item does NOT shrink by giving it `flex-shrink:0`.

- **flex-basis** allows you to specify the “initial” dimensions of a particular item. This is similar in concept to the `width` property, except that `flex-basis` is more flexible (e.g., if you change the `flex-direction` you don’t also have to change from `width` to `height`). Note that this value can be an dimensional measurement (absolute units like `100px`, or a relative unit like `25%`).

In practice, using percentages for the `flex-basis` will let you easily size the columns of your layout.




Figure 7.8: Top: visual example of `flex-grow`. Bottom: visual example of `flex-shrink`. Notice how much extra “space” each item has after the text content.

There is also a *shortcut property* `flex` that allows you to specify all three values at once: give the `flex-grow`, `flex-shrink`, and `flex-basis` values separated by spaces (the second two being optional if you want to use the default values).

The Flexbox framework also provides a number of additional properties that you can specify on the **container** to customize how items of different sizes are organized on the screen:

- **justify-content** specifies how the items should be spread out across the container. Note that items that have `flex-grow:1` will use up the extra space, making this less relevant.
- **align-items** is used to specify “cross-axis” alignment (e.g., the vertical alignment of items for a horizontal row).
- **flex-wrap** is used to have items “wrap around” to the next line when they overflow the container *instead of* shrinking to fit. You can then use the **align-content** property to specify how these “rows” should be aligned within the container.

While it may seem like a lot of options, Flexbox layouts will allow you to easily create layouts (such as multi-column pages) that are otherwise very difficult with the regular box model. Moreover, these layouts will be flexible, and can easily be made **responsive** for different devices and screen sizes.

Resources

- The Box Model (MDN)
- The CSS Box Model (CSS-Tricks)
- A Complete Guide to Flexbox The best explanation of Flexbox properties you'll find.
- A Complete Guide to Grid A similar explanation, but for the Grid framework (not discussed here).

Chapter 8

Responsive CSS

These days the majority of people accessing any web site you build will be using a device with a small screen, such as a mobile phone. But phones come in a wide range of sizes and resolutions, and many people will still access that same site from a laptop or desktop with a much larger monitor (as well as other capabilities, such as a mouse instead of a touchscreen). Different screens may require different visual appearances: for example, a three-column layout would be hard to read on a mobile phone! This poses an interesting *design* dilemma: how do you build one site that looks good and works well on both tiny phones and gigantic desktop monitors?

The modern solution to this problem is **Responsive Web Design**, which involves using CSS to specify *flexible* layouts that will adjust to the size of the display: content can be laid out one column on a small mobile screen, but three columns on a large desktop.

This chapter discusses CSS techniques used to create **responsive** web sites. These techniques underlie popular CSS frameworks, so it is important to understand them even if you rely on such tools.

8.1 Mobile-First Design

Responsive design is often framed as a technique to “make it *also* work on mobile”. This approach feels easy since websites are usually developed and tested on desktops, and follows from the software principle of *graceful degradation* (systems should maintain functionality as portions break down, such as the “capability” of having screen real estate).

But since websites are more likely to be visited on mobile devices, a better approach is to instead utilize **mobile-first design**. This is the idea that you




Figure 8.1: Example of a responsive web site on multiple devices, with content and layout changing at each size. From responsivedesign.is.

should develop a website so its content and purpose is effectively presented on mobile devices (e.g., the most restricted in terms of screen size, capabilities, etc). Only once this base level of functionality is in place should you add features to make it *also* look good on larger devices such as desktops. This approach is also known as *progressive enhancement*: provide the core functionality, and then add “extra” features as more capabilities become available.

Rather than viewing mobile devices as “losing” features, look at desktops as “gaining” features! This will help you to focus on better supporting more common mobile devices.

Remember: working on your personal machine doesn’t ensure that it will work on anyone else’s!

Mobile-First Design Principles

While there is no magic formula for designing websites to support mobile devices, there are a few general principles you should follow:

- **Layout:** On mobile devices, blocks of content should stack on top of each other, rather than sitting side by side in columns—mobile devices want to only scroll on one axis. **fixed** content should be kept to a minimum, as it reduces the amount of scrollable screen real-estate.

As you gain more screen space on desktops, you will *want* to break content up into columns or otherwise constrain its width so that it doesn't stretch to ridiculous lengths; this helps with readability.

- **Media:** Small screens don't have enough space to necessitate very large, high-resolution images and video. Moreover, large images have large file sizes, and so will take a long time to download on slow mobile connections (not to mention eating away at limited data plans). Use compressed or lower-resolution images on mobile, and consider using background colors or linear gradient fills instead of background images. You can use higher-resolution media (and more of it!) on desktops, which usually have higher bandwidth available for downloading.

Page bloat is a real problem. You don't need huge images... or possibly any images at all!

- **Fonts:** Make sure to use a large enough font that it is readable on small screens... but don't make headings or callout text *too* large so that you lose that precious real estate! You can make them more styled and prominent on desktop, where there is room for such flourishes. Be sure to use relative units to accommodate mobile user preferences and screen size variation. Also remember that special web fonts you may be downloading will also take up extra bandwidth!
- **Navigation:** Site navigation links take up a lot of room on small screens and may end up wrapping to multiple lines. Use small tab bars, or menu icons (e.g., the "hamburger icon") to show complex menus on command. Most CSS frameworks provide some kind of collapsible navigation for mobile devices.
- **Input and Interaction:** Tap/click targets need to be large-enough on mobile to select using a finger, especially for people with poor eyesight or thick fingers. Tiny icons placed right next to one another, or one-word hyperlinks are difficult to select accurately. Specifying a data type on form fields (e.g., email address, phone number, date, integer) also generates optimized on-screen keyboards, making data entry much easier.
- **Content:** For some sites, you may even want to adjust what content is shown to mobile users as opposed to desktop users. For example, a phone number might become a large telephone icon with a `tel:` hyperlink on mobile phones, but simply appear as a normal telephone number on desktop displays.

Specifying Viewport

Mobile web browsers will do some work on their own to adjust the web page in response to screen size—primarily by “shrinking” the content to fit. This often produces the effect of the website being “zoomed out” and the user enlarging

the web page to a readable size and then scrolling around the page to view the content. While it may “work” it is not an ideal user interaction—this behavior can also interfere with attempts to be explicit about how webpages should adjust to the size of the screen.

To fix this, you need to specify the viewport size and scale by including an appropriate `<meta>` element in your HTML:

```
<head>
  <meta charset="utf-8"> <!-- always need this -->
  <meta name="viewport" content="width=device-width, initial-scale=1, shrink-to-fit=

  <!-- more head elements, including <link> ... -->
</head>
```

Including this element will keep the text size from adjusting to the browser’s width (though *very* narrow browsers may still shrink the text).

While not technically part of any web standard, the `viewport` meta element was introduced by Safari and is supported by most mobile browsers. The `content` attribute for the above meta tag sets 3 properties for mobile browsers: `width` (how big the viewport should be, specified to be the size of the device screen), `initial-scale` (how much to initially “zoom” the page, specified to be 1x or no zoom), and `shrink-to-fit` (tells Safari 9.0+ not to shrink the content to fit).

You should include the `viewport` meta tag in all of your responsive pages. Make it part of your default HTML template!

8.2 Media Queries

In order to define a webpage with a responsive appearance, you need to be able to *conditionally* change the applied style rules depending on the size of the screen (or browser window). We can conditionally apply CSS rules by using **media queries**. A *media query* is a bit like an `if` statement in CSS: it specifies a *condition* and the rules (*selectors* and *properties*) that should apply when that condition is true.

Media queries have syntax similar to the following:

```
/* A normal CSS rule, will apply to all screen sizes */
body {
  font-size: 14px;
}

/* A Media Query */
@media (min-width: 768px)
```

```

{
    /* these rules apply ONLY on screens 768px and wider */

    /* a normal CSS rule */
    body {
        font-size: 18px;
        background-color: beige;
    }
    /* another CSS rule */
    .mobile-call-icon {
        display: none; /* don't show on large displays */
    }
}

```

A media query is structurally similar to a normal CSS rule. The “selector” is written as `@media`, indicating that this is a media query not a normal CSS rule. The `@media` is then followed by a *query expression*, somewhat similar to the boolean expressions used in `if` statements. Expressions are written in parentheses, with the media feature to check, followed by a colon (`:`), followed by the value to check against. There are no *relational operators* (no `>` or `<`) in media queries, so you use media features with names such as `min-width` and `max-width` (to represent `width > x` and `width < x` respectively).

- Media feature comparisons are not strict inequalities, so `min-width: 1000px` can be read as “width greater than or equal to 1000px”.
- It is also possible to use the *logical operators* `and` and `not` to combine media feature checks. You can produce an “or” operator using a group selector (a comma `,`).

```

/* style rules for screens between 768px and 992px */
@media (min-width: 768px) and (max-width: 992px) { }

/* style rules for screens larger than 700px OR in landscape orientation */
@media (min-width: 700px), (orientation: landscape) { }

```

The `@media` rule is followed by a `{ }` block, inside of which are listed *further regular CSS rules*. These rules will *only* be applied if the `@media` rule holds. If the `@media` rule does not apply, then these “inner” rules will be ignored. These “inner” rules can utilize all the selectors and properties value outside of media queries—think of them as mini “conditional” stylesheets!

You need to put *full rules* (including the selector!) inside of the media query’s body. You can’t just put a property, because the browser won’t know what elements to apply that property to.

Media queries follow the same ordering behavior as other CSS rules: **the last rule on the page wins**. In practice, this means that media queries can be

used to specify conditional rules that will *override* more “general rules”. So in the example above, the page is set to have a `font-size` of 14px using a rule that will apply on any screens. However, on larger screens, the media query will also apply, overriding that property to instead make the default font size 18px.

Following a **mobile-first approach**, this means that your “normal” CSS should define the styling for a the page on a mobile device. Media queries can then be used to add successive sets of rules that will “replace” the mobile styling with properties specific to larger displays.

```
/* on small mobile devices, the header has a purple background */
header {
    font-size: 1.2rem;
    background-color: mediumpurple;
}

/* on 768px OR LARGER displays */
@media (min-width: 768px) {
    header {
        font-size: 1.5rem; /* make the header larger font on larger displays */
    }
}

/* on 992px OR LARGER displays */
@media (min-width: 992px) {
    header {
        background-image: url('../img/banner.png') /* use background image */
    }
}
```

- In this example, the `<header>` is given a simple purple background and default font size. When loaded on mobile devices, this is the only rule that applies (the media queries aren’t valid), so that is all the styling that occurs.
- But on devices 768px or wider (like a tablet), the first media query is activated. This will then run a second rule that applies to the `<header>`, overriding the font to be larger (1.5rem)—it’s as if we had listed those two `header` rules one after another, and the later one wins. But the first rule continues to apply, making the background purple.
- Finally, a device 992px or wider (like a desktop computer), will cause *both* of the media queries to execute. (since a device whose width is greater than 992px is ALSO greater than 768px). Thus the `<header>` will be given a purple background and default font size, which will then be overridden to be a larger font by the first media query. The second media query will add an additional property (a banner image background), which will combine with the previous purple background (e.g., if the banner has any transparency). So on large displays, there will be a banner background on

top of purple, with text in a larger (1.5rem) font.

This structure (starting with “default” mobile rules and then using media queries with *increasingly larger* min-width values) produces an effective mobile-first approach and clean way of organizing how the appearance will change as the screen gets larger. Note that while you can define as many media queries as you want, most professionals define only a few that match the common breakpoints between phone, tablet, and desktop screen widths. Since the media queries need to come **after** the mobile rules, they are often included at the end of the stylesheet—give all the mobile rules first, then all the media queries that modify them. (You can and should put multiple rules in each media query).

(Alternatively, it’s also reasonable to organize your stylesheets based on page “section”, with the mobile rules for that section given before the media queries for that section. E.g., after all your rules for creating the page header, put the media query with variations for desktop headers. Use whatever organization makes your code readable and maintainable, and leave plenty of comments to guide the reader!)

Example: Responsive Flexbox

As another example of using media queries to produce a responsive website, consider how they can combine with Flexbox to produce a single-column layout on mobile devices, but a multi-column layout on larger displays. Because a Flexbox layout is just a property applied to existing elements, we can effectively “turn on” the flex layout by using media queries.

Consider some simple HTML:

```
<div class="row">
  <div class="column">column ONE content</div>
  <div class="column">column TWO content</div>
  <div class="column">column THREE content</div>
  <div class="column">column FOUR content</div>
</div>
```

By default (without any style rules applied to the `.row` or `.column` classes), the four inner `<div>` elements (as block elements) will stack on top of each other. This is the behavior you want on mobile narrow screens, so now additional CSS or Flexbox usage is needed.

In order to turn these divs into columns on larger displays, introduce a media query that applies the flex layout to the `.row` (which acts as the *container*), thereby lining up the `.column` elements (which act as the *items*)

```
/* on devices 768px OR WIDER */
@media (min-width: 768px)
{
```

```

    .row { /* row becomes a flexbox container */
        display: flex;
    }

    .column { /* column becomes a flexbox item */
        flex-grow: 1; /* make the columns grow equally to fill the row */
    }
}

```

You can also add an additional media query at another breakpoint so that the layout starts out stacked, then switches to two columns on medium-sized displays (leading to a two-by-two grid), and *then* switches to four columns on large displays:

```

/* on devices 768px OR WIDER */
@media (min-width: 768px)
{
    .row { /* row is a flexbox container */
        display: flex;
        flex-wrap: wrap; /* wrap extra items to the next "line" */
    }

    .column { /* column is a flexbox item */
        flex-basis: 50%; /* columns take up 50% of parent by default */
        flex-grow: 1;
    }
}

@media (min-width: 1200px) {
    .column {
        flex-basis: auto; /* columns are automatically sized based on content */
    }
}

```

In this case, the first media query (768px + or medium-sized displays) applies the `flex` layout and specifies that the items should `wrap` if they overflow the container... which they will, since each item has a default `flex-basis` size of 50% of the container. This will cause each of the 4 items to take up 50% of the parent, wrapping around to the next line.

Then when the screen is larger (1200px + or large-size displays), the second media query is applied and *overrides* the `flex-basis` so that it will automatically calculate based on the content size, rather than being 50% of the parent. That way as long as the columns fit within the parent, they will all line up in a row (they have the `flex-grow` property to make them equally spread out).

The best way to get a feel for how this works is to see it in action: see **this**

CodePen for an example of the above behavior—resize the browser and watch the layout change!

Media rules are a powerful and declarative way to create a single page that looks great on everything from a small mobile touchscreen to a large desktop monitor with a mouse, and form the foundation for responsive CSS frameworks that can help you easily create fantastic looking pages.

Resources

- [Responsive Web Design Basics \(Google\)](#)
- [Using media queries \(MDN\)](#)

Chapter 9

CSS Frameworks

The previous chapters have covered how to use CSS to make your page both stylish and responsive. While it is certainly possible to implement an entire site by creating your own CSS rules, that can quickly get tedious: often you’d like a page to have a “standard” set of rules (because the browser’s style-less appearance) and then customize those rules to your liking.

For this reason, most professional web developers utilize an existing **CSS Framework** instead. A CSS Framework is a stylesheet (a `.css` file!) that contains a large list of pre-defined rules that you can apply to your page. CSS frameworks provide a number of benefits:

- *Applies attractive default styling to all HTML elements:* CSS frameworks make your pages instantly look better through a bunch of element selector rules. Frameworks provide pleasant default fonts, line spacing, spacing, and link styling without any extra effort on your part.
- *Provides style classes for common UI components:* framework stylesheets will also include CSS classes you can add to your markup to easily include badges, in-page tabs, drop-down buttons, multi-column layouts, and more. Frameworks enable you to add style your page by specifying CSS classes, rather than needing to define multiple CSS rules for a single effect.

CSS frameworks are thus designed to make your life easier and your development more efficient—while still enabling you to provide your own customizations and styling with all of the power of CSS.

There are many different CSS frameworks to choose from. Some popular ones include:

- **Bootstrap** is the most commonly used CSS framework on the web. Sometimes called “Twitter Bootstrap”, it was originally created at Twitter to enforce some consistency among their internal tools, but was released as an open-source project in 2011. Its popularity has benefits and drawbacks:

it's very well tested, documented, and supported, but it's also so prevalent that its default look has become cliché.

The latest release of Bootstrap, v4, is currently in beta. This version includes an update look and feel, and makes it relatively easy to create your own customized version. Bootstrap 4 also uses Flexbox as a foundation for its grid system, offering some better performance. For this and other reasons, Bootstrap 4 does **not support** IE 9 or earlier browsers. For older browsers, you'd want to use Bootstrap 3.

This text discusses Bootstrap 4 as an example. Be careful that you are looking at information about the right version if you seek out documentation or help online!

- **Foundation** is Bootstrap's chief rival, and has a reputation for being more ahead-of-the-curve than Bootstrap (introducing new features sooner). For example, it was the first framework to use a responsive mobile-first design, and provided a Flexbox-based grid long before Bootstrap. Foundation has most of the same UI elements as Bootstrap, but with a different look and feel (that can be customized through a web-based tool).
- **Material Components for the Web (MCW)** is an official implementation of Google's Material Design visual language. This is the look-and-feel found in most Google products and Android applications. Material Design is very opinionated so MCW is very difficult to customize. The MCW style class names are also very verbose, as they follow the Block, Element, Modifier (BEM) naming scheme.
- **Materialize** is the other popular Material Design implementation. This is an open-source project, so it is not provided nor supported by Google. However, it is structurally similar to Bootstrap, making it easy to learn and popular among people who know that framework.
- **normalize.css** is not so much a full framework as a small utility that performs browser normalization (also called a "reset"). **normalize.css** standardizes the quirks across browsers and making it so that the same HTML and CSS is displayed more consistently and without errors. Note that most other frameworks include some version of this package (for example, Bootstrap contains a fork called Reboot).

If you choose **not** to use a framework, you should still include **normalize** to make your sites consistent!

Other frameworks exist as well, each with its own benefits and drawbacks (e.g., Skeleton is designed to be as lightweight as possible, but at the cost of features).

9.1 Using a Framework

The important thing to realize is that a CSS framework is **just a stylesheet with a bunch of rules that someone else wrote for you**. There's nothing magic about them. You can look at the stylesheet and see all the rules that have been defined, and it's all stuff you could have written yourself (though some of it can be quite complex). But those rules have been crafted by professionals and tested on a wide array of browsers to ensure consistent results, so it's a good idea to build on top of them rather than trying to re-invent the wheel.

For example, you can view the `normalize.css` file at <https://necolas.github.io/normalize.css/7.0.0/normalize.css> and see that it looks exactly like the CSS you might write on your own!

Including a Framework

Because a CSS framework is *just a CSS file*, you include one in your page using a `<link>` element just like with any other stylesheet:

```
<head>
  <!--... other elements here...-->

  <link rel="stylesheet"
        href="https://maxcdn.bootstrapcdn.com/bootstrap/4.0.0-beta/css/bootstrap.css">
  <link rel="stylesheet" href="css/my-style.css">
</head>
```

Important! Note that you link the CSS framework *before* you link your own stylesheet. Remember that CSS is read top-to-bottom, and that the last rule on the page wins. By putting your stylesheet second, any rules you define will *override* the rules specified in the framework, thereby allowing you to continue to customize the appearance.

- For example, `bootstrap.css` gives the `<body>` a color of `#212529` (very dark gray); you may instead want to have a different default text color::

```
/* my-style.css */
body { color: blue; }
```

If you included your stylesheet *first* in the `<head>`, then your rule would be applied initially but would be overridden by the `bootstrap.css` file, causing your rule to seem to have no effect! Thus you should **always** put your own custom CSS as the *last* `<link>` in the page, after any CSS frameworks.

It is most common to include the **minified** versions of CSS frameworks: these are usually named as **.min.css** files (e.g., `bootstrap.min.css`). A minified file

is one that has extra comments and white space removed, creating a file with a much smaller size (remember that each space and line break is a character so takes up a byte! Bootstrap goes from 160 kilobytes to about 127 kilobytes when minimized). Because web pages have to be downloaded, smaller pages will be downloaded faster, and thus will be “quicker” for the user. While an unminimized CSS file won’t “break the bank”, the extra white space doesn’t provide anything that you need.

- You can almost always convert between the minimized and non-minimized file just by changing the file name between `file.css` and `file.min.css`.

There are a few different ways of accessing a CSS framework’s stylesheet in order to link it into your page:

Linking to a CDN

The easiest option is usually to link to a Content Delivery Network (CDN), which is a web service intended to quickly serve files commonly used by multiple websites. To do this, you simply provide the URL for the file on the CDN’s servers as the `href` of your link (as in the example above).

CDNJS is a CDN for packages and libraries operated by Cloudflare, and is a good place to start if you’re unsure of the CDN link for a packages. Other packages may have different CDN hosting (e.g., Bootstrap recommends a link from its own BootstrapCDN).

In addition to being simple, CDNs also provide a distinct speed benefit: CDNs replicate their content to machines in several regions of the world, and use dynamic Domain Name Service (DNS) resolution to steer users to the machine nearest them. So a user in Australia might download the Bootstrap CSS from a server located in Singapore, while a user in France might get the same CSS file from a server located in Ireland. This helps increase the speed that the files are downloaded.

Additionally, CDNs allow a site to take advantage of **browser caching**: browsers will save previously downloaded versions of a file to avoid downloading it again. If multiple sites all use the CDN version of Bootstrap, then the browser only has to download that file the first time you visit one of those sites. With popular frameworks like Bootstrap, it’s highly likely that your user has already visited a site that links to Bootstrap’s CDN version, and thus the Bootstrap files are already in the user’s browser cache.

Finally, CDNs allow the frameworks to be accessed as a service, meaning that developers may be able to quickly patch bugs (if the fix doesn’t break existing code) without the developer having to do anything. (This feature can also be seen as a drawback, as you need to trust the library developer not to break anything or accidentally include malicious code).

The only real disadvantage to linking to a CDN is that it won't work when you are offline. If you commonly do your development offline, or if you are building a web application that is meant to run offline, you'll need to utilize a different method.

Downloading the Source

If you can't link to a CDN, you'll need to download the source file directly to your computer and saving it as a source code file in your project (often in a `lib/` folder). These files can be found on the framework's website, or often in their GitHub repo (look for the `.css` file, or for a `dist/` distribution folder).

In addition to being available offline, downloading CSS files directly provides a few other advantages:

- Frameworks may allow you to customize their content even before you include the file: selecting only the components you need, and adjusting base styling properties before downloading your customized source file.
- If you use a CSS pre-compiler such as SASS, you can often modify the source-code mix-ins to more easily provide your own customizations and override the default fonts, colors, sizes, etc. simply by changing some variables.
- If you use a build system such as gulp or webpack, then you can often combine *all* of your CSS files into a single file. This reduces the total number of files that need to be downloaded, which can increase the page load speed—especially on slower mobile networks.

Note that it is also possible to utilize a **package manager** such as npm to download and manage the source for frameworks. This has the advantage of enabling you to avoid adding large extraneous source files to your code repo, as well as managing any dependencies for those libraries.

Recall that you can install a package using npm, saving the dependency into your project's `package.json` file:

```
# install bootstrap
npm install --save bootstrap@4.0.0-beta
```

(This explicitly loads the beta of Bootstrap 4. See below for details).

This will install the library's source code into your project's `node_modules/` folder. You will then need to dig around inside that source folder for the `.css` file you wish to link (it is in no way consistent for all frameworks)

```
<!-- notice the RELATIVE path to the file -->
<link rel="stylesheet" href="node_modules/bootstrap/dist/css/bootstrap.min.css">
```

Because you've saved the dependency in your `package.json` file, you can use the `.gitignore` file to exclude the `node_modules/` folder from your code repo,

having new developers install the required framework via `npm install`. You can also use `npm` to easily upgrade the packages (such as when Bootstrap 4 leaves beta).

Important: if you are publishing a site to GitHub Pages, you **must** add the `node_modules` folders to your code repo so that it will get uploaded to GitHub and the file will be present on the server! For this reason, I recommend just using a CDN when working with GitHub Pages.

9.2 Bootstrap

This section discusses *some* of the features and uses of the Bootstrap CSS framework. Note that many other frameworks are used in a similar fashion (though the specific classes and components may be different).

You include the Bootstrap framework by linking its `.css` file to your page (whether from a CDN or a local file). You can get the link for the latest version from the Bootstrap Home Page.

```
<link rel="stylesheet"
      href="https://maxcdn.bootstrapcdn.com/bootstrap/4.0.0-beta/css/bootstrap.min.c
      integrity="sha384-/Y6pD6FV/Vv2HJnA6t+vslU6fwYXjCFtcEpHbNJ0lyAFsXTsjBbfaDjzALeQ
      crossorigin="anonymous">
```

(This `<link>` example has 2 additional (optional) attributes: `integrity` gives a cryptographic hash of so the browser can make sure you didn't download a malicious file by mistake, and `crossorigin` which says that no credentials are sent to the remote server—see the AJAX chapter for a more detailed discussion).

When you first include the Bootstrap framework, you'll immediately see the effects (usually for the better!): the default font has changed, margins are different, and so forth. The Bootstrap CSS file contains a number of rules that apply to elements without an extra work, including the `<body>` and headers. See the documentation for details about the default styling (as well as the entire "Content" section of the documentation).

Utility Classes

In addition to the base styling rules, Bootstrap also includes a huge number of rules that can be applied to elements by giving them a particular CSS **class**. These utility classes are what gives Bootstrap its power and flexibility.

For example, Bootstrap provides text utilities that allow you to style text just by giving it a CSS class, rather than needing to write a new rule for that element.

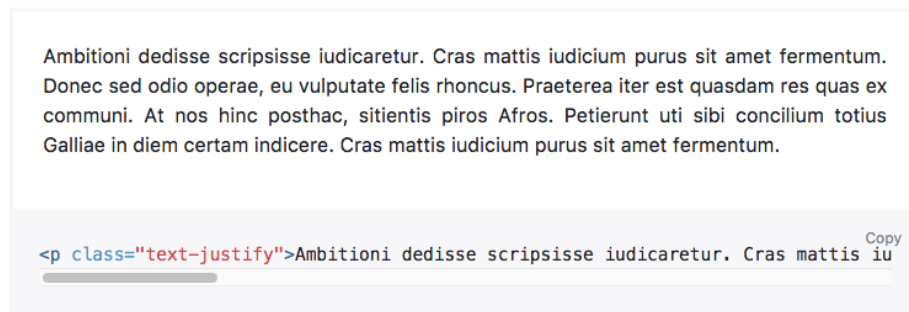


Figure 9.1: The `text-justify` utility class will justify a block of text. Example from the Bootstrap documentation.

While it would be possible to write such a rule yourself, Bootstrap provides these classes as “built-in”, so you can focus on the semantics of the content rather than the CSS.

Many elements need to be given an additional utility class to indicate that they should use Bootstrap styling. For example, in order to give a `<button>` a Bootstrap style, it needs to have the class `btn`. This allows you to utilize Bootstrap’s styling, but to not be forced to “override” that look and feel if you want to do something different.

```
<button>Plain button</button>
<button class="btn">Bootstrap button</button>
<button class="btn btn-primary">Primary colored button</button>
```




Figure 9.2: Bootstrap Button classes.

The last button in this example has an additional utility class `btn-primary` that specifies which color style it should use. Bootstrap provides a number

of pre-defined color styles, named after semantic meaning (e.g., `primary` or `warning`):



Figure 9.3: Bootstrap Button colors.

These same color classes can be applied to text elements (e.g., `text-primary`), element backgrounds (e.g., `bg-primary`), alert blocks (e.g., `alert-warning`), and more. Modifying the SASS source code for Bootstrap will allow you to easily customize these colors while keeping the theme consistent.

Note that you can even use the `btn` class to style other elements (e.g., `<a>` links) so that they *look* like buttons!

```
<a class="btn btn-primary" href="#">A link</a>
```

“Stacking” utility classes like this is quite common in Bootstrap; it is not unusual for an element to have 4 or 5 different utility classes to fully specify its appearance.

Bootstrap is filled with utility classes, both for generic styling (e.g., borders or position), and for specific types of content (e.g., lists, tables, forms). Check out the documentation and examples for details on how to style particular elements and achieve specific effects.

Bootstrap also includes utility classes specifically to support screen readers. For example, the `sr-only` class can be used to style an element so it is only perceivable to screen readers.

Utility classes do **not** mean you never have to write any CSS rules! While Bootstrap provides classes to handle lots of “common” styling, you’ll almost always need to add your own customizations on top of it to produce the appearance you want.

Responsive Design

A major feature of Bootstrap (and other CSS frameworks) is the ability to support **mobile-first, responsive webpages**. Many Bootstrap utility classes use *media queries* to cause them to respond to the device’s screen size.

Bootstrap’s responsive utilities utilize a predefined set of responsive breakpoints: screen widths that mark the “boundaries” between different size devices.

Screen Width	Abbreviation	Approximate Device Size
< 576px	(default)	Extra-small devices (portrait phones)

Screen Width	Abbreviation	Approximate Device Size
576px	sm	Small devices (landscape phones)
768px	md	Medium devices (tablets)
992px	lg	Large devices (desktops)
1200px	xl	Extra-large devices (large desktops)

All of these size specifications are “greater than or equal”, following a mobile first approach. Thus a responsive utility that applies to a medium (`md`) screen will also apply to large (`lg`) and extra-large (`xl`) screens.

Responsive utilities are named with the size’s abbreviation to specify what screen size that utility should apply on. For example, the `float-left` class will cause an element to float to the left on *any* sized screen (it has no size specification). On the other hand, the `float-md-left` class will cause the element to float **only on medium or larger screens**. Similarly, the `float-lg-left` class would cause the element to float only on *large or larger* screens. This would allow you to have an image that floats to the side on a bigger screen, but is flowed with the rest of the content when there isn’t enough room.

- The breakpoint abbreviations are consistent across Bootstrap (and even shared by other frameworks such as Materialize), and usually comes after the name of property to change but before the value to give it. For example, the `d-none` class gives an element a `display:none` on all devices, while `d-sm-none` would apply on small or larger screens.

Many (but not all!) Bootstrap classes support responsive breakpoints. The most common class to do this is the `container` class. Containers are the most basic layout element in Bootstrap, and are required for doing more complex layouts (such as Grids, below). The `container` class gives the element a fixed width that includes some padding on the sides, so that the content isn’t right up against the edge of the window. If the size of the window shrinks to the point where the padding would disappear, the `container` responsively changes size so that there remains padding even on the smaller screen!

- See this Code Pen for an example. if you slowly resize the browser, you’ll see the text “jump” to a smaller width as you pass each responsive breakpoint.
- `.container` elements are usually direct children of the `<body>` (e.g., `<header>`, `<main>`). However, if you want a background color or image to “full bleed” to the edges of the viewport, you can put the `.container` inside the another element that is styled with the background, as in the CodePen example.
- Bootstrap also provides an alternate `container-fluid` class that always has the same amount of padding around the content, with the content “reflowing” based on the browser size.

In general, every page you style with Bootstrap will have a Container (usually a `.container`) to hold the content.

The Grid

After Containers, most common use of the responsive utility classes in CSS frameworks such as Bootstrap is to provide a **grid layout**, similar to the multi-column layouts created with Flexbox in the previous chapters. The Bootstrap Grid system allows you to specify these kinds of complex, responsive layouts without needing to implement your own Flexbox containers and items.

You specify that a set of elements should be part of a “grid” by putting them inside a `.row` element (the `row` class ensures that the elements are lined up). A `.row` acts a lot like a “flex-container” (and in fact does have `display: flex`), where each of its children will be “flex items”. Note that the `.row` must be a *direct* child of a Container (`container` or `container-fluid`); this ensures that grid will be positioned and spaced within its parent correctly.

```
<div class="container"> <!-- container for the grid -->
  <div class="row"> <!-- a row of items -->
    <div class="col-6">Item 1</div>
    <div class="col-4">Item 2</div>
    <div class="col">Item 3</div>
  </div>
</div>
```

You specify how wide each item in the `.row` is by specifying the number of “columns” it spans. Think of the grid as containing **12** “slots”:

Each item in the grid is given a class that indicates how many of these slots it takes up (similar to merging cells in a spreadsheet):

These classes are named in the format `col-#`, where the `#` is how many column slots it should take. For example a `.col-4` will take up 4 columns, while a `.col-6` will take up 6 columns. A class of just `col` will take up an *equal share of the remaining columns* (similar to what happens with the `flex-grow` property!)

Thus the above HTML will produce elements with the following layout:

- Notice that “Item 1” takes up half (6/12) of the grid, “Item 2” takes up a third (4/12), and “Item 3” takes up the remainder (2/12).

You can also make this grid layout responsive by using responsive utility versions of the `col-#` classes. For example, a `.col-md-4` would take up 4 columns *only on medium or larger* screens, while a `.col-lg-6` would take up 6 columns *only on large or larger* screens. Note that a child of the `.row` that doesn’t have a column size will effectively act as its own row (on its own line).

Finally, note that you can specify *multiple responsive utilities* on each grid item! For example, you can specify that items should take up 6 columns on small or larger screens (`col-sm-6`), only but 3 columns on medium or larger screens (`col-3-sm`):

```
<div class="container">
  <div class="row">
    <div class="col-sm-6 col-md-3">Item 1</div>
    <div class="col-sm-6 col-md-3">Item 2</div>
    <div class="col-sm-6 col-md-3">Item 3</div>
    <div class="col-sm-6 col-md-3">Item 4</div>
  </div>
</div>
```

(See a CodePen of this code in action!)

The ordering of Bootstrap’s CSS means that the `col-sm-6` class will get overridden by the `col-md-3` class (since `md` is a more specific size than `sm`), so you can specify both classes and the correct one will apply—the order of the class names in the element doesn’t matter. Note that the “*or larger*” part of the responsive breakpoint definition also means that you don’t need to specify a column width for *each* different breakpoint, only for the smallest level at which the utility should apply (e.g., we don’t need `col-lg-` since it isn’t different than `col-md-`).

Bootstrap includes a number of utilities that can be used to customize how grids of content are displayed, including all of the features of Flexbox. For more details and options, see the Grid documentation.

Components

In addition to its basic and responsive utilities, Bootstrap provides a number of CSS classes that let you easily create more complex **components** or widgets: alerts, cards, navbars, collapsible boxes, image carousels, pop-up modals, and more. Each of these components is discussed in the documentation, and has its own particular element structure and CSS classes you need to use.

For example, you can create a card (a styled content container) by creating an element with the `card` class and `card-body` child. You can then include other nested elements with classes such as `card-img-top`, `card-title`, or `card-text` to produce further styled effects:

```
<div class="card"> <!-- the card -->
  
  <div class="card-body"> <!-- the "body" of the card -->
    <h4 class="card-title">A Card Title</h4>
    <p class="card-text">Some card text that appears below the card title.</p>
```

```

        <a href="#" class="card-link">Action Link 1</a>
        <a href="#" class="card-link">Action Link 2</a>
    </div>
</div>

```

(See this code in [CodePen](#)).

Many of these components include some form of **interactivity**: for example, the navbar will “collapse” into a hamburger menu on small devices, while modals add in-page pop-up dialogs.

CSS is not capable of adding this level of interactivity—that requires **JavaScript**. Thus to support these components, you also need to include the Bootstrap JavaScript library (as well as its dependency).

```

<!-- dependency: jQuery -->
<script src="https://code.jquery.com/jquery-3.2.1.slim.min.js"></script>
<!-- the Bootstrap JavaScript library -->
<script src="https://maxcdn.bootstrapcdn.com/bootstrap/4.0.0-beta/js/bootstrap.min.js"></script>

```

See the chapter on JavaScript for more details on including JavaScript libraries. Note that not all components or utilities require JavaScript.

However, with most Bootstrap components, you don’t actually need to write any JavaScript code. Instead, you just need to provide some particular data-HTML attributes, and the library will take care of everything else.

- A data attribute is an normal HTML attribute that begins with the text `data-` (e.g., `data-thingamabob` or `data-my-attribute`). The HTML specification explicitly says that attributes that begin with `data-` are legal but not part of the standard, thus enabling them to be used for storing app-specific information. They effectively act as custom “variables” for elements that do not violate the HTML standard.

For example, consider the below code for defining a modal and a button to open it:

```

<!-- note that the `id` attribute is set to "exampleModal" -->
<div id="exampleModal" class="modal fade" tabindex="-1" role="dialog" aria-labelledby="exampleModalLabel">
  <div class="modal-dialog" role="document">
    <div class="modal-content">
      <div class="modal-header">
        <h5 id="exampleModalLabel" class="modal-title">Modal title</h5>
        <button type="button" class="close" data-dismiss="modal" aria-label="Close">
          <span aria-hidden="true">&times;</span>
        </button>
      </div>
      <div class="modal-body">Some content shown in the body of the modal</div>
      <div class="modal-footer">

```

```

        <button type="button" class="btn btn-secondary" data-dismiss="modal">Close</button>
        <button type="button" class="btn btn-primary">Save changes</button>
      </div>
    </div>
  </div>
</div>

<!-- data-toggle must be set to "modal" and data-target must be set to
the value of the `id` attribute on the element that starts your modal markup -->
<button type="button" class="btn btn-primary" data-toggle="modal" data-target="#exampleModal">

```

(Try it out in CodePen).

There is a lot to unpack in this code, including styling classes and ARIA attributes (e.g., `role`, `aria-labelledby`, and `aria-hidden`). The important part is that the root `.modal` element is *hidden* (`display:none`) when the page first loads. But when you click on the `<button>`, the included JavaScript library sees that it has a `data-toggle="modal"` attribute, and determines that clicking the button should show the modal. The JavaScript uses the `data-target` attribute on that button to determine *which* modal to open (finding the modal with that `id`), and then shows it. The `data-dismiss` attribute on the first `<button>` inside the modal allows it to be closed.

The best way to understand these components is to start from the examples in the documentation, and then consider each element and attribute one at a time. Read the documentation!

The JavaScript included in these CSS frameworks allows you to easily add interactive features, but **they do not mix well with more advanced JavaScript frameworks like React**. To work with React, you'd need to switch instead to a library such as `reactstrap`, which uses the Bootstrap 4 CSS stylesheet, but re-implement the JavaScript portions to fit into the React framework.

Resources

- [Bootstrap 4 Documentation](#)
- [Materialize Documentation](#)
- [Foundation Documentation](#)
- [Material Components for the Web Documentation](#)
- [normalize.css Documentation](#)
- [Bootstrap 3 Tutorial \(w3schools\)](#) For an older version of bootstrap, but many of the components and examples are similar.

This chapter was borrowed to a large extent from a tutorial by David Stearns.

Chapter 10

JavaScript Fundamentals

HTML and CSS are *markup languages*, used to annotate websites to describe their meaning and appearance. But they are not “full” (i.e., Turing Complete) programming languages—they don’t have variables, control structures, or other features required for the computer to execute an *algorithm*. So to make *interactive* websites or complex web applications, you need a complete programming language. And the language used by browsers is **JavaScript**. This chapter introduces the fundamentals of the JavaScript language as used in web development, focusing on variables, data types, and basic control structures. Functions are introduced, but addressed in more detail in the next chapter.

Note course assumes you already are already familiar with introductory programming in a imperative language such as Java. It focuses on how JavaScript differs from languages such as Java, rather than foundational programming concepts. If you need review on those topics, you can check out a basic programming tutorial or review your notes from a previous programming course.

10.1 Programming with JavaScript

JavaScript is a **high-level, general-purpose programming language**, allowing you to declare instructions for the computer in an almost-human readable way (similar to Java). JavaScript is an *imperative* language, so you write algorithms as step-by-step instructions (lines of code) using JavaScript’s syntax, and the computer will interpret these instructions in order to execute the algorithm. Browsers are able to download scripts written in JavaScript, executing them line-by-line and using those instructions to manipulate what content is displayed.

Indeed, JavaScript is an **interpreted language**, in that the computer (specifically a *JavaScript Interpreter*) translates the high-level language into machine

language *on the fly at runtime*. The interpreter will read and execute one line of code at a time, executing that line before it even begins to consider the next. This is in contrast with *compiled languages* like C or Java that have the computer do the translation in one pass (at compile-time), and then only execute the program after the whole thing is converted to machine language.

- This on-the-fly translation means that interpreted languages like JavaScript are usually slower than compiled languages, but not enough that we'll notice (web browsers include *highly* optimized interpreters). The bigger drawback is that without the compile step, program errors appear at runtime rather than compile time, making them more difficult to catch and fix. As such, JavaScript developers make heavy use of various linting tools (which flag common problems that can lead to runtime errors), as well as automated testing systems to check against a variety of inputs.

Many JavaScript editors do this kind of error checking—for example, VS Code provides syntax error checking out of the box, with more detailed linting available through extensions. See this article for more code features that support programming in JavaScript, including IntelliSense and type checking.

Although JavaScript was designed for and is most commonly used within web browsers (which contain their own JavaScript Interpreters), it can also be executed on the command-line by using Node.js, allowing JavaScript to be a fully general language. Both techniques are described in this chapter.

History and Versions

The JavaScript language was developed by Brendan Eich (the co-founder of Mozilla) in 1995 while he was working for Netscape. The original prototype of the language was created in 10 days... a fact which may help explain some of the quirks in the language.

- Despite the names, *JavaScript* and the *Java* language have nothing to do with one another (and are in fact totally separate programming languages used in drastically different contexts). JavaScript was named after Java as a marketing ploy to cash in on the popularity of the latter. It is the programming equivalent of the Transmorphers movie.
- The JavaScript language is officially an implementation of ECMAScript (named for the European Computer Manufacturers Association) specification. Hence versions of JavaScript are labeled with the prefix “ES”. For example, this chapter includes describes ES5 (JavaScript 5) syntax.

Like HTML and CSS, the JavaScript language continues to be developed and refined through new versions. Each new version of JavaScript includes additional syntax shorts, specialized keywords, and additional core functions, but

otherwise are compatible with previous code. The main limitation on utilizing new JavaScript features is whether the *interpreters* found in web browsers are able to support them.

This course primarily utilizes syntax and features for ES5, which was introduced in 2009 and today is supported by all modern browsers (i.e., IE 10 or later). Later chapters will describe features of ES6 (also called ES2015), which was introduced in 2015 and works on many browsers (e.g., Microsoft Edge but not IE 11). ES7 was finalized in 2016 but is still not reliably supported.

- Note that Microsoft no longer supports earlier versions of IE, but a small percentage of desktops still run Windows XP with IE 8. If you must support these browsers, you can often use a “transpiler” like Babel to translate modern JavaScript source code into an earlier version.

10.2 Running JavaScript

There are two primary ways of executing code written in JavaScript:

In the Browser

JavaScript scripts are most commonly executed in a web browser as part of the browser rendering a web page. Since browsers render HTML content (in `.html` files), JavaScript scripts are included in that HTML by using a `<script>` tag and specifying the *relative* path to the a file containing the code (usually a `.js` file) to execute. When the browser renders the HTML (reading top to bottom) and gets to that tag, it will download and execute the specified script file using the JavaScript interpreter:

```
<!-- execute the script.js file -->
<script src="path/to/my/script.js"></script>
```

- Notice that the `<script>` element is *not* empty! It is possible to include JavaScript code directly inside the tag, but this is considered bad practice (keep concerns separated!) and should only be used for quick tests.

The `<script>` tag can be included anywhere in an HTML page. Most commonly it is either placed in the `<head>` in order for the script to be executed *before* the page content loads, or at the very end of the `<body>` in order for the script to be executed *after* the page content loads (and thus allows the JavaScript to immediately interact with the loaded HTML elements). We will most commonly be putting `<script>` tags at the end of the `<body>`:

```
<!DOCTYPE html>
<html>
<head>
```

```

<!-- include here to run before the page appears -->
<script src="js/script.js"></script>
</head>
<body>
  ... content ...

  <!-- include here to run "after" html appears -->
  <!-- we will usually do this -->
  <script src="js/script.js"></script>
</body>
<html>

```

- Most browsers also support adding a **defer** attribute to the `<script>` element that tells the browser to defer downloading and running the script until the page is fully loaded and rendered to the screen. With this attribute, you can put the `<script>` element in the `<head>` without any issues. However, this attribute is not supported in IE 9 or earlier; it is thus still best practice to put the `<script>` element at the end of the body.

While JavaScript most commonly is used to manipulate the web page content and is thus pretty obvious to the user, it *also* can produce “terminal-like” output—including printed text and **error messages**. This output can be viewed through the **JavaScript Console** included as a *developer tool* in the Chrome browser (other browsers include similar tools):




Figure 10.1: Accessing the developer console in Chrome.

Important: You should *always* have the JavaScript Console open when developing JavaScript code, since this is where any error messages will appear!

On the Command-line with Node.js

It is also possible to execute JavaScript code on the command-line using Node.js. Node is a **command-line runtime environment**—that is, a JavaScript inter-

preter that can be run on the command-line.

With Node installed on your machine, you can use the **node** terminal command to start an *interactive Node session*. This will allow you to type JavaScript code directly in the terminal, and your computer will interpret and execute each line of code:




Figure 10.2: An interactive Node session.

You can enter one line of code at a time at the prompt (>). This is a nice way to experiment with the JavaScript language or to quickly run and test some code.

- You can exit this session by typing the `quit()` command, or hitting `ctrl-z` (followed by Enter on Windows).
- Note that the JavaScript console in the Chrome Developer tools provides the same interactive functionality.

It is also possible to run entire scripts (`.js` files) from the command-line by using the Node command and specifying the script file you wish to execute:

```
node my-script.js
```

This allows you author entire programs in JavaScript (e.g., writing the code using VS Code), and then executing them from the command-line. This process is more common when doing *server-side* web development (such as implementing a web server using Node), but we will still use it for practice and testing.

10.3 Writing Scripts

Unlike Java or C, JavaScript has no `main()` method that “starts” the program.

Instead, a JavaScript script (a `.js` file) contains a sequence of **statements** (instructions) that the interpreter executes *in order, one at a time, top to bottom*—just as if you had typed them into an interactive session one after another. Each line can declare a variable or function, which will then be available to the next statement to use or call. Thus with a JavaScript program, you should

think of the *sequence* of steps you want to be performed, and write lines of code in that order.

- In a way, you can think of the entire file as the “body” of a `main()` method... except that this body will contain further function declarations.

For example:

```
/* script.js */  
/* This is the ENTIRE contents of the file! */  
console.log("Hello world!"); //this is executed first  
console.log("I'm doing JavaScript!"); //this is executed second
```

The above example code demonstrates the `console.log()` function, which is JavaScript’s equivalent to Java’s `System.out.println()`—the output will be shown in the JavaScript console (in the Developer Tools). Thus we talk about “logging out” values in JavaScript, instead of “printing” values in Java.

- The `console.log()` function is technically a `log()` method belonging being called on a *global* `console` object. Globals and objects will be discussed in more detail below.



Figure 10.3: Don’t forget the semicolons!

As in Java, each statement should end with a semicolon (`;`) marking the end of that statement. But unlike Java, JavaScript can be forgiving if you forget a semicolon. The JavaScript interpreter tries to be “helpful” and will often assume that statements end at the end of a line if the next line “looks like” a new statement. However, it occasionally screws up in horrible ways—and so best practice as a developer is to **always include the semicolons**.

Strict Mode

ES5 includes the ability for JavaScript to be interpreted in **strict mode**. Strict mode is more “strict” about how the interpreter understands the syntax: it is less likely to assume that certain programmer mistakes were intentional (and so try to run the code anyway). For example, in *strict mode* the interpreter will produce an *Error* if you try and use a variable that has not yet been defined, while without strict mode the code will just use an *undefined* value. Thus working in strict mode can help catch a lot of silly mistakes.

You declare that a script or function should be executed in strict mode by putting an *interpreter declaration* at the top:

```
'use strict';
```

- This is not a String (or even JavaScript code!), but rather a *declaration* to the interpreter about how that interpreter should behave.

ALWAYS USE STRICT MODE! It will help avoid typo-based bugs, as well as enable your code to run more efficiently.

10.4 Variables

Variables in JavaScript are **dynamically typed**. This means that variables are not *not* declared as a particular type (e.g., `int` or `String`), but instead take on the data type (`Number`, `String`, etc.) of the *value* currently assigned to that variable. As we don’t specify the data type, JavaScript variables are declared using the `let` keyword:

```
let message = 'Hello World'; //a String
console.log(typeof message); //=> 'string'

let shoeSize = 7; //a Number
console.log(typeof shoeSize); //=> 'number'
```

- The `typeof` operator will return the data type of a variable. It is not widely used outside of debugging.
- As in Java, JavaScript variables should be given descriptive names using camelCase.

****Pro Tip:**** Even though variables in JavaScript loosely typed, the data type of a value is still important! In order to help keep track of the type of each variable in JavaScript, include the type in the variable name. For example: `textString`, `wordsArray`, `totalNum`, `itemStr`, etc..

Declared variables have a default value of *undefined*—a value representing that the variable has no value. This is somewhat similar to `null` in Java (though

JavaScript *also* as a null value that is not commonly used):

```
//create a variable (not assigned)
let hoursSlept;
console.log(hoursSlept); //=> undefined
```

Note that the `let` keyword is in fact new syntax introduced with ES6: older versions of JavaScript used the `var` keyword instead (and you will see `var` in most existing examples and tutorials). The difference is that variables declared with `let` are “**block scoped**”, meaning they are only available within the *block* (the `{}`) in which they are defined. This is the same way variables are scoped in Java. Variables declared with `var`, on the other hand, are “functionally scoped” so are available anywhere within the *function* in which they are defined. This means that you could declare a variable within an `if` block, and that variable would continue to be available outside that block! Thus `let` allows for cleaner, more efficient code, and with less bugs.

- `let` is the one ES6 feature supported by IE 11, meaning it can be used with most all current browsers. However, if you do need to support an older browser (e.g., IE 10, Safari 9.3, Android 4.4), you should transpile your code or stick to using `var`

Along with `let`, JavaScript variables can also be declared using the `const` keyword to indicate that they are *constant* (similar to what the `final` keyword does in Java). A `const` variable is block scoped, but can only be assigned once:

```
const ISCHOOL_URL = 'https://ischool.uw.edu'; //declare constant
ISCHOOL_URL = 'https://example.com'; //TypeError: Assignment to constant variable.
```

Basic Data Types

JavaScript supports many of the same basic data types as Java and other languages:

- **Numbers** are used to represent numeric data (JavaScript does not distinguish between integers and floats). Numbers support the same *mathematical* and operators as Java. Common mathematical functions can be accessed through in the built-in `Math` global (similar to Java’s `Math` class).

```
let x = 5;
typeof x; //'number'
let y = x/4;
typeof y; //'number'

//numbers use floating point division
console.log( x/4 ); //1.25

//use the Math.floor() function to do integer division
```

```
console.log( Math.floor(x/4) ); //1

//other common Math functions available as well
console.log( Math.sqrt(x) ); //2.23606797749979
```

- **Strings** can be written in either single quotes (') or double quotes ("), but most style guidelines recommend single quotes—just be consistent! Strings can be concatenated as in Java:

```
let name = 'Joel';
let greeting = 'Hello, my name is '+name; //concatenation
```

Strings also support many methods for working with them. As with Java, JavaScript strings are *immutable*, so most string methods return a new, altered string.

```
let message = 'Hello World';
let shouted = message.toUpperCase();
console.log(shouted); //=> 'HELLO WORLD'
```

- **Booleans** (true and false) in JavaScript work the same way as in Java, can be produced using the same *relational operators* (e.g., <, >=, !=), and support the same *logical operators* (e.g., &&, ||, and !):

```
//conjunction (and)
boolOne && boolTwo

//disjunction (or)
boolOne || boolTwo

//negation (not)
!boolOne //not
```

See Type Coercion and Control Structures below for more on working with Booleans in JavaScript.

Arrays

JavaScript also supports **arrays**, which are *ordered, one-dimensional sequences of values*. As in Java, JavaScript Arrays are written as literals inside square brackets []. Individual elements can be accessed by (0-based) *index* using **bracket notation**.

```
//an array of names
let names = ['John', 'Paul', 'George', 'Ringo'];

//an array of numbers (can contain "duplicate" values)
```

```

let numbers = [1, 2, 2, 3, 5, 8];

//arrays can contain different types (including other arrays!)
let things = ['raindrops', 2.5, true, [3, 4, 3]];

//arrays can be empty (contain no elements)
let empty = [];

//access using bracket notation
console.log( names[1] ); // "Paul"
console.log( things[3][2] ); // 3

numbers[0] = '343'; //assign new value at index 0
console.log( numbers ); // [343, 2, 2, 3, 5, 8]

```

Note that it is possible to assign a value to *any* index in the array, even if that index is “out of bounds”. This will *grow* the array (increase its length) to include that index—intermediate indices will be given values of *undefined*. The *length* of the array (accessed via the *.length* attribute) will always be the index of the “last” element + 1, even if there are fewer defined values within the array.

```

let letters = ['a', 'b', 'c'];
console.log(letters.length); // 3
letters[5] = 'f'; //grows the array
console.log(letters); // [ 'a', 'b', 'c', , , 'f' ]
                      //blank spaces are undefined
console.log(letters.length); // 6

```

Arrays also support a variety of methods that can be used to easily modify their elements, similar to the *ArrayList* class in Java:

```

//Make a new array
let array = ['i','n','f','o'];

//add item to end of the array
array.push('343');
console.log(array); //=> ['i','n','f','o','343']

//combine elements into a string
let str = array.join('-');
console.log(str); //=> "i-n-f-o-343"

//get index of an element (first occurrence)
let oIndex = array.indexOf('o'); //=> 3

//remove 1 element starting at oIndex
array.splice(oIndex, 1);

```



```
console.log(array); //=> ['i','n','f','343']
```

Objects

The most generic and useful data type in JavaScript is the Object data type. An **Object** is a lot like an array in that it is in that it is a (one-dimensional) sequence of values that are all stored in a single variable. However, rather than using *integers* as the index for each element, an Object uses *Strings*. Thus Objects are *unordered* sequences of **key-value pairs**, where the keys (called “**properties**”) are arbitrary Strings and the values are any data type—each property can be used to *look up* (reference) the value associated with it.

- This is a lot like a real-world dictionary or encyclopedia, in which the words (keys) are used to look up the definitions (values). A phone book works the same way (the names are the keys, the phone numbers are the values). A JavaScript Object provides a *mapping* from properties to values.
- JavaScript Objects are similar data structures to Java *HashMaps*, Python *dictionaries*, R *lists*, or even an HTML `<dl>` descriptive list! More generally, these are known as *associative arrays* (they are arrays that “associate” a key and a value). As with other associative arrays, JavaScript Objects are implemented as hash tables, making data access very fast.

Objects are written as literals inside curly braces `{}`. Property-value pairs are written with a *colon* (`:`) between the property name and the value, and each element (pair) in the Object is separated by a *comma* (`,`). Note that the property names do *not* need to be written in quotes if they are a single word (the quotes are implied—properties are always Strings):

```
//an object of ages (explicit Strings for keys)
//The `ages` object has a `sarah` property (with a value of 42)
let ages = {'sarah':42, 'amit':35, 'zhang':13};

//different properties can have the same values
//property names with non-letter characters must be in quotes
let meals = {breakfast:'coffee', lunch: 'coffee', 'afternoon tea': 'coffee'}

//values can be of different types (including arrays or other objects!)
let typeExamples = {number:12, string:'dog', array:[1,2,3]};

//objects can be empty (contains no properties)
let empty = {}
```

Important note: Objects are an **unordered** collection of key-value pairs! Because you reference a value by its *property name* and not by its position (as you

do in an array), the exact ordering of those elements doesn't matter—the interpreter just goes immediately to the value associated with the property. This almost means that when you log out an Object, the order in which the properties are printed may not match the order in which you specified them in the literal.

- If you `console.log()` an object, you will usually see a nicely format version of that object. However, if you try to convert that object to a *string* (e.g., via concatenation: `let output = "my object: "+myObject`), you will instead be presented with the “string version” of that object: `[object Object]`. This is not an array, but the String version of an object (similar to the hash you get when you try and print a Java array directly). You can instead use the fact that `console.log` accepts *multiple parameters* to output an object with a leading string:

```
let myObject = {a:1, b:2}

//convert object to string, won't log nicely
console.log("My object: " + myObject); //=> My object: [object Object]

//log the object directly
console.log("My object ", myObject); //=> My object {a: 1, b: 2}
```

Despite the name, JavaScript Objects have nothing to do with Object-Oriented Programming. However, they can be used to represent “things” in the same way as a Java object, with each property acting like an “attribute” (instance variable):

```
//an object representing a Person (spacing is for readability; white-space is ignored)
let person = {
  first_name: 'Alice',
  last_name: 'Smith',
  age: 40,
  pets: ['rover', 'fluffy', 'mittens'], //value is an array
  favorites: { //value is another object
    music: 'jazz',
    food: 'pizza',
    numbers: [12, 42] //value is an array
  }
}
```

Accessing Objects

Similar to arrays, Object values can be access via **bracket notation**, specifying the *property name* as the index. If an object does not explicitly have a property value, accessing that key produces `undefined` (the property's value is `undefined`).

```

let favorites = {music: 'jazz', food: 'pizza', numbers: [12, 42]};

//access variable
console.log( favorites['music'] ); //'jazz'

//assign variable
favorites['food'] = 'cake'; //property name is a string
console.log( favorites['food'] ); //'cake'

//access undefined key
console.log( favorites['language'] ); //undefined
favorites['language'] = 'javascript'; //assign new key and value

//access nested values
console.log( favorites['numbers'][0] ); //12

//use a variable as the "key"
let userInputTopic = 'food'; //pretend this value is supplied dynamically
console.log(favorites[userInputTopic]); //'cake'

```

Additionally, Object values can also be accessed via **dot notation**, as if the properties were *public attributes* of a class instance. This is often simpler to write and to read: remember to read the `.` as an `'s'`!

```

let favorites = {music: 'jazz', food: 'pizza', numbers: [12, 42]};

//access variable
console.log( favorites.music ); //'jazz'

//assign variable
favorites.food = 'cake';
console.log( favorites.food ); //'cake'

//access undefined key
console.log( favorites.language ); //undefined
favorites.language = 'javascript'; //assign new key and value

//access nested values
console.log( favorites.numbers[0] ); //12

```

- The only advantage to using *bracket notation* is that you can specify property names as variables or the results of an expression. But overall, the recommendation is to use *dot notation* unless the property you wish to access is dynamically determined.

It is possible to get an *array* of an object's keys calling the `Object.keys()` method and passing in the object you wish to get the keys of. Note that an

equivalent function for values is not supported by most browsers; a better approach is to iterate through the keys to identify all the values.

```
let ages = {sarah:42, amit:35, zhang:13};  
let keys = Object.keys(ages); // [ 'sarah', 'amit', 'zhang' ]
```

Arrays of Objects

As noted above, both arrays and objects can have values of any type—including other arrays or objects! The ability to nest objects inside of objects is incredibly powerful, and allows us to define arbitrarily complex information structurings (schemas). Indeed, most data in computer programs—as well as public information available on the web—is structured as a set of nested maps like this (though possibly with some level of abstraction).

One of the most common forms of nesting you'll see is to have an **array of objects** where each object has *the same properties* (but different values). For example:

```
//an arbitrary list of people's names, heights, and weights  
let people = [  
  {name: 'Ada', height: 58, 'weight': 115},  
  {name: 'Bob', height: 59, 'weight': 117},  
  {name: 'Chris', height: 60, 'weight': 120},  
  {name: 'Diya', height: 61, 'weight': 123},  
  {name: 'Emma', height: 62, 'weight': 126}  
]
```

This structure can be seen as a list of **records** (the objects), each of which have a number of different **features** (the key-value pairs). This list of feature records is in fact a common way of understanding a **data table** like you would create as an Excel spreadsheet:

name	height	weight
Ada	58	115
Bob	59	117
Chris	60	120
Diya	61	123
Emma	62	126

Figure 10.4: An data table.

Each object (record) acts as a “row” in the table, and each property (feature) acts as a “column”. As long as all of the objects share the same keys, this array of objects *is* a table!

Type Coercion

As mentioned above, variables in JavaScript are *dynamically typed*, and thus have a data type of the value **currently** assigned to that variable. JavaScript variables are able to “change type” by having a different type of data assignment to them:

```
let myVariable = 'hello'; //value is a String
myVariable = 42; //value is now a Number
```

Unlike some other dynamically typed variables, JavaScript will not throw errors if you try to apply operators (such as + or <) to different types. Instead, the interpreter will try to be “helpful” and **coerce** (convert) a value from one data type into another. While this process is similar to how Java will automatically cast data types into Strings (e.g., “hello”+4), JavaScript’s type coercion can produce a few quirks:

```
let x = '40' + 2;
console.log(x); //=> '402'; the 2 is coerced to a String
let y = '40' - 4;
console.log(y); //=> 36; can't subtract strings so '40' is coerced to a Number!
```

JavaScript will also attempt to coerce values when checking for equality with ==:

```
var num = 10
var str = '10'

console.log(num == str) //true, the values can be coerced into one another
```

In this case, the interpreter will coerce the Number 10 into the String '10' (since numbers can always be made into Strings), and since those Strings are the same, determines that the variables are equal.

In general this type of automatic coercion can lead to subtle bugs. Thus you should instead always use the === operator (and its partner !==), which checks both value *and* type for equality:

```
var num = 10
var str = '10'

console.log(num === str) //false, the values have different types
```

JavaScript will do its best to coerce any value when compared. Often this

means converting values to Strings, but it will also commonly convert values into *booleans* to compare them. So for example:

```
//compare an empty String to the number 0
console.log( '' == 0 ); //true; both can be coerced to a `false` value
```

This is because both the empty string `''` and `0` are considered “**falsey**” values (values that can be coerced to `false`). Other falsy values include `undefined`, `null`, and `NaN` (not a number). All other values will be coerced to `true`.

For more examples of the horror of JavaScript coercion, see this video (about 1:20 in).

10.5 Control Structures

JavaScript control structures have a similar syntax to those in Java or C. For example, a JavaScript **if statement** is written as:

```
if(condition){
    //statements
}
else if(alternativeCondition) {
    //statements
}
else {
    //statements
}
```

The **condition** can be any expression that evaluates to a Boolean value. But since any value can be *coerced* into Booleans, you can put any value you want inside the `if` condition. This is actually really useful—since `undefined` is a falsy value, you can use this coercion to check if a variable has been assigned or not:

```
//check if a `person` variable has a `name` property
if(person.name){
    console.log('Person does have a name!');
}
```

In the above example, the condition will only coerce to `true` if `person.name` is defined (*not undefined*) and is not empty. If somehow the variable has not been assigned (e.g., the user didn’t fill out the form completely), then the condition will not be true.

- While using `person.name !== undefined` is an equivalent expression, it is more idiomatic to utilize the coercion and have plain variable as the conditional expression.

Additionally, JavaScript provides a ternary conditional operator that lets you write a simple if statement as a single expression:

```
let x; //declare variable
if(condition) {
  x = 'foo';
} else {
  x = 'bar';
}

//can be condensed into:
let x = condition ? 'foo' : 'bar';
```

- This expression is read as “*if condition then* resolve to 'foo' *else* resolve to bar”. Each part (the condition, the “if true” result, and the “if false” result) can be any expression, though you should keep them simple to avoid confusion.

JavaScript also supports **while** loop (for indefinite iteration) and **for** loops (for definite iteration) similar to Java. The only difference is that because JavaScript variables are dynamically typed, the *loop control variables* are not declared with a type:

```
//an example for loop. The `i` is not declared as an int
for(let i=0; i<array.length; i++){
  console.log(array[i]);
}
```

JavaScript *does* have a `for ... in` syntax. However, it doesn't work as you would expect for arrays (it iterates over “enumerable properties” rather than the specific indices), and so should **not** be used with arrays. ES6 also introduces a `for ... of` syntax for iterating over arrays, but this is not supported by all browsers and so is not recommended. Instead, the current best practice is to use the above `for` loop, or better yet the `forEach()` method described in the next chapter.

- If you need to iterate over the keys of an object, use the `Object.keys()` method to get an array to loop through!

10.6 Functions

And of course, JavaScript includes **functions** (named sequences of statements used to *abstract* code). JavaScript functions are written using the following syntax:

```
//A function named `makeFullName` that takes two arguments
//and returns the "full name" made from them
```

```
function makeFullName(firstName, lastName) {  
    //Function body: perform tasks in here  
    let fullName = firstName + " " + lastName;  
  
    // Return: what you want the function to output  
    return fullName;  
}  
  
// Call the makeFullName function with the values "Alice" and "Kim"  
// Assign the result to `myName`  
let myName = makeFullName("Alice", "Kim") // "Alice Kim"
```

- Functions are defined by using the **function** keyword (placed before the name of the function) instead of Java's **public static**, don't declare a return type (since the language is dynamically typed), and don't indicate types for the parameters. Otherwise, JavaScript functions have identical syntax to Java functions.
- If a function lacks a **return** value, then that function returns the value **undefined**.

As in other languages, function arguments in JavaScript are implicitly declared *local* variables. However, in JavaScript **all arguments are optional**. Any argument that is *not* passed a specific value will be **undefined**. Any passed in value that does not have a variable declared for its position will not be assigned to a variable.

```
function sayHello(name) {  
    return "Hello, "+name;  
}  
  
//expected: argument is assigned a value  
sayHello("Joel"); // "Hello, Joel"  
  
//argument not assigned a value (left undefined)  
sayHello(); // "Hello, undefined"  
  
//extra arguments (values) are not assigned to variables, so are ignored  
sayHello("Joel", "y'all"); // "Hello, Joel"
```

- If a function has an argument, that doesn't mean it got a value. If a function lacks an argument, that doesn't mean it wasn't given a value!

In addition to this basic structure, JavaScript functions are frequently used for *functional programming*, as described in the next chapter.

Resources

As the language used for web programming, JavaScript may have more freely available online learning resources than any other language! Some of my favorites include:

- [A Re-Introduction to JavaScript](#) a focused tutorial on the primary language features
- [You Don't Know JS](#) a free textbook covering all aspects of the JavaScript language. Very readable and thorough, with lots of good examples.
- [JavaScript for Cats](#) a gentler introduction for “Scaredy-Cats”
- [MDN JavaScript Reference](#) a complete documentation of JavaScript, including tutorials
- [w3Schools JavaScript Reference](#) a slightly more friendly reference for the language
- [Google's JavaScript Style Guide](#)

Chapter 11

Functional Programming in JS

Despite its name, the JavaScript language was based more on Scheme than it was on Java. Scheme is a **functional programming language**, which means it follows a programming paradigm centered on *functions* rather than on *variables*, objects, and statements as you’ve done before (known as *imperative programming*). An alternative to object-oriented programming, functional programming provides a framework for thinking about how to give instructions to a computer. While JavaScript is not a fully functional language, it does support a number of functional programming features that are vital to developing effective and interactive systems. This chapter introduces these functional concepts.

11.1 Functions ARE Variables

Normally you’ve considered functions as “named sequences of instructions”, or groupings of lines of code that are given a name. But in a functional programming paradigm, functions are first-class citizens—that is, they are “things” (values) that can be organized and manipulated *just like variables*.

In JavaScript, **functions ARE variables**:

```
//create a function called `sayHello`  
function sayHello(name) {  
    console.log("Hello, "+name);  
}  
  
//what kind of thing is `sayHello` ?  
console.log(typeof sayHello); //=> 'function'
```

Just like `let x = 3` defines a variable for a value of type `number`, or `let msg = "hello"` defines a variable for a value of type `string`, the above `sayHello` function is actually a variable for a *value* of type `function`!

Important: we refer to a function by its name *without* the parentheses!

The fact that functions **are** variables is the core realization to make when programming in a functional style. You need to be able to think about functions as **things** (nouns), rather than as **behaviors** (verbs). If you imagine that functions are “recipes”, then you need to think about them as *pages from the cookbook* (that can be bound together or handed to a friend), rather than just the sequence of actions that they tell you to perform.

And because functions are just another type of variable, they can be used **anywhere** that a “regular” variable can be used. For example, functions are values, so they can be assigned to other variables!

```
//create a function called `sayHello`  
function sayHello(name) {  
    console.log("Hello, "+name);  
}  
  
//assign the `sayHello` value to a new variable `greet`  
let greet = sayHello;  
  
//call the function assigned to the `greet` variable  
greet("world"); //logs "Hello world"
```

- It helps to think of functions as just a special kind of array. Just as *arrays* have a special syntax `[]` (bracket notation) that can be used to “get” a value from the list, *functions* have a special syntax `()` (parentheses) that can be used to “run” the function.

Anonymous Functions

Functions are values, just like arrays and objects. And just as arrays and objects can be written as literals which can be *anonymously* passed into functions, JavaScript supports **anonymous functions**:

```
var array = [1,2,3]; //named variable (not anonymous)  
console.log(array); //pass in named var  
console.log( [4,5,6] ); //pass in anonymous value  
  
//named function (normal)  
function sayHello(person){  
    console.log("Hello, "+person);  
}
```

```
//an anonymous function (with no name!)
//(We can't reference this without a name, so writing an anonymous function is
//not a valid statement)
function(person) {
    console.log("Hello, "+person);
}

//anonymous function (value) assigned to variable
//equivalent to the version in the previous example
let sayHello = function(person) {
    console.log("Hello, "+person);
}
```

- You can think of this structure as equivalent to declaring and assigning an array `let myVar = [1,2,3]...` just in this case instead of taking the anonymous array (right-hand side) and giving it a name, we're taking an *anonymous function* and giving it a name!

Thus you can define named functions in one of two ways: either by making an explicitly named function or by assigning an anonymous function to a variable:

```
//these produce the same function
function foo(bar) {}
let foo = function(bar) {}
```

The only difference between these two constructions is one of *ordering*. When the JavaScript interpreter is initially reading and parsing the script file, it will put variable and function declarations into memory before it executes any of the file. In effect, JavaScript will seem to “move” variable and function declarations to the top of the file! This process is called **hoisting** (declarations are “hoisted” to the top of the script). Hoisting *only* works for named function declarations: assigning an anonymous function to a variable will not hoist that function’s definition:

In practice, you should **always** declare and define functions before you use them (put them all at the top of the file!), which will reduce the impact of hoisting and allow you to use either construction.

11.2 Object Functions

Moreover, functions are values, so they can be assigned as values of *object properties* (since object properties are like name-spaced variables):

```
//an object representing a dog
let dog = {
    name: 'Sparky'
```

```
    breed: 'mutt'
  }

  //assign an anonymous function to the `bark` property
  dog.bark = function(){
    console.log('woof!');
  }

  //call the function
  dog.bark(); //logs "woof!"
```

- Again, this is just like how you can assign an array as an object's property. With an array value you would use bracket notation to use it's "special power"; with a function value you use parentheses!

This is how we can create an equivalent of "member functions" (or methods) for individual objects: the `dog` object now has a function `bark()`!

- Similar to Java, you can refer to the *object on which a function is called* using the keyword `this`. Note that the manner in which the `this` variable is assigned can lead to some subtle errors when using callback functions (below). For more details, see the chapter on ES6 features.

11.3 Callback Functions

Finally, functions are values, so they can be *passed as parameters to other functions*!

```
//create a function `sayHello`
function sayHello(name){
  console.log("Hello, "+name);
}

//a function that takes ANOTHER FUNCTION as an argument
//this function will call the argument function, passing it "world"
function doWithWorld(funcToCall){
  //call the given function with an argument of "world"
  funcToCall("world");
}

doWithWorld(sayHello); //logs "Hello world";
```

In this case, the `doWithWorld` function will *execute* whatever function it is given, passing in a value of "world". (You can kind of think of this as similar to having a function that accesses the 'world' property of a given object).

- **Important note:** when we pass `sayHello` as an argument, we don't put any parentheses after it! Putting the parentheses after the function name *executes* the function, causing it to perform the lines of code it defines. This will cause the expression containing the function to *resolve* to its returned value, rather than being the function value itself. It's like passing in the baked cake rather than the recipe page.

```
function greet() { //version with no args for clarity
  return "Hello";
}

//log out the function value itself
console.log(greet); //logs e.g., [Function: greet], the function

console.log(greet()); //logs "Hello", which is what `sayHello()` resolves to
```

A function that is passed into another is commonly referred to as a **callback function**: it is an argument that the other function will “call back to” and execute when needed.

```
function doLater(callback) {
  console.log("I'm waiting a bit...");
  console.log("Okay, time to work!");
  callback(); // "call back" and execute that function
}

function doHomework() {
  // ...
};

doLater(doHomework);
```

Functions can take more than one callback function as arguments, which can be a useful way of *composing* behaviors.

```
function doTogether(firstCallback, secondCallback){
  firstCallback(); //execute the first function
  secondCallback(); //execute the second function
  console.log('at the same time!');
}

function patHead() {
  console.log('pat your head');
}

function rubBelly() {
  console.log('rub your belly');
```

```
}

//pass in the callbacks to do them together
doTogether(patHead, rubBelly);
```

This idea of *passing functions as arguments to other functions* is at the heart of functional programming, and is what gives it expressive power: we can define program behavior primarily in terms of the behaviors that are run, and less in terms of the data variables used. Moreover, callback functions are vital for supporting **interactivity**: many built-in JavaScript functions take in a callback function that specifies what should occur at some specific time (e.g., when the user clicks a button).

Often a callback function will be defined just to be passed into a single other function. This makes naming the callback somewhat redundant, and so it is more common to utilize **anonymous callback functions**:

```
//name anonymous function by assigning to variable
let sayHello = function(name){
    console.log("Hello, "+name);
}

function doWithWorld(funcToCall){
    funcToCall("world");
}

//pass the named function by name
doWithWorld(sayHello);

//pass in anonymous version of the function
doWithWorld(function(name){
    console.log("Hello, "+name);
});
```

- In a way, we’ve just “copy-and-pasted” the anonymous value (which happens to be a function) into the `doWithWorld()` call—just as you would do with any other anonymous variable type.
- Look carefully at the location of the closing brace `}` and parenthesis `)` on the last line. The brace ends the definition of the anonymous function value (the first and only parameter to `doWithWorld`), and the parenthesis ends *the parameter list* of the `doWithWorld` function. You need to include both for the syntax to be valid!
 - And since anonymous functions can be defined within other anonymous functions, it is not unusual to have lots of `})` lines in your code.

Closures

Functions are values, so not only can then be passed as parameters to other functions, they can also be *returned as results* of other functions!

```
//This function produces ANOTHER FUNCTION
//which greets a person with a given greeting
function makeGreeterFunc(greeting){
  //explicitly store the param as a local variable (for clarity)
  let localGreeting = greeting;

  //A new function that uses the `greeting` param
  //this is just a value!
  let aGreeterFunc = function(name){
    console.log(localGreeting+" "+name);
  }

  return aGreeterFunc; //return the value (which happens to be a function)
}

//Use the "maker" to create two new functions
let sayHello = makeGreeterFunc('Hello'); //says 'Hello' to a name
let sayHowdy = makeGreeterFunc('Howdy'); //says 'Howdy' to a name

//call the functions that were made
sayHello('world'); //"Hello world"
sayHello('Dave'); //"Hello Dave"
sayHowdy('world'); //"Howdy wold"
sayHowdy('partner'); //"Howdy partner"
```

- In this example, we’ve defined a function `makeGreeterFunc` that takes in some information (a greeting) as a parameter. It uses that information to create a new function `aGreeterFunc`—this function will have different behavior depending on the parameter (e.g., it can say “Hello” or “Howdy” or any other greeting given). We then return this new `aGreeterFunc` so that it can be used later (outside of the “maker” function).

When we then call the `makeGreeterFunc()`, the result (a function) is assigned to a variable (e.g., `sayHello`). And because that result is a function, we can call it with a parameter! Thus `makeGreeterFunc` acts a bit like a “factory” for making other functions, which can then be used where needed.

The most significant part of this example is the *scoping* of the `greeting` variable (and its `localGreeting` alias). Normally, you would think about `localGreeting` as being scoped to `makeGreeterFunc`—once the maker function is finished, then the `localGreeting` variable should be lost. However,

the `greeting` variable was in scope when the `aGreeterFunc` was created, and thus *remains in scope* (available) for that `aGreeterFunc` even after the maker function has returned!

This structure in which a function “remembers” its context (the in-scope variables around it) is called a **closure**. Even though the `localGreeting` variable was scoped outside of the `aGreeterFunc`, it has been *enclosed* by that function so it continues to be available later. Closures are one of the most powerful yet confusing techniques in JavaScript, and are a highly effective way of saving data in variables (instead of relying on global variables or other poor programming styles). They will also be useful when dealing with some problems introduced by Asynchronous Programming

11.4 Functional Looping

Another way that functional programming and callback functions specifically are utilized is to *replace loops with function calls*. For a number of common looping patterns, this can make the code more *expressive*—more clearly indicative of what it is doing and thus easier to understand. Functional looping was introduced in ES5.

To understand functional looping, first consider the common `for` loop used to iterate through an array of objects:

```
let array = [{...}, {...}, {...}];

for(let i=0; i<array.length; i++){
    let currentItem = array[i]; //convenience variable for current item

    //do something with current item
    console.log(currentItem);
}
```

While this loop may be familiar and fast, it does require extra work to manage the loop control variable (the `i`), which can get especially confusing when dealing with nested loops (and nested data structures are very common in JavaScript!)

As an alternative, you can consider using the Array type’s `forEach()` method:

```
let array = [{...}, {...}, {...}];

//function for what to do with each item
function printItem(currentItem){
    console.log(currentItem);
}
```

```
//print out each item  
array.forEach(printItem);
```

The `forEach()` method goes through each item in the array and executes the given *callback function*, passing that item as a parameter to the callback. In effect, it lets you specify “what to do with each element” in the array as a separate function, and then “apply” that function to each elements.

- `forEach()` is a *built-in* method for Arrays—similar to `push()` or `indexOf()`. For reference, the “implementation” of the `forEach()` function looks something like:

```
Array.forEach = function(callback) { //define the Array's forEach method  
    for(let i=0; i<this.length; i++) {  
        callback(this[i], i, this);  
    }  
}
```

In effect, the method does the job of managing the loop and the loop control variable for you, allowing you to just focus on what you want to do for each item.

- The *callback function* give to the `forEach()` will be executed with up to three argument (in order): (1) the *current item* in the array, (2) the *index* of the item in the array, and (3) the *array itself*. This means that your callback can contain up to three arguments, but since all arguments are optional in JavaScript, it can also be used with fewer—you don’t need to include an argument for the index or array if you aren’t utilizing them!

While it is possible to make a named callback function for `forEach()`, it is *much* more common to use an **anonymous callback function**:

```
//print each item in the array  
array.forEach(function(item){  
    console.log(item);  
})
```

- This code can almost be read as: “take the *array* and *forEach* thing execute the *function* on that *item*”.
- This is similar in usage to the enhanced for loop in Java.

Map

JavaScript provides a number of other functional loop methods. For example, consider the following “regular” loop:

```
function square(n) { //a function that squares a number
    return n*n;
}

let numbers = [1,2,3,4,5]; //an initial array

let squares = []; //the transformed array
for(let i=0; i<numbers.length; i++){
    let transformed = square(numbers[i]); //call our square() function
    squares.push(transformed); //add transformed to the list
}
console.log(squares); // [1, 4, 9, 16, 25]
```

This loop represents a **mapping** operation: it takes an original array (e.g., of numbers 1 to 5) and produces a *new* array with each of the original elements transformed in a certain way (e.g., squared). This is a common operation to apply: maybe you want to “transform” an array so that all the values are rounded or lowercase, or you want to *map* an array of words to an array of their lengths, or you want to *map* an array of values to an array of `` HTML strings. It is possible to make all these changes using the above code pattern: create a new empty array, then loop through the original array and push the transformed values onto that new array.

However, JavaScript also provides a *built-in array method* called **map()** that directly performs this kind of mapping operation on an array without needing to use a loop:

```
function square(n) { //a function that squares a number
    return n*n;
}

let numbers = [1,2,3,4,5]; //an initial array

//map the numbers using the `square` transforming function
let squares = numbers.map(square);

console.log(squares); // [1, 4, 9, 16, 25]
```

The array’s `map()` function produces a **new** array with each of the elements transformed. The `map()` function takes as an argument *a callback function* that will do the transformation. The callback function should as an argument the element to transform, and **return** a value (the transformed element).

- Callback functions for `map()` will be passed the same three arguments as the callback functions for `forEach()`: the element, the index, and the array.

And again, the `map()` callback function (e.g., `square()` in the above example)

is more commonly written as an anonymous callback function:

```
let numbers = [1,2,3,4,5]; //an initial array
let squares = numbers.map(function(item){
    return n*n;
});
```

Filter

A second common operation is to **filter** a list of elements, removing elements that we don't want (or more accurately: only keeping elements that we *DO* want). For example, consider the following loop:

```
function isEven(n) { //a function that determines if a number is even
    let remainder = n % 2; //get remainder when dividing by 2 (modulo operator)
    return remainder == 0; //true if no remainder, false otherwise
}

let numbers = [2,7,1,8,3]; //an initial array

let evens = []; //the filtered array
for(let i=0; i<numbers.length; i++){
    if(isEven(numbers[i])){
        evens.push(numbers[i]);
    }
}
console.log(evens); //[2, 8]
```

With this **filtering** loop, we are *keeping* the values for which the `isEven()` function returns `true` (the function determines “what to let in” not “what to keep out”; a whitelist), which we do by appending the “good” values to a new array.

Similar to `map()`, JavaScript arrays include a *built-in method* called **`filter()`** that will directly perform this filtering:

```
function isEven(n) { //a function that determines if a number is even
    return (n % 2) == 0; //true if no remainder, false otherwise
}

let numbers = [2,7,1,8,3]; //an initial array

let evens = numbers.filter(isEven); //the filtered array

console.log(evens); //[2, 8]
```

The array's `filter()` function produces a **new** array that contains only the

elements that *do match* a specific criteria. The `filter()` function takes as an argument a *callback function* that will make this decision. The callback function takes in the same arguments as `forEach()` and `map()`, and should return `true` if the given element should be included in the filtered array (or `false` if it should not).

- And again, we usually use *anonymous callback functions* for `filter()`:

```
let numbers = [2,7,1,8,3]; //an initial array
let evens = numbers.filter(function(n) { return (n%2)==0; }); //one-liner!
```

(Since JavaScript ignores whitespace, we can compact simple callbacks onto a single line. ES6 and Beyond describes an even more compact syntax for such functions).

Because `map()` and `filter()` are both called on and produce arrays, it is possible **chain** them together, calling subsequent methods on each returned value:

```
let numbers = [1,2,3,4,5]; //an initial array

//get the squares of EVEN numbers only
let filtered = numbers.filter(isEven);
let squares = filtered.map(square);
console.log(squares); //[4, 16, 36]

//or in one statement, using results anonymously
let squares = numbers.filter(isEven)
                      .map(square);
console.log(squares); //[4, 16, 36]
```

This structure can potentially make it easier to understand the code's intent than using a set of nested loops or conditionals: we are taking `numbers` and then *filtering* for the evens and *mapping* to squares!

Reduce

The third important operation in functional programming (besides *mapping* and *filtering*) is **reducing** an array. Reducing an array means to *aggregate* that array's values together, transforming the array elements into a single value. For example, summing an array is a *reducing* operation (and in fact, the most common one!): it reduces an array of numbers to a single summed value.

- You can think of `reduce()` as a *generalization* of the `sum()` function found in many other languages—but rather than just adding (+) the values together, `reduce()` allows you to specify what operation to perform when aggregating (e.g., multiplication).

To understand how a *reduce* operation works, consider the following basic loop:

```
function add(x, y) { //a function that adds two numbers
  return x+y;
}

let numbers = [1,2,3,4,5]; //an initial array

let runningTotal = 0; //an accumulated aggregate
for(let i=0; i<numbers.length; i++){
  runningTotal = add(runningTotal, numbers[i]);
}
console.log(runningTotal); //15
```

This loop **reduces** the array into an “accumulated” sum of all the numbers in the list. Inside the loop, the `add()` function is called and passed the “current total” and the “new value” to be combined into the aggregate (*in that order*). The resulting total is then reassigned as the “current total” for the next iteration.

The *built-in array method* **`reduce()`** does exactly this work: it takes as an argument *a callback function* used to combine the current running total with the new value, and returns the aggregated total. Whereas the `map()` and `filter()` callback functions each usually took 1 argument (with 2 others optional), the `reduce()` callback function requires **2** arguments (with 2 others optional): the first will be the “running total” (called the **accumulator**), and the second will be the “new value” to mix into the aggregate. (While this ordering doesn’t influence the summation example, it is relevant for other operations):

```
function add(x, y) { //a function that adds two numbers
  return x+y;
}

let numbers = [1,2,3,4,5]; //an initial array

let sum = numbers.reduce(add);
console.log(sum); //15
```

The `reduce()` function (not the callback, but `reduce()` itself) has a second optional argument *after* the callback function representing the initial starting value of the reduction. For example, if we wanted our summation function to start with 10 instead of 0, we’d use:

```
//sum starting from 10
let sum = numbers.reduce(add, 10);
```

Note that the syntax can be a little hard to parse if you use an anonymous callback function:

```
//sum starting from 10
numbers.reduce(function(x, y){
```

```

    return x+y;
  }, 10); //the starting value comes AFTER the callback!

```

The *accumulator* value can be any type you want! For example, you can have the starting value be an empty object `{}` instead of a number, and have the accumulator use the current value to “update” that object (which is “accumulating” information).

To summarize, the `map()`, `filter()`, and `reduce()` operations work as follows:

```

map([🐱, 🍌, 🐔, 🍌], cook)
=> [🍌, 🍌, 🍌, 🍌]

filter([🍌, 🍌, 🍌, 🍌], isVegetarian)
=> [🍌, 🍌]

reduce([🍌, 🍌, 🍌, 🍌], eat)
=> 🤔

```

Figure 11.1: Map, filter, reduce explained with emoji. Not valid JavaScript syntax.

All together, the **map**, **filter**, and **reduce** operations form the basic for a functional consideration of a program. Indeed, these kinds of operations are very common when discussing data manipulations: for example, the famous MapReduce model involves “mapping” each element through a complex function (on a different computer no less!), and then “reducing” the results into a single answer.

11.5 Pure Functions

This section was adapted a tutorial by Dave Stearns.

The concept of *first-class functions* (functions are values) is central to any functional programming language. However, there is more to the functional programming paradigm than just callback functions. In a fully functional programming language, you construct programs by combining small, reusable, **pure** functions

that take in some inputted data, transform it, and then return that data for future use. *Pure* functions have the following qualities:

- They operate only on their inputs, and make no reference to other data (e.g., variables at a higher scope such as *globals*)
- They never modify their inputs—instead, they always return new data or a reference to an unmodified input
- They have no *side effects* outside of their outputs (e.g., they never modify variables at a higher scope)
- Because of these previous rules, they always return the same outputs for the same inputs

A functional program sends its initial input state through a series of these pure functions, much like a plumbing system sends water through a series of pipes, filters, valves, splitters, heaters, coolers, and pumps. The output of the “final” function in this chain becomes the program’s output.

Pure functions are almost always easier to test and reason about. Since they have no side-effects, you can simply test all possible classes of inputs and verify that you get the correct outputs. If all of your pure functions are well-tested, you can then combine them together to create highly-predictable and reliable programs. Functional programs can also be easier to read and reason about because they end up looking highly *declarative*: it reads as a series of data transformations (e.g., take the data, then filter it, then transform it, then sort it, then print it), with the output of each function becoming the input to the next.

Although some functional programming zealots would argue that all programs should be written in a functional style, it’s better to think of functional programming as another tool in your toolbox that is appropriate for some jobs, and not so much for others. Object-oriented programming is often the better choice for long-running, highly-interactive client programs, while functional is a better choice for short-lived programs or systems that handle discrete transactions (like many web applications). It’s also possible to combine the two styles: for example, React components can be either object-oriented or functional, and you often use some functional techniques within object-oriented components.

- Indeed, the latest versions of Java—a highly object-oriented language—add support for functional programming!

If you are interested in doing more serious functional programming in JavaScript, there are a large number of additional libraries that can help support that:

- Lodash (and its more pure variant *lodash/fp*)
- Ramda
- Lazy.js

There are also many languages that were designed to be functional from the get-go, but can be compiled down into JavaScript to run on web browsers. These include Clojure (via ClojureScript) and Elm.

Resources

- Functional Programming in JavaScript a *fantastic* interactive tutorial for learning functional programming in JS
- Higher Order Functions a chapter from the online textbook *Eloquent JavaScript*
- Scope & Closures an online textbook with an extremely detailed explanation of scoping in JavaScript

Chapter 12

Document Object Model (DOM)

The primary purpose of using JavaScript in a web page is to make that page **interactive**: the JavaScript language is used to program logical decisions that will effect what is shown on the page. It does this primarily by *changing the HTML rendered by the browser*. For example, JavaScript can be used to change the text inside a `<p>`, add additional `` elements to a list, or to give a `<div>` a new CSS `class` attribute. The programmatic representation of the HTML elements currently being shown by the browser is known as the **Document Object Model (DOM)**. In web programming JavaScript code is used to modify the DOM (HTML elements) currently being shown by the browser in response to user input, thereby making the page interactive. This chapter introduces the Document Object Model and how to use JavaScript to manipulate it through user-driven interaction.

12.1 The DOM API

As you should recall from Chapter 3, HTML elements can be nested, allowing us to consider a webpage as a “**tree**” of elements:

- A **tree** is a hierarchical data structure, where each element (called a *node*) contains references to *child* elements. Following the arboreal metaphor, the “start” of the tree is called the *root node*, hierarchical sequences of nodes are called *branches*, and a *node* that does not have any children is called a *leaf*.

Considering a web page’s content to be a tree of HTML elements is one way to *model* (represent) the structure of that information. This particular model




Figure 12.1: An example DOM tree (a tree of HTML elements).

of a web *document* (as a tree of *object* nodes) is called the **Document Object Model**, or **DOM** for short. In many ways the DOM *is* the HTML (though the HTML rendered in the browser, not the `.html` source code you’ve written)! Thus we can refer to the web page’s content as “the DOM”, and an HTML element as a “DOM element”.

- Note that even “plain text” content (e.g., what is inside a `<p>` tag) are considered nodes in the DOM tree—they are “text content” nodes (instead of “element nodes”).

Moreover, the DOM also provides an **Application Programming Interface (API)** which allows computer *applications* to *programmatically* (e.g., through JavaScript code) *interact* with it: accessing and manipulating the tree of elements. As you may recall from previous courses, an API is often a set of *functions* and *variables* that can be used give instructions to a program. The DOM API is no different: it is a group of functions you can call and variables (usually Object properties) you can adjust to change the rendered web content. You write code to call these functions in order to make a page interactive.

Global Variables

You can programmatically access the API in JavaScript by utilizing a set of global variables. **Global variables** are variables that are “globally” scoped: they are available anywhere in the program (not just within a particular function).

An important programming style rule is to **minimize** the use of global variables. Try to avoid creating too many new globals yourself!

Global variables in JavaScript are almost always *Objects* that have methods as their values. For example, the JavaScript language itself provides a global **Math** object that has includes a number of function properties (e.g., `sqrt()`, `floor()`, etc.).

```
console.log( typeof Math ); //=> 'object'
console.log( typeof Math.sqrt ); //=> 'function'
console.log( Math.sqrt(25) ); //=> 5
```

- In fact, the `console` object is *another* global variable provided by the JavaScript runtime (whether inside the browser or inside Node.js)!

The *web browser* also provides a number of global variables that you can use. For example **window** is a global object that represents the browser itself. You can use this object to get information about the browser:

```
/* example properties */
let width = window.innerWidth; //viewport width
let height = window.innerHeight; //viewport height
```

```

var url = window.location.href; //url for this page

/* example functions */
window.alert("Boo!"); //show a popup alert. Do not use this.
window.scrollTo(0, 1000); //scroll to a position
window.setTimeout(callback, 1000); //execute callback after an delay
window.setInterval(callback, 1000); //execute callback repeatedly after interval

```

While these examples are included for completeness, most `window` functions are rarely used and should be avoided. Popups with the `window.alert()` function are inelegant, interrupt the user's actions, and produce a bad user experience—you should instead use in-window alerting options instead (such as showing a `<p class="alert">`). Browser control functions such as `scrollTo()` are non-standard and can vary drastically across systems and platforms. Proceed with caution when using `window` functions!

12.2 DOM Manipulation

While `window` represents the Browser, the **DOM** itself is represented by the **document** global object—`document` is the DOM (the *current* HTML rendered in the browser). You access properties and call methods of this object in order to manipulate the content displayed in the browser!

Referencing HTML Elements

In order to manipulate the DOM elements in a page, you first need to get a *reference* to the element you want to change—that is, you need a variable that refers to that element. You can get these variable references by using one of the `document` “selector” functions:

```

//element with id="foo"
let fooElem = document.getElementById('foo');

//elements with class="row"
let rowElems = document.getElementsByClassName('row'); //note the plural!

//<li> elements
let liElems = document.getElementsByTagName('li'); //note the plural!

/*easiest to select by reusing CSS selectors! */
let cssSelector = 'header p, .title > p'; //a string of a CSS selector
//selects FIRST element that matches css selector

```

```
let elem = document.querySelector(cssSelector);

//matches ALL elements that match css selector
let elems = document.querySelectorAll(cssSelector);
```

- The `document.querySelector()` is *by far* the most flexible and easy to use of these methods: it can easily do the same as all the other methods (just put in an id, class, or element selector). **You should always use `querySelector()`.**
- Note that the methods that return multiple nodes (e.g., `querySelectorAll`) return a `NodeList` object. While this is like an array (you can access elements via index through bracket notation and it has a `.length` property), it is **not** an array: meaning it doesn't support methods like `forEach()` and `map()` across all browsers. If you need to iterate through a `NodeList`, you should use a regular `for` loop. But in practice, you're much more likely to only work with single elements at a time.

Modifying HTML Elements

Once you have a reference to an element, you access properties and call methods on that object in order to modify its state in the DOM—which will in turn modify how it *currently* is displayed on the page. Thus by modifying these objects, you are dynamically changing the web page's content!

Important: setting these properties do not change the `.html` source code file! Instead, they just change the *rendered DOM elements* (think: the content stored in the computer's memory rather than in a file). If you refresh the page, the content will go back to how the `.html` source code file specifies it should appear—unless that also loads the script that modifies the DOM. What is shown on the page is the HTML with the JavaScript modifications added in.

Changing Content

You can use JavaScript to access and modify the **content** of a DOM element (e.g., the stuff between the start and close tags):

```
//get a reference to the FIRST <p> element
let elem = document.querySelector('p');

console.log(elem); //to demonstrate

let text = elem.textContent; //the text content of the elem
elem.textContent = "This is different content!"; //change the content
```

```
let html = elem.innerHTML; //content including HTML
elem.innerHTML = "This is <em>different</em> content!"; //interpreted as HTML
```

The `textContent` property of the element refers to *all* of the content, but considered as “plain text” this means that it is considered a “safe” property: you can assign strings that contain’s HTML code (e.g., `Hello`), but that code will be escaped and not interpreted as HTML (instead the `<` and `>` will be written out as if you had used HTML entities). The `.innerHTML` property, on the other hand, is “not safe”: any HTML included in the String you assign to it will be converted into DOM elements. This makes it not a great property to use unless you are absolutely certain the content came from a trusted source.

- The `innerHTML` property should be used primarily for including *inline* elements such as `` or ``. For more complex HTML content, it is cleaner code (separation of concerns!) to explicitly create new elements—see below for details.
- You can “clear” the content of an element by setting it’s content to be an empty string (`''`):

```
let alertElem = document.querySelector('.alert');
alertElem.textContent = ''; //no more alert!
```

Changing Attributes

You can also change the **attributes** of individual elements. Each attribute defined in the HTML specification is typically exposed as a *property* of the element object:

```
//get a reference to the `#picture` element
let imgElem = document.querySelector('#picture');

//access the attribute
console.log( imgElem.src ); //logs the source of the image

//modify the attribute
imgElem.src = 'my-picture.png';
```

You **cannot** access `element.class` or `element.style` attributes directly in this way; see below for specifics on changing the CSS of an element.

You can alternatively modify element attributes by using the methods `getAttribute()` (passing it which attribute to access) and `setAttribute()` (passing it which attribute to modify and what value to assign to that attribute):


```

let imgElem = document.querySelector('#picture');
imgElem.setAttribute('src', 'my-other-picture.png'); //set the src

console.log( imgElem.getAttribute('src') ); //=> 'my-other-picture.png'

//the `hasAttribute()` method returns a boolean.
let isThick = document.querySelector('svg rect')
                .hasAttribute('stroke-width'); //chained anonymous variables

```

These methods will let you interact with attributes that are *not* defined by the HTML spec, such as `data-` attribute. However, they *don't* work with certain element attributes (such as the `value` attribute of an `<input>` element). Other elements may have their own special DOM properties: see the DOM Documentation for a list of HTML interfaces.

12.2.0.1 Changing Element CSS

It is possible to modify the **CSS classes** (and even inline styling) of an element. But rather than using the `class` property like with other attributes, you instead access the **`className`** property. On modern browsers (IE 10 or later), this property supports methods `.add()` and `.remove()` for adding and removing classes from the list:

```

//access list of classes
let classList = elem.classList;

//add a class
elem.classList.add('small'); //add a single class
elem.classList.add('alert','alert-warning'); //add multiples classes (not on IE)

//remove a class
elem.classList.remove('small');

//"toggle" (add if missing, remove if present)
elem.classList.toggle('small');

```

- While IE 10+ does support these methods, it doesn't support *multiple arguments* for them (so you can't add multiple classes in a single method call). If you need to support older browsers (including any version of IE), you can instead modify the `.className` property as if it were a String:

```

//fallback for IE (all)
var classes = elem.className;
classes += ' ' + 'sweet sour'; //modify the string (append!)
elem.className = classes;      //reassign

```

The `classList` methods work perfectly on Microsoft Edge.

It is also possible to access and modify individual CSS properties of elements through the DOM element's `style` property. `.style` references an Object whose keys are the CSS property names (but written in *camelCase* instead of *kabob-case*)

```
let h1 = document.querySelector('h1');
h1.style.color = 'green';
h1.style.backgroundColor = 'black'; //not `.background-color`
```

In general, you should modify element CSS by changing the class of the element, rather than specific style properties.

Modifying the DOM Tree

In addition to modifying the individual DOM elements, it is also possible to access and modify the *DOM tree itself*! That is, you can create new elements and add them to the tree (read: webpage), remove elements from the tree, or pluck them out of the tree and insert them somewhere else!

First, note that each JavaScript DOM element has *read-only* properties referring to its parent, children, and sibling elements:

```
<main>
  <section id="first-section">
    <p>First paragraph</p>
    <p>Second paragraph</p>
  </section>
  <section id="second-section"></section>
</main>
```

```
//get reference to the first section
let firstSection = document.querySelector('#first-section');

//get reference to the "parent" node
let main = firstSection.parentElement;
console.log(main); //<main>...</main>

//get reference to the child elements (2 paragraphs)
let paragraphs = firstSection.children;
console.log(paragraphs.length); //2
console.log(paragraphs[0]); //<p>First paragraph</p>

//get reference to the the next sibling
let sectionSection = firstSection.nextElementSibling;
console.log(secondSection); //<section id="second-section"></section>
```

- Note that these properties only deal with *HTML elements*—text content nodes are ignored. You can instead use equivalent properties `parentNode` and `childNodes` to also consider text content nodes.

SVG content doesn't support `parentElement`, but does support `parentNode`.

You can also call methods to create and add new HTML DOM elements to the tree. The `document.createElement()` function is used to create a new HTML element. However this element is *not* created a part of the tree (after all, you haven't specified where it would put into the page)! Thus you need to also use a method such as `appendChild` to add that new element as a child of another element:

```
//create a new <p> (not yet in the tree)
let newP = document.createElement('p');
newP.textContent = "I'm new!";

//create Node of textContent only (not an HTML element, just text)
let newText = document.createTextNode("I'm blank");

let main = document.querySelector('main');
main.appendChild(newP); //add element INSIDE (at end)
main.appendChild(newText); //add the text inside, AFTER the <p>

//add anonymous new node BEFORE element. Parameters are: (new, old)
main.insertBefore(document.createTextNode("First!"), newP);

//replace node. Parameters are: (new, old)
main.replaceChild(document.createTextNode('boo'), newText);

//remove node
main.removeChild(main.querySelector('p'));
```

The `appendChild()` method is considered a cleaner approach than just modifying the `innerHTML` property, as it allows you to adjust the DOM tree without erasing what was previously there. A common practice is to use `document.createElement()` to create a *block* element, then set the `innerHTML` of that element to its content (which can include *inline* elements), and then use `appendChild` to add the new block element to the tree at the desired location.

Accessibility

Whenever you learn a new technology, you should ask: **how does this affect accessibility?** With the JavaScript code modifying the rendered DOM, it is possible that the content of a page will change *after* it has been read by a screen

reader. And while a sighted user will likely be able to see the change visually, a screen reader has no way of knowing that something on the page is different unless you tell it.

You can let screen readers know that an element in a page may have its content change *in the future* by making that element into an ARIA Live Region. Live regions are “watched” by assistive technologies, and whenever the content changes they will speak the new content to the reader as if it were being read for the first time.

You make an element into a live region by giving it the `aria-live` attribute:

```
<div aria-live="polite">
  This content can change!
</div>
```

The value assigned to the `aria-live` attribute is the “politeness level”, which specifies the priority by which the screen reader should read the change. The most common option (that you should almost always use) is “`polite`”, which indicates that the changed text will be read only once the user has *paused* whatever is currently being read. A “`polite`” alert doesn’t interrupt the currently being read text or description, but instead will be injected when there is a break (if the current reading goes on for too long, then the new content will not be spoken).

- The other option is “`assertive`”, which indicates that the new content should be spoken as soon as it changes, possibly interrupting other content. This should only be used for important information (like alerts, warnings, or errors), as it can interrupt the user’s flow in ways that are very disorienting. In short: *always be polite!*

12.3 Listening for Events

In order to make a page **interactive** (that is, able to change in response to user actions), you need to be able to respond to *user events*. Whenever a user interacts with a computer, the operating system announces that interaction as an **event**—the *event* of a button being clicked, the *event* of the mouse being moved, the *event* of a keyboard key being pressed, etc. These events are **broadcast** to the entire system, allowing any application (including the browser) to “respond” the occurrence of the event, such as by executing a particular JavaScript function.

Thus in order to respond to user actions (and the *events* those actions generate), we need to define a function that will be executed **when the event occurs**. You will define a function as normal, but the function will not get called by you as a particular step in your script. Instead, the function you specify will

be executed *by the system* when an event occurs, which will be at some indeterminate time in the future. This process is known as event-driven programming. It is also an example of **asynchronous programming**: in which statements are not executed in a single order one after another (“synchronously”), but may occur “out of order” or even at the same time! (For more about working with asynchronous programming, see Chapter 14).

In order for your script to respond to user events, you need to *register an event listener*. This is a bit like following someone on social media: you specify that you want to “listen” for updates from that person, as well as what you want to do when you “hear” some news from that person.

- Specifying that you want Slack to notify you when your name is mentioned is another good analogy!

The DOM API allows you to register an event listener by call the **.addEventListener()** on a selected element (e.g., on the element that you want to listen to). This method takes two arguments: a string representing what kind of event you want to listen for, and a *callback function* to execute when you hear that event:

```
//a (named) callback function
function onClickCallback() {
  console.log("You clicked me!");
}

//get a reference to the element we want to "listen" to
let button = document.querySelector('button');

//register a listener for 'click' events
button.addEventListener('click', onClickCallback);
```

- When the button is clicked by the user, it will “shout” out a ‘click’ event (“I was clicked! I was clicked!”). Because you have set up a listener (an alert/notification) for such an occurrence, your script will be able to do something—and that something that it will do is run the specified callback function.

It’s like you handed someone a recipe and told them “when I call you, bake this cake!”

- It is **much** more common to use an *anonymous function* as the callback:

```
let button = document.querySelector('button');
button.addEventListener('click', function() {
  console.log("You clicked me!");
});
```

- Note that this listener *only* applies to that particular button—if you wanted to respond to a different button, you’d need to register a sepa-

rate listener! Also, as the method name implies, it is possible to add multiple listeners (callbacks) to the same element for the same event: all of them will be executed “at once”.

The event callback will be passed in a single argument: an object representing the “*event*” that occurred. (Since all parameters are optional in JavaScript, and it wasn’t used in the above example, it wasn’t included in the callback definition). This event includes information such as where the event occurred (in x,y coordinates), what DOM element caused the event, and more:

```
elem.addEventListener('click', function(event) {  
    //get who was clicked;  
    let clickedElem = event.target; //target property of the event  
    console.log(clickedElem);  
});
```

Also note that sometimes you want to stop the “normal” results of an event from occurring. For example, perhaps you don’t want a button to do its normal button thing (such as submitting a form), and instead want to provide your own custom behavior. To support this, you can “interrupt the event” by calling the following methods on the event:

```
submitBtn.addEventListener('click', function(event) {  
  
    event.preventDefault(); //don't do normal behavior  
    event.stopPropagation(); //don't pass the event to parents  
  
    //..do custom behavior here  
  
    return false; //don't do normal behavior OR propagate! (for IE)  
})
```

Types of Events

There are numerous different events that you can listen for, including:

- Mouse Events such as **'click'**. The `event.offsetX` and `event.offsetY` will provide (x,y) coordinates for the clicks location *relative to the target element*; you can use `clientX/Y` for coordinates relative to the browser window, or `pageX/Y` for coordinates relative to the document (regardless of scrolling). See this post for details, and this page for an example.

Other mouse events include **'dblclick'** (double-click), **'mousedown'** (mouse button is pressed down, may be held), **'hover'** (mouse hover), **'mouseenter'** (mouse moves over element), **'mousemove'** (mouse moves over element), and **'mouseleave'** (mouse moves off of element).

- Keyboard Events such as **'keydown'**. The `event.key` property is used to determine *which key* was pressed, giving a predefined key value you can check:

```
elem.addEventListener('keydown', function(event){
    if(event.key === 'ArrowUp'){
        console.log("Going up!")
    }
    //...
});
```

The `event` object also has properties to check if any “modifier keys” such as shift, control, or meta (Windows/command) are held when the event occurs.

Note that you almost always want to respond to the `'keydown'` and `'keyup'` events; the `'keypressed'` event is sent later and only applies to non-modifier keys.

- Window Events are event created by the `window` global, which we are also able to register event listeners on! For example, the `'resize'` event can be used to identify when the window has changed size (e.g., if you want to make the content responsive as well as the CSS):

```
window.addEventListener("resize", function() {
    //...
});
```

(See the documentation for advise on using this callback)

Additionally, the `window` global defines a special event callback that occurs when the web page has finished loading. You can assign your own function to this callback to run code only *after* the webpage has loaded (e.g., for scripts specified in the `<head>`):

```
window.onload = function() {
    //...do stuff once page is ready (e.g., run the rest of your code)
}
```

Style guideline: always register event listeners in the JavaScript—do *not* utilize the HTML attributes such as `onClick`. This is to help keep concerns separated: the HTML should not need to know anything about the JavaScript that is utilized (since the browser may not even support JavaScript!), but it’s okay for the JavaScript to rely on and modify the HTML.

Event-Driven Programming

In a typical web program event callback functions can occur repeatedly, over and over again (e.g., every time the user clicks a button). This makes them

potentially act a bit like the body of a `while` loop. However, because these callbacks are *functions*, any variables defined within them are **scoped** to that function, and will not be available on subsequent executions. Thus if you want to keep track of some additional information (e.g., how many times the button was clicked), you will need to use a variable declared *outside* of the function (e.g., a **global**). Such variables can be used to represent the **state** (situation) of the program, which can then be used to determine what behavior to perform when an event occurs, following the below pattern:

```
//pseudocode
WHEN an event occurs {
    check the STATE of the program;
    DECIDE what to do based on that state;
    UPDATE the state as necessary for the next event;
}
```

For example:

```
let clickCount = 0; //keep track of the "state" (global)
document.querySelector('button').addEventListener('click', function() {
    if(clickCount > 10) { //decide what to do
        console.log("I think you've had enough");
    }
    else {
        clickCount++; //change state (+1)
        console.log('You clicked me!');
    }
});
```

- These “state” variables can be global, or can simply be declared within a containing function as a closure. State variables are often objects, with individual values stored as the properties. This provides a name-spacing feature, and helps to keep the code from being cluttered with many variables.

Resources

- What is the DOM? (CSS-Tricks)
- Document Object Model reference (MDN) complete DOM reference
- JavaScript HTML DOM reference (w3c)
- Introduction to Events (MDN)

Chapter 13

JavaScript Libraries

Many web programmers encounter the same programming requirements as they develop interactive web sites. For example, many websites will need to wrangle and manipulate objects and arrays, show standard components (such as modal windows or collapsible content), or do similar forms of complex DOM manipulation. Since one of the main principles of software development is *reuse*, developers will often make the code solving these problems available for others to use in the form of **libraries and frameworks**. These are publicly release script files that you can download and include in your project, allowing you to more quickly and easily develop complex applications. This is the amazing thing about the open-source community: people solve problems and then make those solutions available to others.

Modern web applications make *extensive* use of external libraries, whether by integrating many different libraries into a single app, or by relying on a particular framework (which may itself be comprised of many different libraries)! This chapter describes how to include and utilize external JavaScript libraries in your web page, and presents the popular jQuery library as an example.

- For a sense of scale, the **npm** package manager's directory lists almost *half a million* different JavaScript libraries!

Note that external scripts are generally referred to as either *libraries* or *frameworks*. However, these terms are not entirely interchangeable. A **library** is a set of behaviors (functions) that you are able to utilize and call within your code. For example, **Lodash** (described below) is a library that provides utility functions you can use. You call a library's code at your whim. A **framework**, on the other hand, provides a set of code into which you insert your *own* behaviors, either by subclassing provided components or by specifying your own callback functions. The framework calls *your* code at its whim. Martin Fowler refers to this as an "*Inversion of control*". Frameworks often seem to be easier to use (they do more with less work on your part!), but can be hard to customize

to achieve your exact goals. Libraries may be harder to use, but can likely be deployed exactly as needed. This chapter focuses mainly on *libraries*; React (detailed in later chapters) is more of a framework.

13.1 Including a Library

Just like you can include multiple CSS files in a page with multiple `<link>` elements, you can include multiple JavaScript scripts with multiple `<script>` elements.

```
<script src="alpha.js"></script> <!-- run this script first -->
<script src="beta.js"></script> <!-- run this script second -->
```

Importantly, script files are executed in the order in which they appear in the DOM. This matters because all scripts included via the `<script>` element are all run within the same namespace. This means variables and functions declared in one script file are also available in later scripts—it's almost as if all the script files have been combined into a single file.

```
/* alpha.js */
let message = "Hello World";

/* beta.js */
console.log(message); //=> "Hello World"
                        //variable was defined in previous script!
```

The order matters: if `beta.js` were included *before* `alpha.js`, then when it is run the `message` variable won't have been defined yet, causing an error.

Other techniques for managing and organizing large numbers of script files will be discussed in Chapter 15.

Like CSS frameworks, JavaScript libraries are thus included in a page by providing a `<script>` tag that references the JavaScript file that contains that library's code. For example, you can include the Lodash of helpful data processing functions:

```
<body>
  ...content

  <!-- include JavaScript files -->
  <script src="https://cdnjs.cloudflare.com/ajax/libs/lodash.js/4.17.4/lodash.min.
  <script src="js/my-script.js"></script>
</body>
```

- Like CSS frameworks, JavaScript libraries can be accessed by referencing the file on a **CDN**, by downloading the source code directly (check for the `.js` file; you may need to unpack a `.zip` file), or by installing the library

via a package manager like `npm`. Refer to Chapter 3 for details—just change `.css` to `.js`!

- Like CSS, JavaScript files can be *minimized* (usually named as `.min.js` files). Minimizing JavaScript files not only removes whitespace, but also often replaces local variable names with shorter, meaningless ones such as `a` or `x`.

You **always** want to load external libraries before your script, so that any variables and functions they define will be available to your code! If you put your script first, then those functions won't be available yet!

Because all JavaScript files share the same namespace, most popular libraries make all of their functions available as properties of a single **global variable**. For example, Lodash provides a global variable called `_` (a single underscore—a “low dash”). You use this variable to call the provided functions:

```
let array = [3,4,3];

let unique = _.uniq(array); //call function of lodash object

console.log(unique); //[3,4], the unique elements!
```

- Once the library is loaded, the globals it provide can be accessed just like built-in globals `window`, `document`, `Math`, etc.
- Lodash uses an otherwise silly variable name so that it's quick and easy to type, and to avoid conflicts with other variables you might create.

Lodash provides lots of useful functions that can be called on the global `_` object. For example:

- `_.random()` to generate a random number within a range
- `_.range()` to generate an array of numbers within a range
- `_.includes()` to determine if a value is in an array (cross-browser!)
- `_.pick()` to produce a new object that has only select properties of another
- `_.merge()` to combine the properties of two objects into a single new object

Note that exactly what global object is provided (and what functions it has) is different for every different library. You need to look through the documentation (often the `README.md` file on Github) for instructions and examples on how to call and utilize that library.

- “Knowing” a library or framework is just about being familiar with what functions you can call from it!

Pro-tip: Visual Studio Code will automatically download the type definitions for any library listed as a dependency in the project's `package.json`. For example, if you include Lodash as a dependency (by installing it with `npm`

`install --save lodash`), then just typing the `_.` will cause VS Code to provide autocomplete options. This makes it a good idea to install JavaScript libraries via `npm`, even if you plan on loading the library from a CDN instead of from a local file.

13.2 Example: jQuery

As an extended example of a JavaScript library, the remainder of this chapter describes **jQuery**—a popular JavaScript library that helps with DOM manipulation. jQuery was designed to provide easier, faster, and more reliable methods for manipulating the DOM in the ways described in the previous chapter. jQuery is one of the most popular libraries used in web development, and is used in around 70% of the most popular websites.

jQuery was developed to help fill in gaps in the JavaScript language and to standardize functionality across browsers. It made it so that you could write manipulate the DOM while writing 10% of the code, and produce scripts that would run on both Firefox *and* Internet Explorer!

However, in the 10+ years since it's release, later versions of the JavaScript language and the DOM API have improved (e.g., by introducing a `querySelector()` function) so that much of what made jQuery special is now standard and widely supported. Indeed, jQuery' approach to interactive webpages is now common practice. Nevertheless, jQuery is still commonly used in web systems (Bootstrap's JavaScript components utilize jQuery), and is popular enough that it you should be familiar with it in the likely chance that you need to engage with a web system build more than 3 years ago.

As with other libraries, you include jQuery by referencing it from a `<script>` tag in your HTML (*before your own script!*):

```
<script src="http://code.jquery.com/jquery-3.2.1.min.js"></script>
```

- Note that version jQuery 3.x is for current browsers. If you need to support an older version (e.g., IE 6-8), you can use version 1.12. Yes, jQuery supports IE 6!

Loading the jQuery library creates a **jQuery** variable in the global scope, meaning you can access this variable and utilize it in your script.

The jQuery variable is actually a *function* that you can call, called the **Selector function**. This function is used to select DOM elements by CSS selector, similar to how `document.querySelector()` works:

```
//selects element with id="foo" (e.g., <p id="foo">)
let fooElem = jQuery('#foo');
```

```
//selects all <a> elements
let allLinksArray = jQuery('a');
```

However, it is much more common to use a provided *alias* for the `jQuery()` function called `$()`. This “shortcut” lets you select elements with a single character (instead of 6)!

```
//selects element with id="foo" (e.g., <p id="foo">)
let fooElem = $('#foo');

//selects all <a> elements
let allLinksArray = $('a');
```

- Similar to Lodash’s `_`, jQuery uses `$` because no one in their right might would name a variable that, so the chance of having a namespace conflict is minimal.

Like `document.querySelector()`, the jQuery selector function handles most all selectors you are familiar with from CSS, as well as some additional useful pseudo-classes:

```
$('#my-div') // by id
$('div') // by type
$('.my-class') // by class
$('header, footer') // group selector
$('nav a') // descendant selector
$('p.red') // scoped selector

$('section:first') // first <section> element
                  // (not a css selector!)
```

Manipulating the DOM

Once you’ve selected some elements, jQuery provides *methods* that perform most of the manipulations you would do using DOM properties:

```
var txt = $('#my-div').text(); //get the textContent
$('#my-div').text('new info!'); //change the textContent

$('#my-div').html('<em>new html!</em>'); //change the HTML

$('svg rect').attr('height'); //get attribute (of all selected)
$('svg rect').attr('height',200); //set attribute (of all selected)
$('svg rect').attr({x:50, y:60}); //set multiple attributes by passing in an object

$('section').addClass('container'); //add class (to all selected)
$('section').removeClass('old'); //remove class (from all selected)
```

```

$('body').css('font-size','24px'); //set css property (of all selected)
$('body').css( {'font-size':'24px',
                'font-family':'Helvetica'} ); //set multiple properties via an object

```

If you compare this to the equivalent operations using the DOM, you'll notice that (a) jQuery is much more concise. jQuery is also more powerful: for example, you can set multiple attributes with a single call to `.attr()` by passing in an *object* containing the attributes you wish to set (the keys are the attribute names). This lets you easily change lots of attributes at once!

Importantly, note that all of these methods apply the change to *all* elements selected by the jQuery selector function. You do not need to use a loop to apply changes to multiple elements; you can just select the elements you wish to modify and change them all at once. This does mean that you may need to be careful about your selector—only select the elements you actually wish to work with!

jQuery also provides methods that allow you to manipulate the DOM tree (e.g., to add, move, or remove elements):

```

//create an element (not yet in the DOM)
let newP = $('<p class="new">This is a new element</p>'); //notice the tag!

//add content to DOM
$('main').append(newP); //add the element INSIDE <main>, at end
$('main').append('<em>new</em>'); //can a create element on the fly!

$('main').prepend('<em>new</em>'); //add new <em> element INSIDE <main>, at beginning

$('main').before('<header>'); //insert new <header> BEFORE <main> (older sibling)
//notice you can omit the closing tag if no content

$('footer').insertAfter('main'); //insert selected (<footer>) AFTER parameter (<main>)
//since the <footer> was selected, it will move!

$('main').wrap('<div class="container"></div>'); //surround the <main> with a .container

$('footer').remove(); //remove selected <footer> element
$('main').empty(); //remove all child elements of <main>

```

The first important thing to note is that you create new elements by providing the HTML content (including the tag `<>`) to the `$()` function. This is distinct from `document.createElement()`, which explicitly does *not* include the `<>` angle brackets. If creating an element with no content, you can even just specify the start tag!

jQuery also provides much more powerful manipulation methods, allowing you

to easily position elements inside, outside, around, and in place of other elements.

Although jQuery will allow you to create arbitrarily complex HTML elements to add to the DOM, you should **avoid** writing large chunks of HTML inside your `.js` file to append via jQuery (e.g., don't add an entire DOM tree at once!). This violates the separation of concerns, and makes your code difficult to read, interpret, and modify (because your HTML is now in two places!). If you do need to dynamically insert large amounts of “hard-coded” HTML, you should instead write that content inside the `.html` file, make it invisible (e.g., `display:none`), and then use jQuery to move and show the element when needed.

Handling Events

jQuery also provides *convenience methods* for registering event listeners:

```
$('#button').click(function(event) {  
    console.log('clicky clicky!');  
  
    //who was clicked  
    let element = $(event.target);  
});
```

- There are equivalent methods for other events: `.mousedown()`, `.keypress()`, etc. If you want to listen for an event that jQuery doesn't provide a convenience method for, you can use the `.on()` function:

```
$('#search-input').on('input', callback);
```

By using a method such as `.click(callback)` rather than `.addEventListener('click', callback)` or `event.on('click', callback)`, you avoid potential hard-to-catch bugs that may get introduced from misspelling `'click'`—rather than having the browser listen for `'clik'` events and then never seeming to respond to your actions, it will instead report an error (`.clik()` is not a known function!)

The `event` parameter passed into the `.click()` function's callback is exactly like the `event` passed to callbacks of `addEventListener()` (though jQuery standardizes cross-browser quirks)—thus you can access the source of the event with the `event.target` property. However, this property refers to a DOM element, *not* a jQuery selector. DOM elements don't support the jQuery methods described above—those are only available to “jQuery selection objects”. Thus if you want to work with a DOM element, you need to “select” it again using the `$()` function, thus providing a jQuery selection that has all the useful jQuery methods.

- Additionally, jQuery assigns the `this` variable inside an event callback to the event's target; this you can equivalently use `$(this)` to select the element.

The previous chapter noted the `window.onload` even listener, which was used to determine when the DOM had finished loading and so was ready to be manipulated. jQuery provides a similar functionality via the `.ready()` listener (usually called on the whole document):

```
$(document).ready(function() { //this need not be an anonymous function
    //program goes here
    console.log('Hello world!');

    //...
});
```

This is a very common pattern: often with jQuery entire programs will be specified inside the `.ready()` callback, so that the `<script>` tags can be placed in the `<head>` and downloaded quickly but still run only when the DOM is available. This pattern is *so* common in fact, that the jQuery selector function can serve as a shortcut to it:

```
//equivalent to the above
$(function() {
    //program goes here
    console.log('Hello world!');

    //...
});
```

- If you pass a *function* rather than a Selector string to the `$()` jQuery function, it will be interpreted as specifying a callback function to run when the document is ready! This is something to be aware of if you look at someone else's code and see it just start with a random `$(function)`.

And more!

This is only the tip of the iceberg for what jQuery can do. For example, jQuery also provides a number of utility functions that can be called on the jQuery (or `$`) global:

```
//check if an item is in an array
$.isArray(4, [3,4,3] );
jQuery.isArray(4, [3,4,3] ); //equivalent, but maybe easier to read

//find an item in an array that matches the filter function
//this is like .filter, but works on old browsers (if right jQuery version)
```



```
$.grep( [3,4,3], function(item) {  
    return item > 3;  
});  
  
//iterate over arrays or objects -- works for either!  
$.each( [3,4,3], function(key, value) {  
    console.log('Give me a '+value);  
});  
  
$.each( {first:'Joel',last:'Ross'}, function(key, value) {  
    console.log(key+' name: '+value);  
});
```

- This can be useful, though JavaScript native functions or Lodash methods will often be easier and (computationally) faster.

But as a final fun example: jQuery also provides functions that allow you to easily produce animated effects!

```
$('#id').fadeIn(1000);    //fade in over 1 second  
$('#id').fadeOut(500);   //fade out over 1/2 sec  
$('#id').slideDown(200); //slide down over 200ms  
$('#id').slideUp(500);   //slide up over 500ms  
$('#id').toggle();       //toggle whether displayed  
  
//custom animation syntax:  
//$(selector).animate({targetProperties}, speed [, doneCallback]);  
$('#box').animate({  
    left: '500px', //make the box fly around!  
    opacity: '0.5',  
    height: '200px',  
    width: '200px'  
}, 1500);
```

Resources

- Lodash Documentation
- jQuery Documentation
- jQuery Learning Center
- jQuery Tutorial (w3c)

Chapter 14

AJAX Requests

JavaScript allows you to dynamically define the content of a web page, generating the DOM at runtime rather than in `.html` source files. One of the primary reasons we would want to dynamically produce a DOM is if the web page's content is based on some **data** that may change over time: for example, the kind of data that is available through a **Web API**. By using JavaScript to render the DOM, you can quickly produce large amounts of HTML needed to display large data sets, sure that you have up-to-date data each time the page loads, and even *automatically refresh the page content* without requiring the user to reload!

This chapter describes how to use JavaScript to dynamically send HTTP Requests to download data (without reloading the page!), as well as how to perform the **asynchronous programming** needed when working with web requests and other time-consuming operations.

Note that this lecture assumes that you have a basic familiarity with RESTful Web APIS, including how to read and access their endpoints. For a review of some of the terminology used in APIS and RESTful requests, see the INFO 201 course reader.

14.1 AJAX

As discussed in Chapter 2, you download data from the Internet by sending an **HTTP Request** and then processing the **response**. In everyday usage, HTTP Requests are normally sent by the *browser* when the user enters a URL or clicks on a link. By default, if you wanted to download new data, you'd need to have the browser send a new request, loading a new page (or reloading the current page) in order to show that result.

To make modern dynamic web pages that display new data without needing to refresh the browser, we use a technique to send HTTP Requests *from JavaScript code* rather than from the browser. This allows us to “by-pass” the browser and get new data (and change the webpage) without reloading it! This technique is referred to as **AJAX** (**A**synchronous **J**avaScript **A**nd **X**ML)—we write code that sends an “request with AJAX” or an “AJAX request”.

Fun fact: The technology used to send AJAX requests was originally developed by Microsoft in the late 90s to support their fledgling web version of the Outlook email/calendar app. The JavaScript functions used to send these requests were included in Internet Explorer as a *non-standard* feature—an example of a browser adding new functionality that it thinks will be useful but that doesn’t work on other platforms. However, AJAX quickly gained popularity (particularly when Google showed off what you could do with it via Gmail and Google Maps), and has since become a standard that is now supported by all browsers. This is how standards come into existence!

XML and JSON

AJAX is called “AJAX” because it was originally designed to request data in XML format. **XML** (**EX**tensible **M**arkup **L**anguage) is a markup language (like HTML) that is used to encode meaning in content in a format that is both human *and* computer readable. The syntax for XML is *exactly* the same as HTML: in fact, HTML can be seen as a “subset” of the language. You can think of XML as “HTML, but you get to make up your own element names!”

```
<!-- Some XML encoding information about a person -->
<person>
  <firstName>Alice</firstName>
  <lastName>Smith</lastName>
  <favorites>
    <music>jazz</music>
    <food>pizza</food>
    <numbers>
      <item>12</item>
      <item>42</item>
    </numbers>
  </favorites>
</person>
```

- The XML language does not define any particular tags the way HTML does; instead it is up to individual applications to determine what tags it will recognize and interpret (and what tags it would see as gibberish)—what is referred to as a XML Schema.

At the time AJAX was first developed, XML was the most common way of encoding generic data for transmission. And because XML is a tree of elements

just like the DOM, similar methods could be used to navigate and extract information from the tree. However, XML is a very *verbose* language: it requires a lot of characters to encode information (meaning that the amount of data being transferred is larger), and traversing an element tree requires a lot of code. As such, JavaScript developers (led by Douglas Crockford) developed an alternative language called **JSON** (**J**ava**S**cript **O**bject **N**otation) that is more compact than XML and can be *directly* parsed into JavaScript objects and arrays:

```
{
  "firstName": "Alice",
  "lastName": "Smith",
  "favorites": {
    "music": "jazz",
    "food": "pizza",
    "numbers": [12, 42]
  }
}
```

JSON format uses a syntax that is almost identical to that for defining Object literals in JavaScript, with a few key differences:

- JSON always defines an Object `{}` at the “top level”.
- JSON object **keys** (which must be strings) *must* be written in double-quotes.
- JSON **values** can only be strings, numbers, booleans (`true` or `false`), arrays (`[]`), or other objects. You cannot include a function in JSON.
- JSON objects and arrays can’t have trailing commas or other extraneous symbols—no comments!

The JavaScript language provides a global object `JSON` (like the global `Math` object) that can be used to convert from encoded *strings* of JSON content (e.g., the above code block as a single string variable `'{"firstName":"Alice"}'`) to JavaScript objects, and vice versa:

```
//convert from Object to encoded String
let personObj = {firstName:"Alice", lastName:"Smith", id:12} //JavaScript object
let personString = JSON.stringify(personObj); //turn object into JSON string
console.log(personString); //=> '{"firstName":"Alice","lastName":"Smith","id":12}'
console.log(typeof personString); //=> 'string'

//convert from encoded String to Object
let favoritesString = '{"music":"jazz", "numbers":[12,42]}'; //a string, not an object!
let favoritesObj = JSON.parse(favoritesString); //turn JSON string into object
console.log(favoritesObj); //=> { music: 'jazz', numbers: [ 12, 42 ] }
console.log(typeof favoritesObj); //=> 'object'
```

- Note that if your JSON string is not properly formatted (e.g., you’re missing a quote), the `JSON.parse()` function will throw a `SyntaxError`.

The exact error in the JSON string can be hard to find; online tools can help show the problem.

JSON has replaced XML as the encoding of choice for working with AJAX requests—however, the technique is still referred as “AJAX” (“AJAJ” isn’t as easy to say!)

14.2 Fetching Data

AJAX support is built into browsers through the included `XMLHttpRequest` global variable (the “xml http thing”). This object provides functions that allow you to send an HTTP request to the server, but the object’s API is **really complex to use**:

An example `XMLHttpRequest`

```
//create a new XMLHttpRequest object
let request = new XMLHttpRequest();

//configure it to do an HTTP GET request for some URL
request.open('GET', 'https://domain.com/data', true);

//add a listener for the "load" event (when the data has been downloaded)
request.addEventListener('load', function() {
    if (request.status >= 200 && request.status < 400) { //check response status
        let data = JSON.parse(request.responseText);
        console.log(data); //do something with the data
    }
});

//listen for "error" events if there was a network error
request.addEventListener('error', function() {
    //handle error...
})

//finally, send the request to the server!
request.send();
```

Instead of needing to understand all that code, developers tended to use functions from external libraries such as jQuery’s `$.getJSON()` or `$.ajax()`:

```
$.getJSON('https://domain.com/data', function(data) {
    //`data` is the already-parsed JSON data
    console.log(data); //do something with the data
});
```

But this requires including the jQuery library in your page, and since the need for jQuery is rapidly going away, other options are now *built in* to modern browsers. In particular, we will utilize the **fetch()** API to easily send AJAX requests for data!

fetch() is an recent standard—so recent that it is not supported by IE (any) or Safari 10. However, we can still use **fetch()** with these browsers by including a **polyfill**—an external library that replicates an existing API in platforms that don’t support it! The **fetch()** polyfill will provide a **fetch()** function to browsers that don’t provide it (leaving other browsers unchanged) that uses the existing XMLHttpRequest without you needing to interact with that object.

It’s easiest to just load the polyfill from a CDN:

```
<!-- put this BEFORE your own script! -->
<script src="https://cdnjs.cloudflare.com/ajax/libs/fetch/2.0.3/fetch.min.js"></script>
```

The **fetch()** function makes it easy to send a request: simply call it and pass in the URL of the data you wish to download:

```
fetch('https://domain.com/data');
```

In some browsers, you will not be able to send an AJAX request when the web page is loading via the **file://** protocol (e.g., by double-clicking on the **.html** file). This is a security feature to keep you from accidentally running an HTML file that contains malicious JavaScript that will download a virus onto your computer. Instead, you should run a local web server such as **live-server** for testing AJAX requests.

14.3 Asynchronous Programming

However, the **fetch()** function does **NOT** directly return the data you want to download! Downloading data off the internet can take a long time: the network connection may be slow and the amount of data to download may be quite large (metadata for the latest 100 tweets from Twitter involves almost 500k of JSON content). Because fetching data may take time, AJAX requests are made **asynchronously** (that’s the “A” in AJAX)—the download will occur *at the same time* that the rest of the code is being executed. Thus the download and the remaining script will *not* be synchronized; they will be “asynchronous”.

```
console.log('About to send request'); //statement 1

//send request for data to the url
fetch(url); //statement 2

console.log('Sent request'); //statement 3
```

```
//The data is actually received sometime later,  
//when the JS interpreter is down here!
```

- In the above example, the JS interpreter will execute statement 1, then statement 2 (the `fetch()` call). It will then precede *immediately* to statement 3 (the second `console.log`), without waiting for the request to finish! The download will continue to occur in the background, and will finish at some point later in the program—though we don’t know exactly when.
- It is best to think of `fetch()` as a function that will just “start to download data”, not one that actually downloads data!

Because `fetch()` is an **asynchronous function** (it’s code is run asynchronously), it returns what is called a **Promise**. A Promise is object that holds a value which may not be available yet—you can think of a Promise as like a placeholder where the result of the asynchronous function call will eventually be stored (it is a “promise” to eventually have some data, though that promise may be kept or broken!).

- Promises are the modern way of handling asynchronous functions, but as part of the ES6 standard they are not yet available to all browsers (specifically: Internet Explorer). So you’ll need to include *another polyfill* to support IE. This is also available from a CDN.

Promises have three possible states: *pending* (the data is downloading), *fulfilled* (the data has finished downloading), or *rejected* (the data failed to download and the promise was “broken”). We are able to specify *callback functions* (similar to event listeners) that occur when a pending Promise is either successfully fulfilled or has been rejected. The “on success” callback function is specified by calling the `.then()` function on the Promise object, and passing the “on success callback” as a parameter:

```
function onSuccessCallback(response) { //what to do when we get the response  
    console.log(response);  
}  
  
//When fulfilled, execute the callback function (which will be passed the response)  
let promise = fetch(url);  
promise.then(onSuccessCallback);  
  
//It is much more common to use anonymous variables/callbacks:  
fetch(url).then(function(response) {  
    console.log(response);  
});
```

The “on success” callback will be passed a single parameter: the **data value** that the Promise was made for (e.g., the data that will eventually be downloaded

from `fetch()`). So when the callback is executed, you will have access to the data! For example, when the `fetch()` Promise is fulfilled, it will pass an object representing the *response* to the HTTP Request:

```
let promise = fetch(url);
promise.then(function(response){
    console.log( response.url ); //a string of where the request was sent
    console.log( response.status ); //the HTTP status code (e.g., 200, 404)
});
```

This response object does have a `body` property that represents the “body” (data content) of the HTTP response. However, that body stored as a “stream” of 0s and 1s, not as a JavaScript object (or even a string you can `JSON.parse()`)! In order to get the body into a format you can use, you will need to “encode” it into a JavaScript object by calling the `.json()` method on it.

- There is also an equivalent `.text()` method to encode a response body into plain text.

Chaining Promises

But there’s a catch: the “encoding” process performed by the `.json()` might take some time (particularly for a large amount of data). So instead of blocking (pausing) the rest of your program while that encoding occurs, the `.json()` method returns *another Promise* as a placeholder for when the encoded body is available! So you will then need to specify a `.then()` callback for *that* Promise as well.

However, a Promise’s `.then()` function has a neat property that makes this easy to do. Calling the `.then()` function on a Promise returns a new Promise as a placeholder for any data produced by the `.then()` function. This promised data will be whatever value is *returned* by the “on success” callback function. This allows you to in effect “chain” `.then()` calls together, each of which can perform some kind of transformation on the data:

```
function makeQuestion(dataString) { //a function to make a string a question
    return dataString + '???';
}

//image a hypothetical asynchronous function `getAsyncString`
//it returns a Promise (placeholder) for a string load from a given source
let originalPromise = getAsyncString(myDataSource);

//when the original promise is fulfilled, call `makeQuestion` on it
//`questionPromise` will be a placeholder for that transformed data
let questionPromise = originalPromise.then(makeUpperCase);
```

```
//when the `questionPromise` is fulfilled, call an anonymous callback on it
//the callback will be passed the transformed ("question") data
questionPromise.then(function(data){
    console.log(data); //data will be a question!
})
```

More commonly, we use *anonymous variables* for subsequent promises, allowing you to chain them together in a way that almost reads like English!

```
getAsyncString()
    .then(makeQuestion)
    .then(function(data){
        console.log(data);
    });
```

But wait there's more! `.then()` also has a special feature where if the “on success” callback function returns a *Promise* (rather than another kind of value), then the “outer” promise will take on the state of that new returned Promise. This means that you can just *return* a Promise from inside a `.then()` callback, and that Promise will be the subject of the subsequent `.then()` call:

```
let outerPromise = getAsyncString(myFirstSource).then(function(firstData){
    //do something with `firstData`

    let newPromise = getAsyncString(mySecondSource); //a second async call!
    return newPromise; //return the promise.
}); //`outerPromise` now takes on the state and data of `newPromise`

outerPromise.then(function(secondData){
    //do something with `secondData`, the data downloaded from `mySecondSource`
});
```

Going back to `fetch()` to bring it all together: since the `.json()` encoding function returns a Promise, you can simply *return* that Promise from the `.then()` callback in order to make it available to subsequent `.then()` calls!

```
fetch(url) //start the download
    .then(function(response) { //when done downloading
        let dataPromise = response.json(); //start encoding into an object
        return dataPromise; //hand this Promise up
    })
    .then(function(data) { //when done encoding
        //do something with the data!!
        console.log(data); //will now be encoded as a JavaScript object!
    });
```

This code example will allow you to download data and encode it into a plain old JavaScript object that you can work with.

Handling Errors

When downloading data from the internet, it is always possible that the HTTP request may fail. The request may be sent to the wrong URL, the client computer may be having connection problems, or the receiving server may be having problems.

In order to deal with inevitable errors, Promises provide a `.catch()` method that is used to specify a callback that should occur if the Promise is *rejected* (an error occurs). This callback function will be passed an Error object that contains details about the error.

```
fetch(url)
  .then(function(response) { //when done downloading
    return response.json(); //second promise is anonymous
  })
  .then(function(data) { //when done encoding
    //do something with the data!!
    console.log(data); //will now be encoded as a JavaScript object!
  })
  .catch(function(err) {
    //do something with the error
    console.error(err); //e.g., show in the console
  });
```

- Importantly, the `.catch()` method will “catch” errors from *all previous* Promises in a `.then()` chain! This means that the above `.catch()` will show both errors in the downloading (`.fetch()`), and errors in the body encoding (`.json()`).
- You will almost always want to show the error to the user in some way, such as by creating an alert element in the DOM.
- The `.catch()` function *also* returns a Promise, so you can continue to chain `.then()` calls after it. These later callbacks will only be executed if no previous Promise has been rejected (that is, there haven’t been any errors yet).

Important: a Promise will only be rejected if there is an actual “Error” in sending the request. If the server replies with a 401 error (e.g., you didn’t have permission to access the resource) or just the message “invalid API key”, that won’t be handled by `.catch()`. From JavaScript’s perspective, the request went through perfectly—it’s not `fetch`’s fault that the data you asked for wasn’t what you actually wanted!

- You can use the `response.status` and `response.ok` properties to check the status of the HTTP response.

As such, you will want to make sure to handle things like bad responses or

unexpected response bodies, both in testing your application (to make sure the request is sent to the correct URL) and when handling any user input.

Resources

- [An Introduction to AJAX for Front-End Developers](#) (tuts+)
- [An Introduction to `fetch\(\)`](#) (Google)
- [JavaScript Promises: an Introduction](#) (Google)

Chapter 15

ES6 and Beyond

Coming soon...